

图数据库解谜

搞清楚什么是图数据库，它能做什么，为什么它做的好？如何学用大厂们都选择了的 Nebula Graph？

古思为

DEVELOPER ADVOCATE @  vesoft

 墨天轮 线上直播
Dec. 9th, 2021

古思为

- Nebula Graph 开发者布道师
- 程序员
- 开源信徒



❏ [wey-gu](#)

❏ [wey_gu](#)

❏ [siwei.io](#)

Overview

- 什么是图？图数据库？
- 为什么需要图数据库？
- 到底图数据库可以做什么？
- 为什么大厂都在用的新一代开源图数据库：Nebula Graph

Overview

- 什么是图？图数据库？
- 为什么需要图数据库？
- 到底图数据库可以做什么？
- 为什么大厂都在用的新一代开源图数据库：Nebula Graph

 [wey-gu](#)

 [wey_gu](#)

 [siwei.io](#)

图数据库简介

什么是图？ 什么是图数据库？ 为什么我们需要一个专门的数据库？



什么是图？



"以图结构、图语义来用点、边、属性来查询、表示存储数据的数据库

什么是图数据库

[wikipedia.org/wiki/graph_database](https://en.wikipedia.org/wiki/Graph_database)

了解更多关于 [什么是图数据库](#)

为什么需要图数据库？

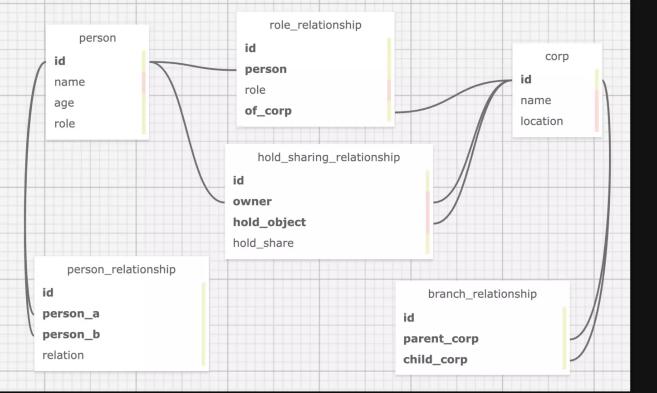
传统数据库

图数据库

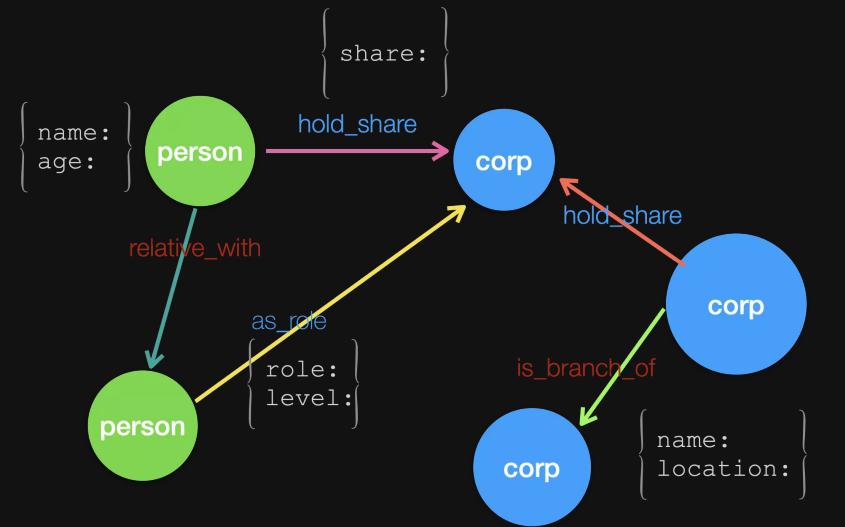
为什么需要图数据库？

传统数据库

图模型的结构



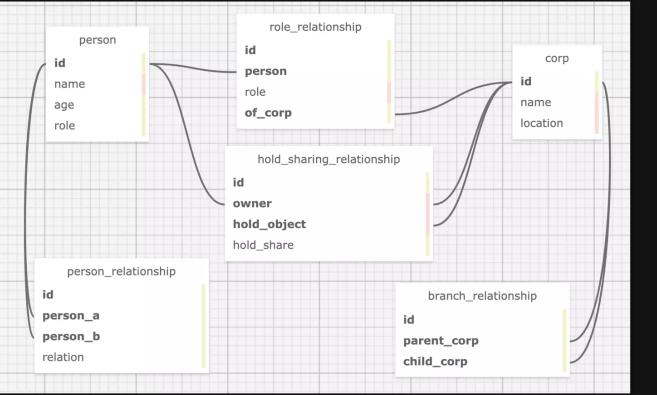
图数据库



为什么需要图数据库？

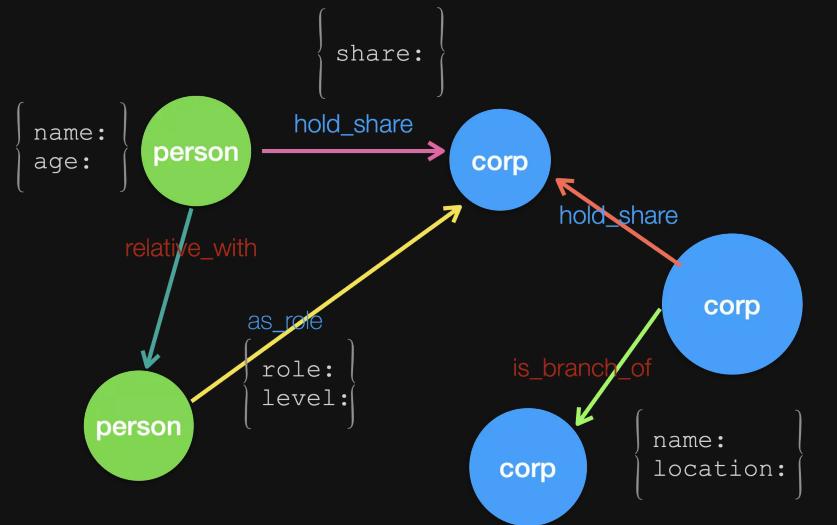
传统数据库

图模型的结构



图数据库

图语义的查询



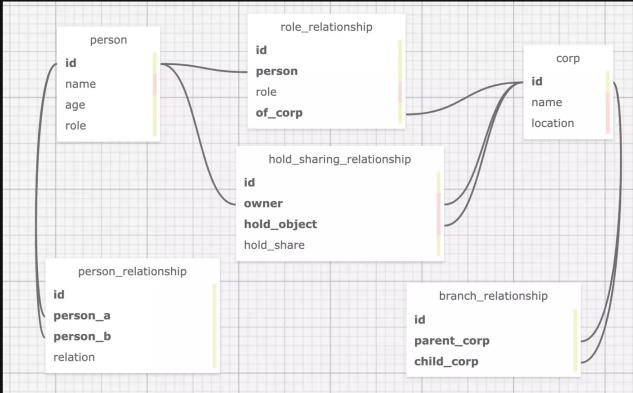
```
SELECT a.id, a.name, c.name
FROM player a
JOIN serve b ON a.id=b.player_id
JOIN team c ON c.id=b.team_id
WHERE c.name IN (SELECT c.name
FROM player a
JOIN serve b ON a.id=b.player_id
JOIN team c ON c.id=b.team_id
WHERE a.name = 'Tim Duncan')
```

```
GO FROM 100 OVER serve YIELD serve._dst AS Team | \
GO FROM $-.Team OVER serve REVERSELY YIELD $$ .player.name;
```

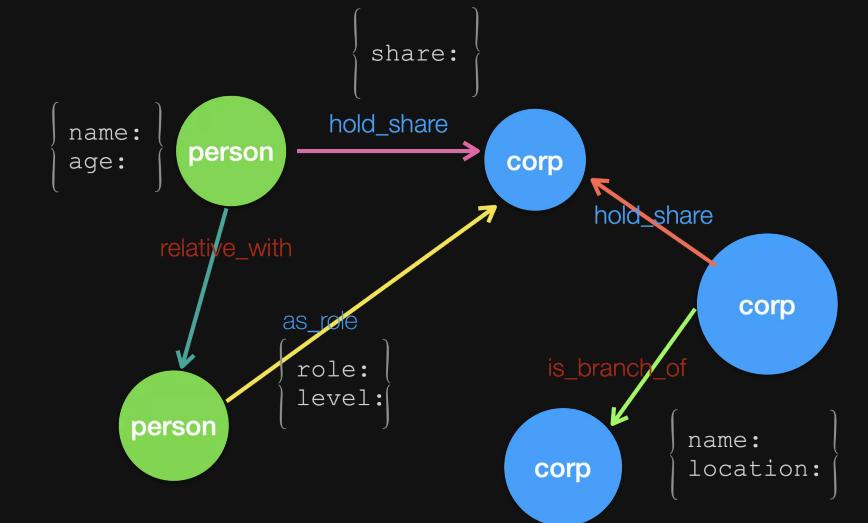
为什么需要图数据库？

传统数据库

图模型的结构



图数据库



图语义的查询

```
SELECT a.id, a.name, c.name
FROM player a
JOIN serve b ON a.id=b.player_id
JOIN team c ON c.id=b.team_id
WHERE c.name IN (SELECT c.name
FROM player a
JOIN serve b ON a.id=b.player_id
JOIN team c ON c.id=b.team_id
WHERE a.name = 'Tim Duncan')
```

```
GO FROM 100 OVER serve YIELD serve._dst AS Team | \
GO FROM $-.Team OVER serve REVERSELY YIELD $$ .player.name;
```

性能

	主要应用场景	2-hop 延时 (~2.5K)	3-hop 延时 (~110K)	4-hop 延时 (~600K)
图数据库	关系遍历	0.01 秒	0.168 秒	1.36 秒
SQL数据库	信息检索	0.016 秒	30 秒	1544 秒

Nebula Graph!

如何发音：['nebjələ]，它有哪些特点？

Nebula Graph 介绍

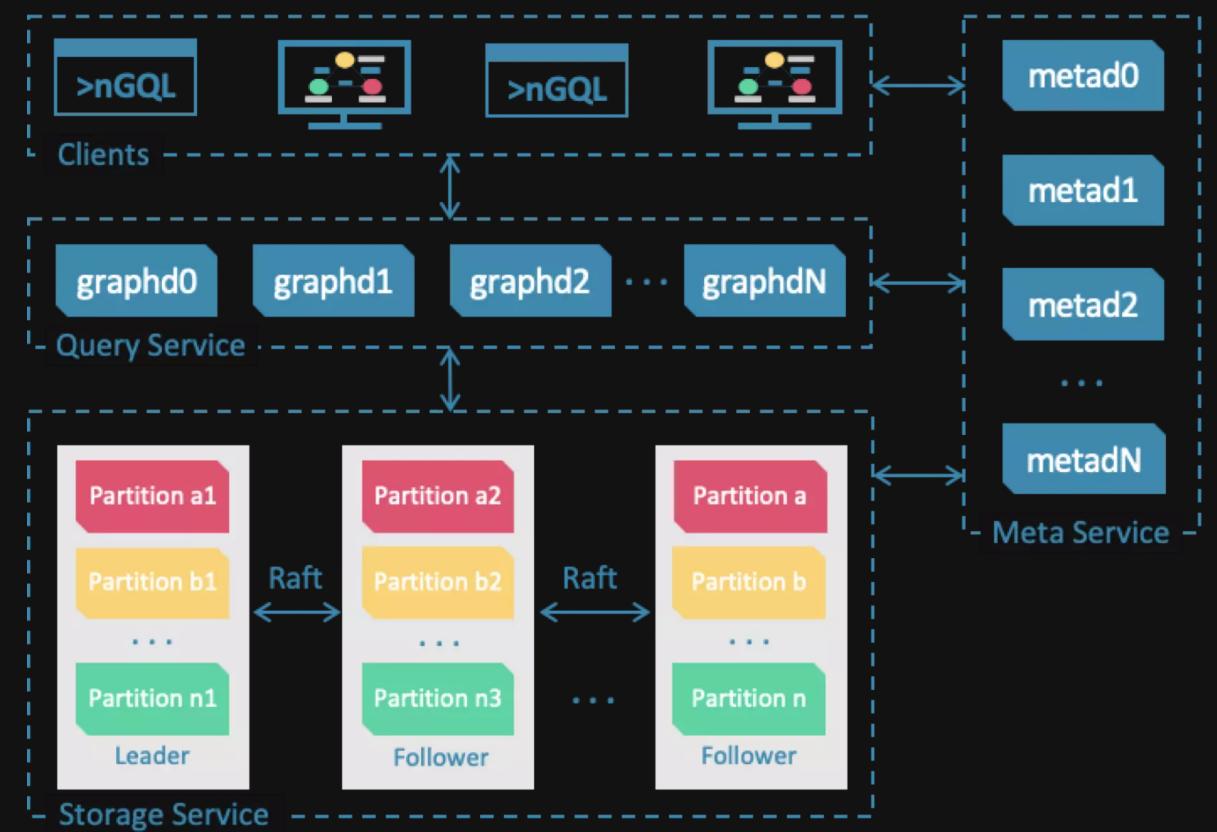
一个可靠的分布式、线性扩容、性能高效的图数据库

世界上唯一能够容纳千亿顶点和万亿条边，并提供毫秒级查询延时的图数据库解决方案

Nebula Graph 介绍

一个可靠的分布式、线性扩容、性能高效的图数据库

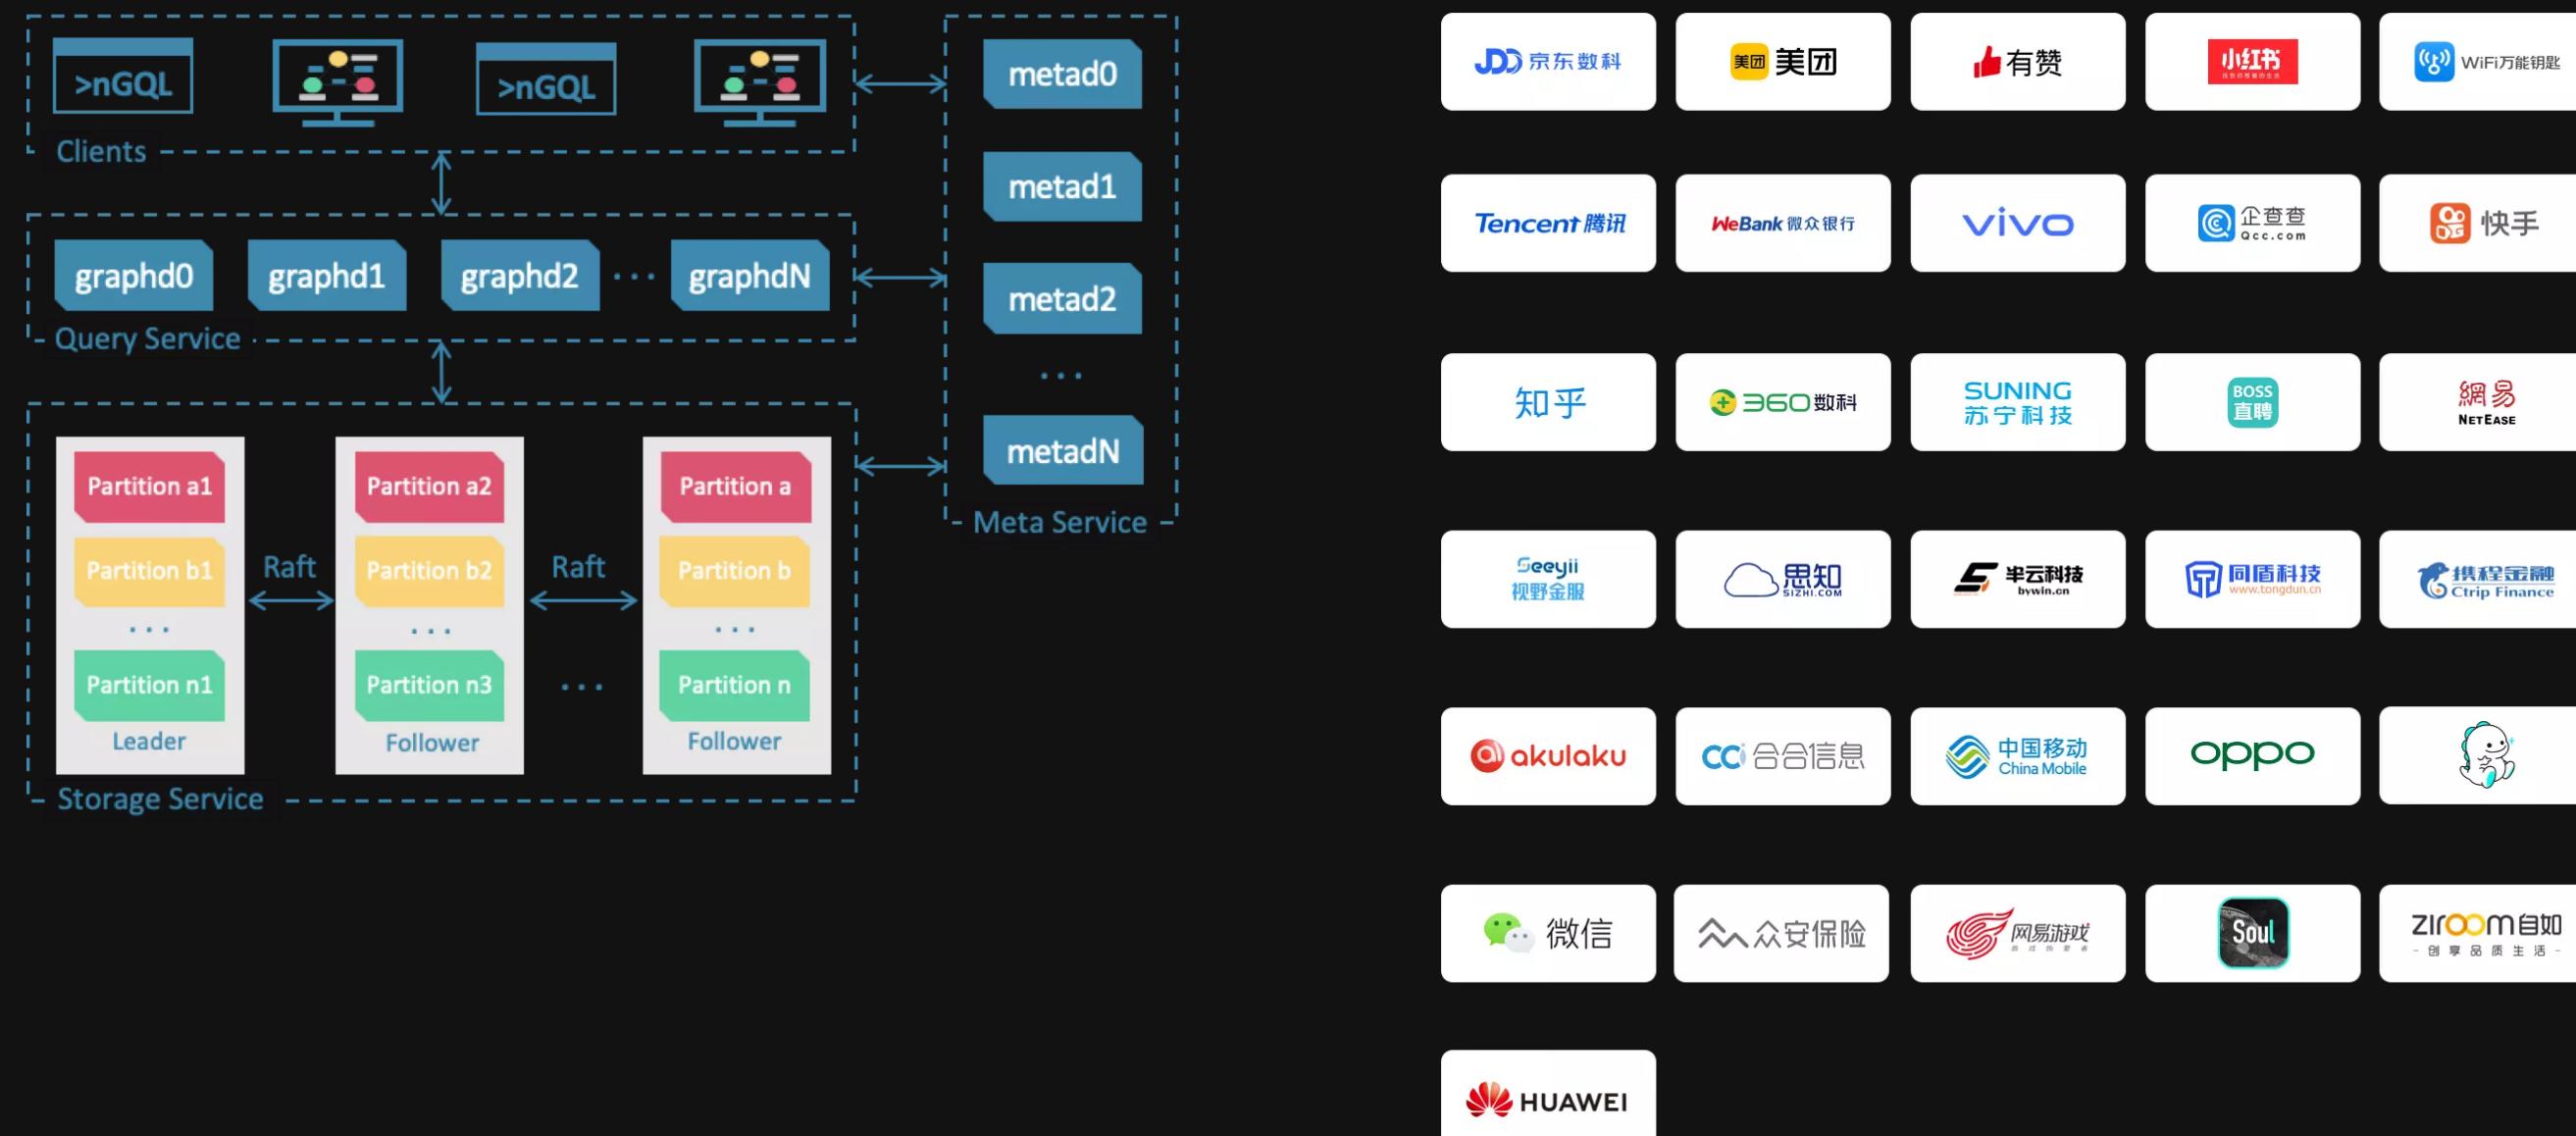
世界上唯一能够容纳千亿顶点和万亿条边，并提供毫秒级查询延时的图数据库解决方案



Nebula Graph 介绍

一个可靠的分布式、线性扩容、性能高效的图数据库

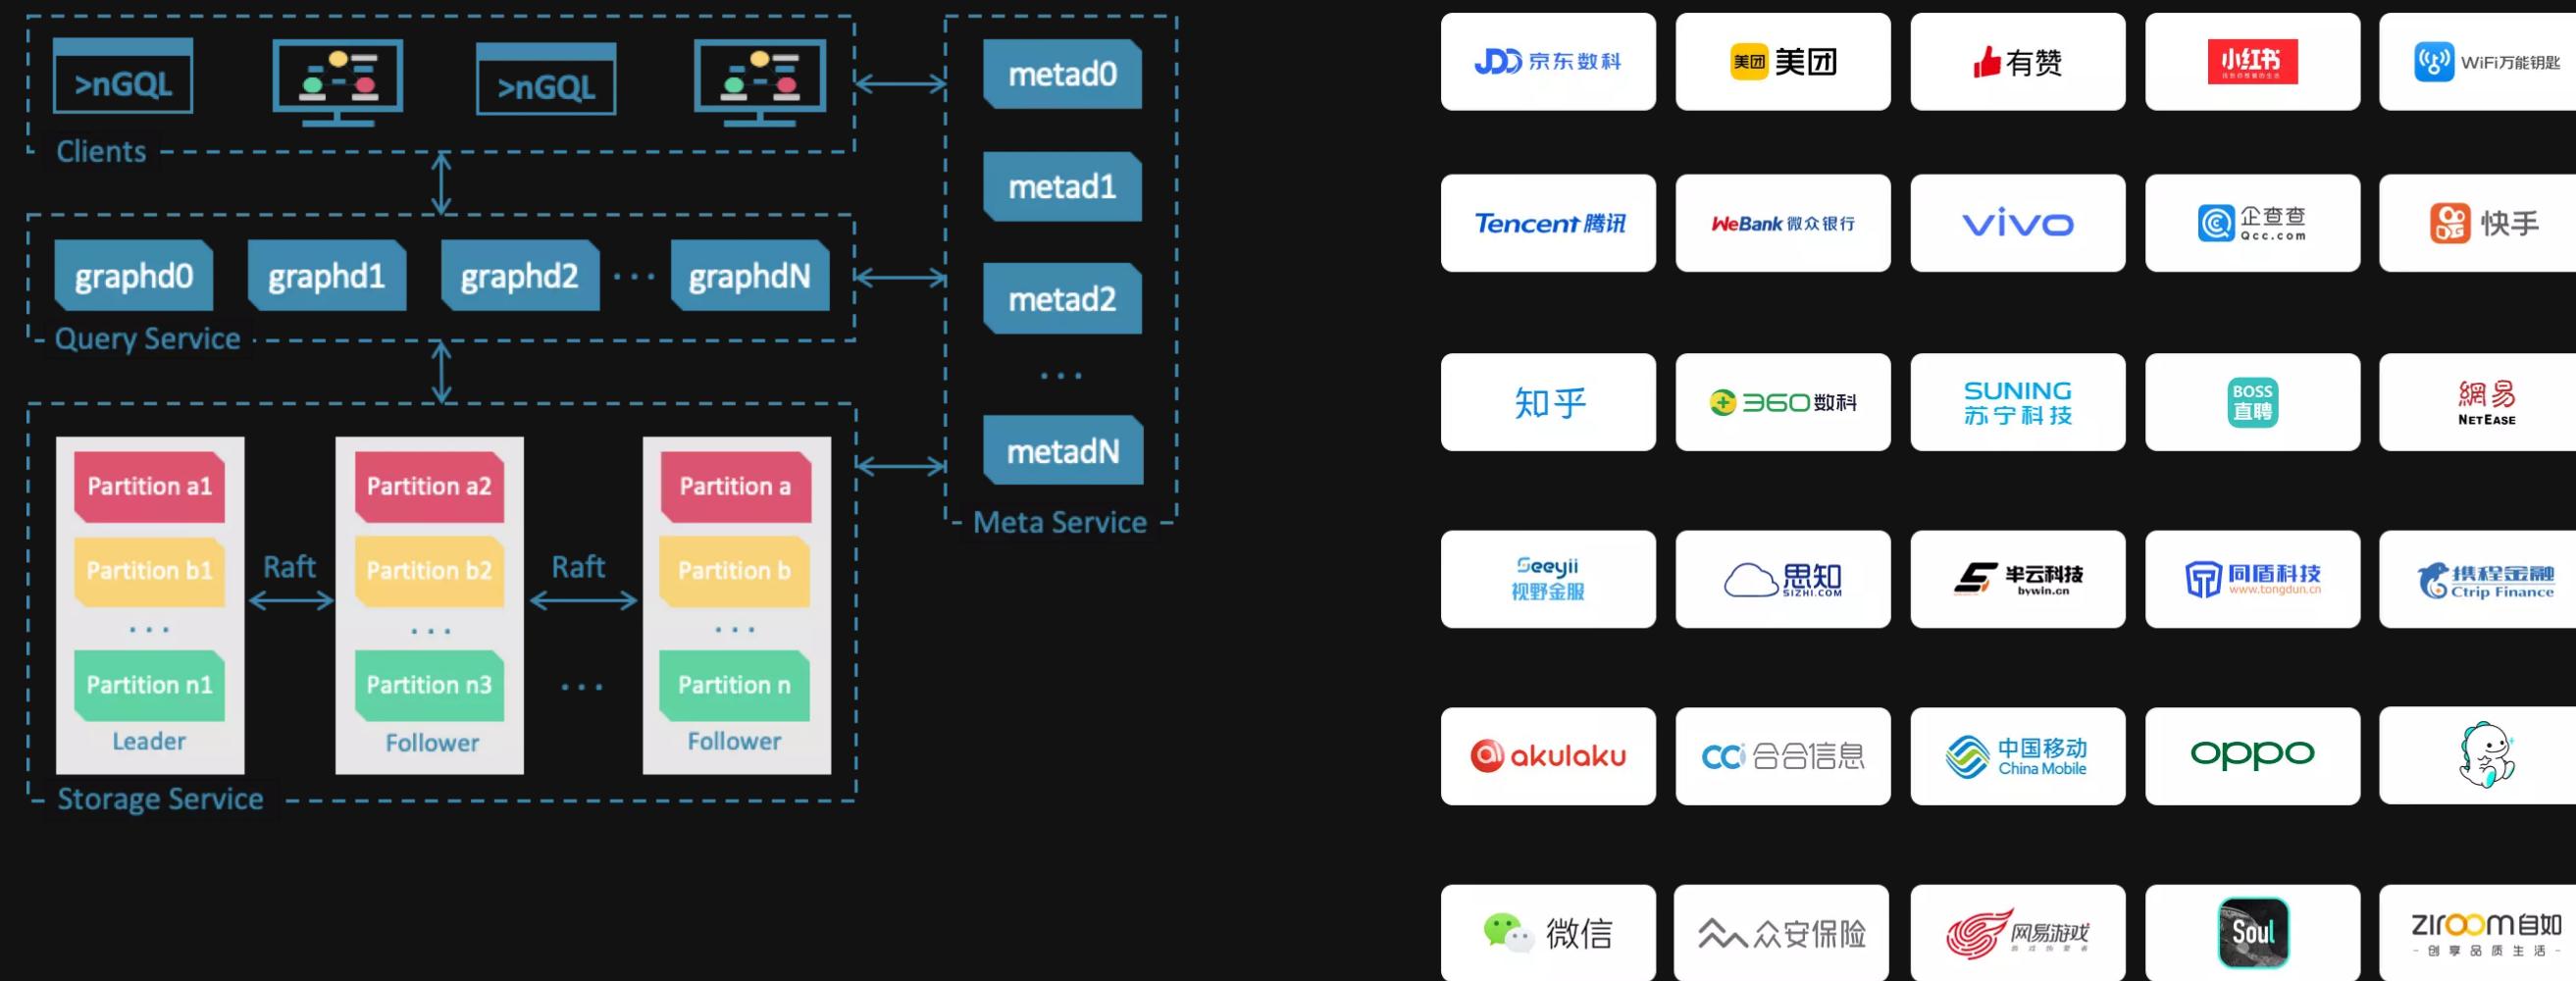
世界上唯一能够容纳千亿顶点和万亿条边，并提供毫秒级查询延时的图数据库解决方案



Nebula Graph 介绍

一个可靠的分布式、线性扩容、性能高效的图数据库

世界上唯一能够容纳千亿顶点和万亿条边，并提供毫秒级查询延时的图数据库解决方案



了解更多 >>>

文档: Nebula 架构

官网: 用户案例

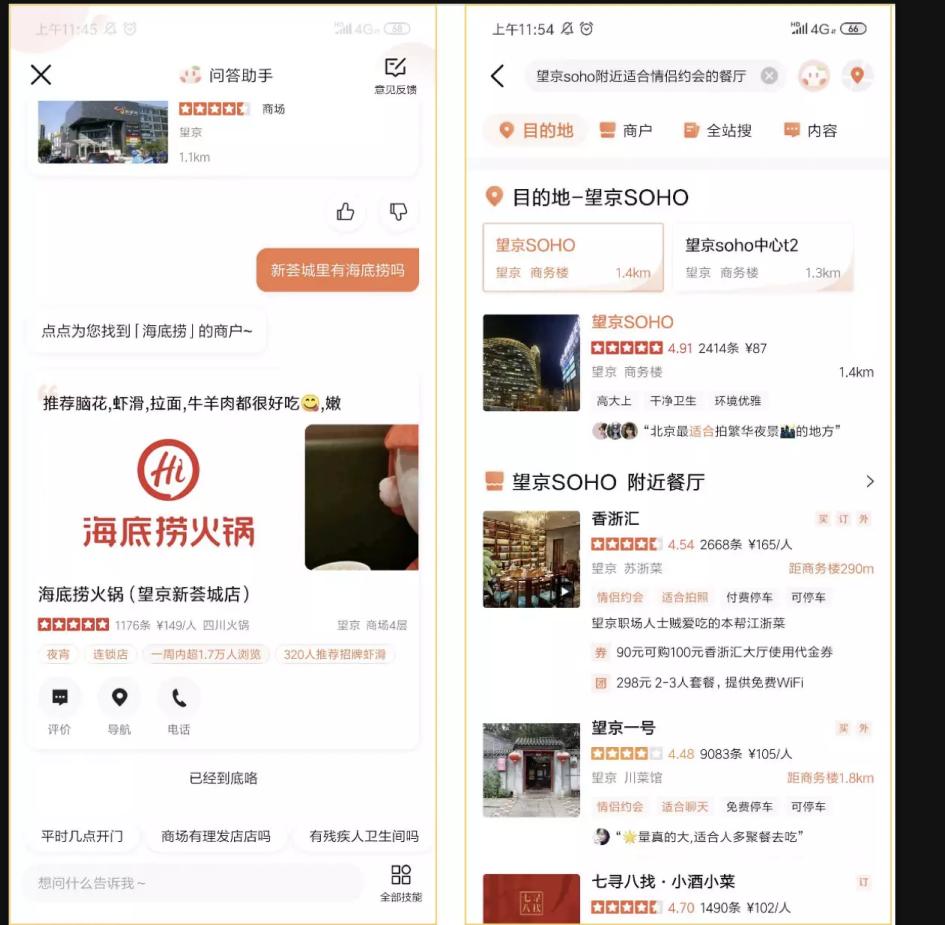


一个贴切的小 Demo

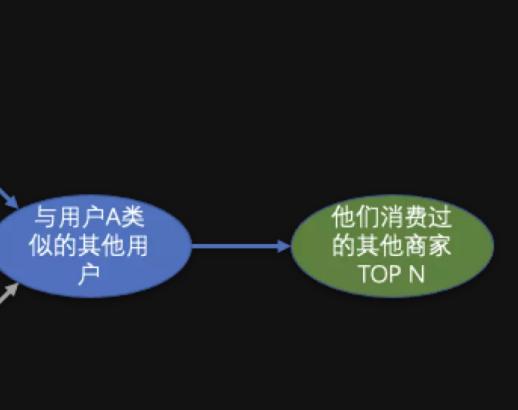
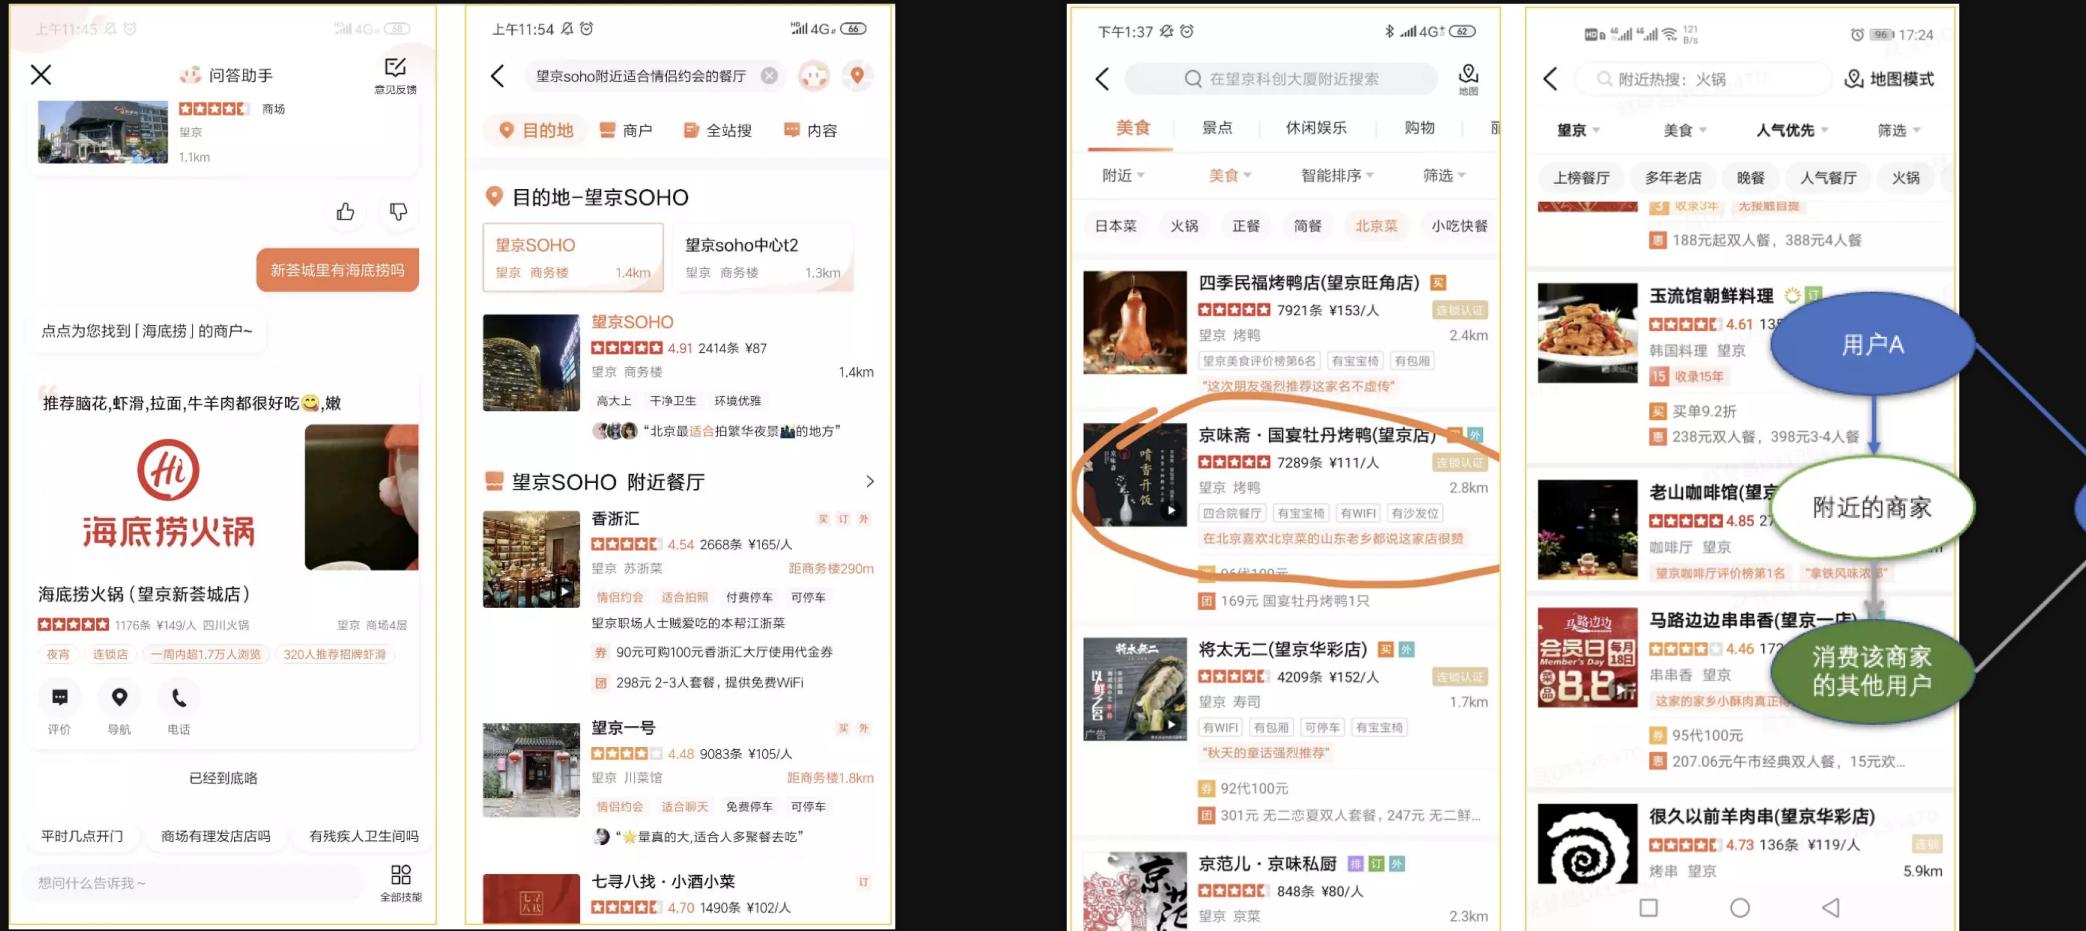
关于漫威电影宇宙 (MCU) 与姚明的故事

图数据库的应用场景

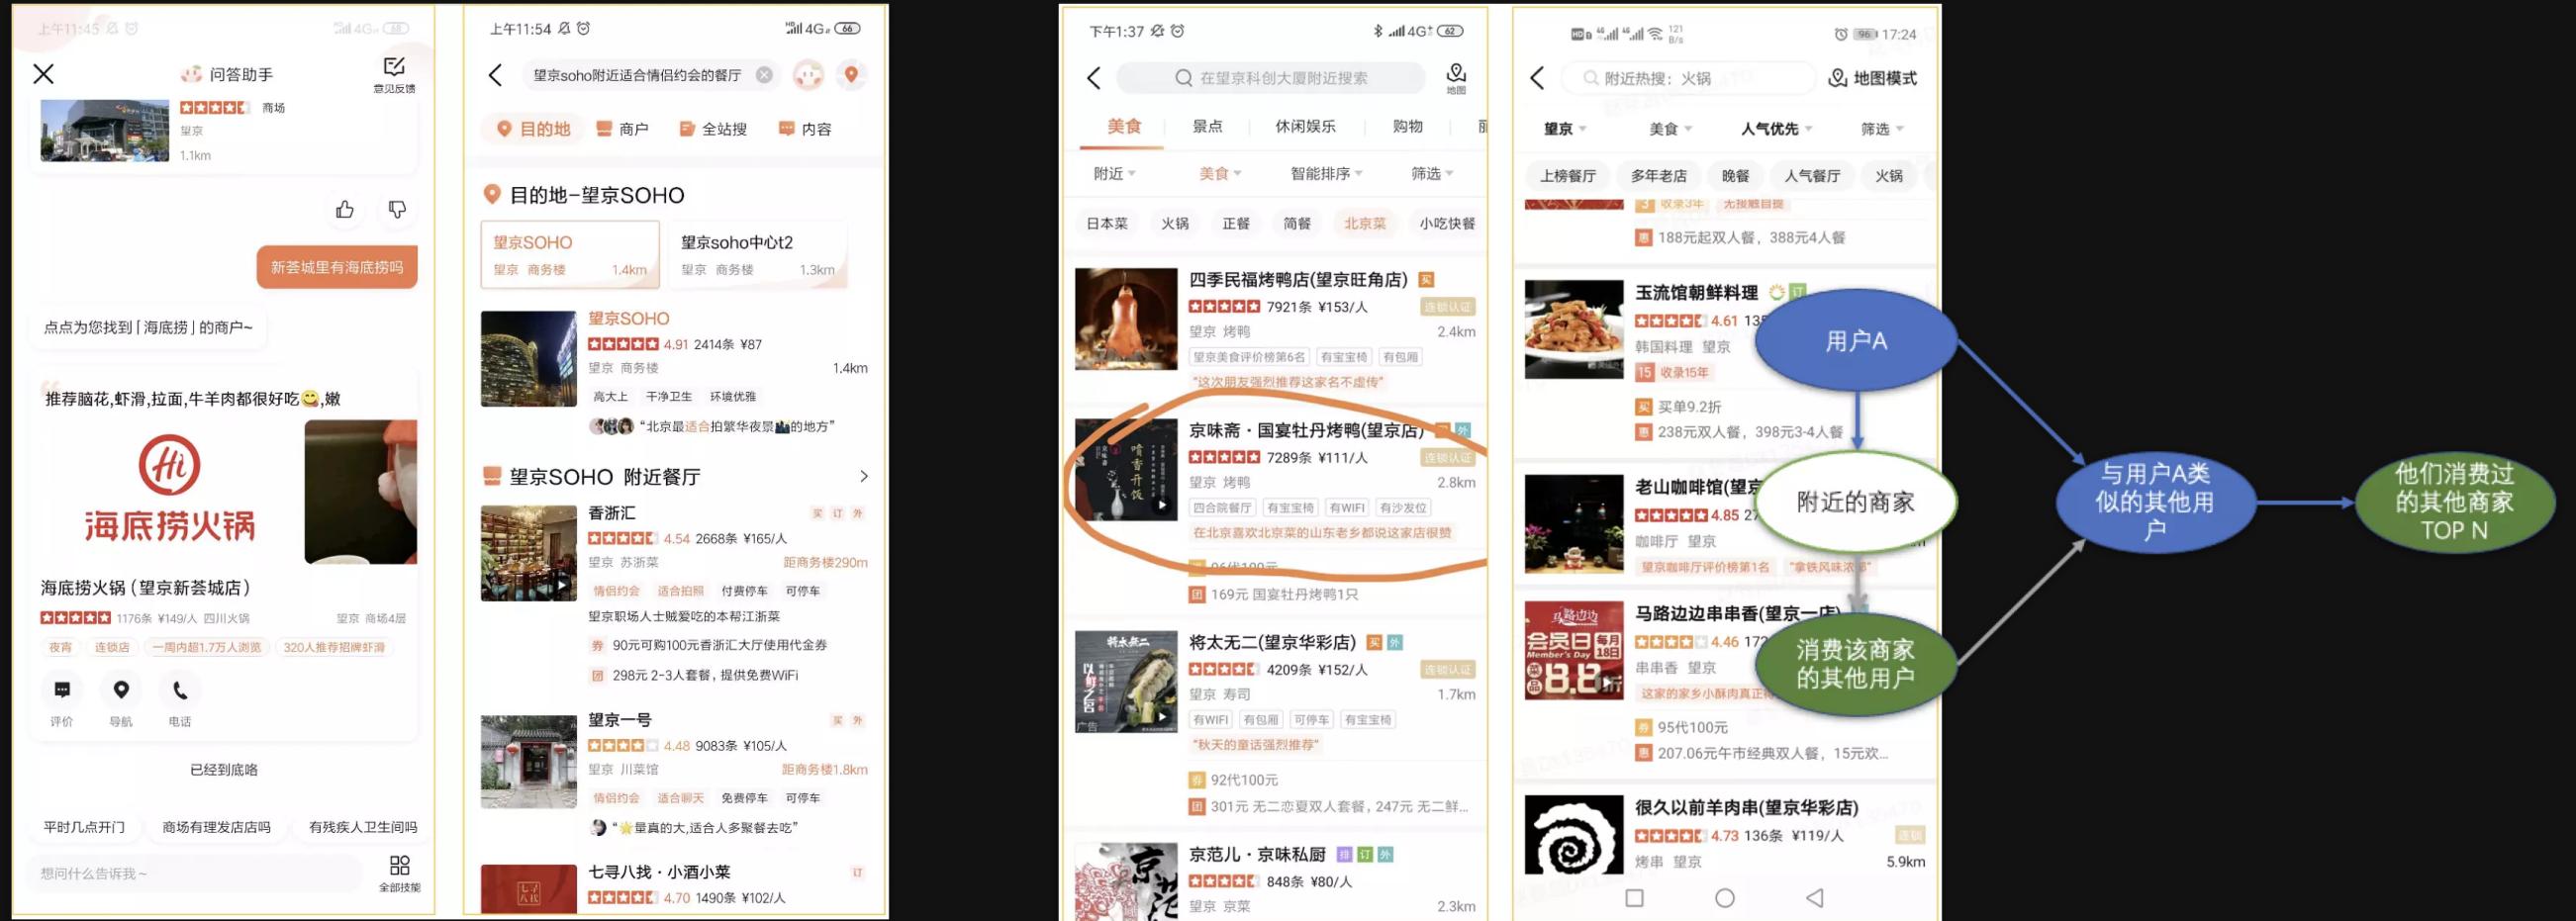
图数据库的应用场景



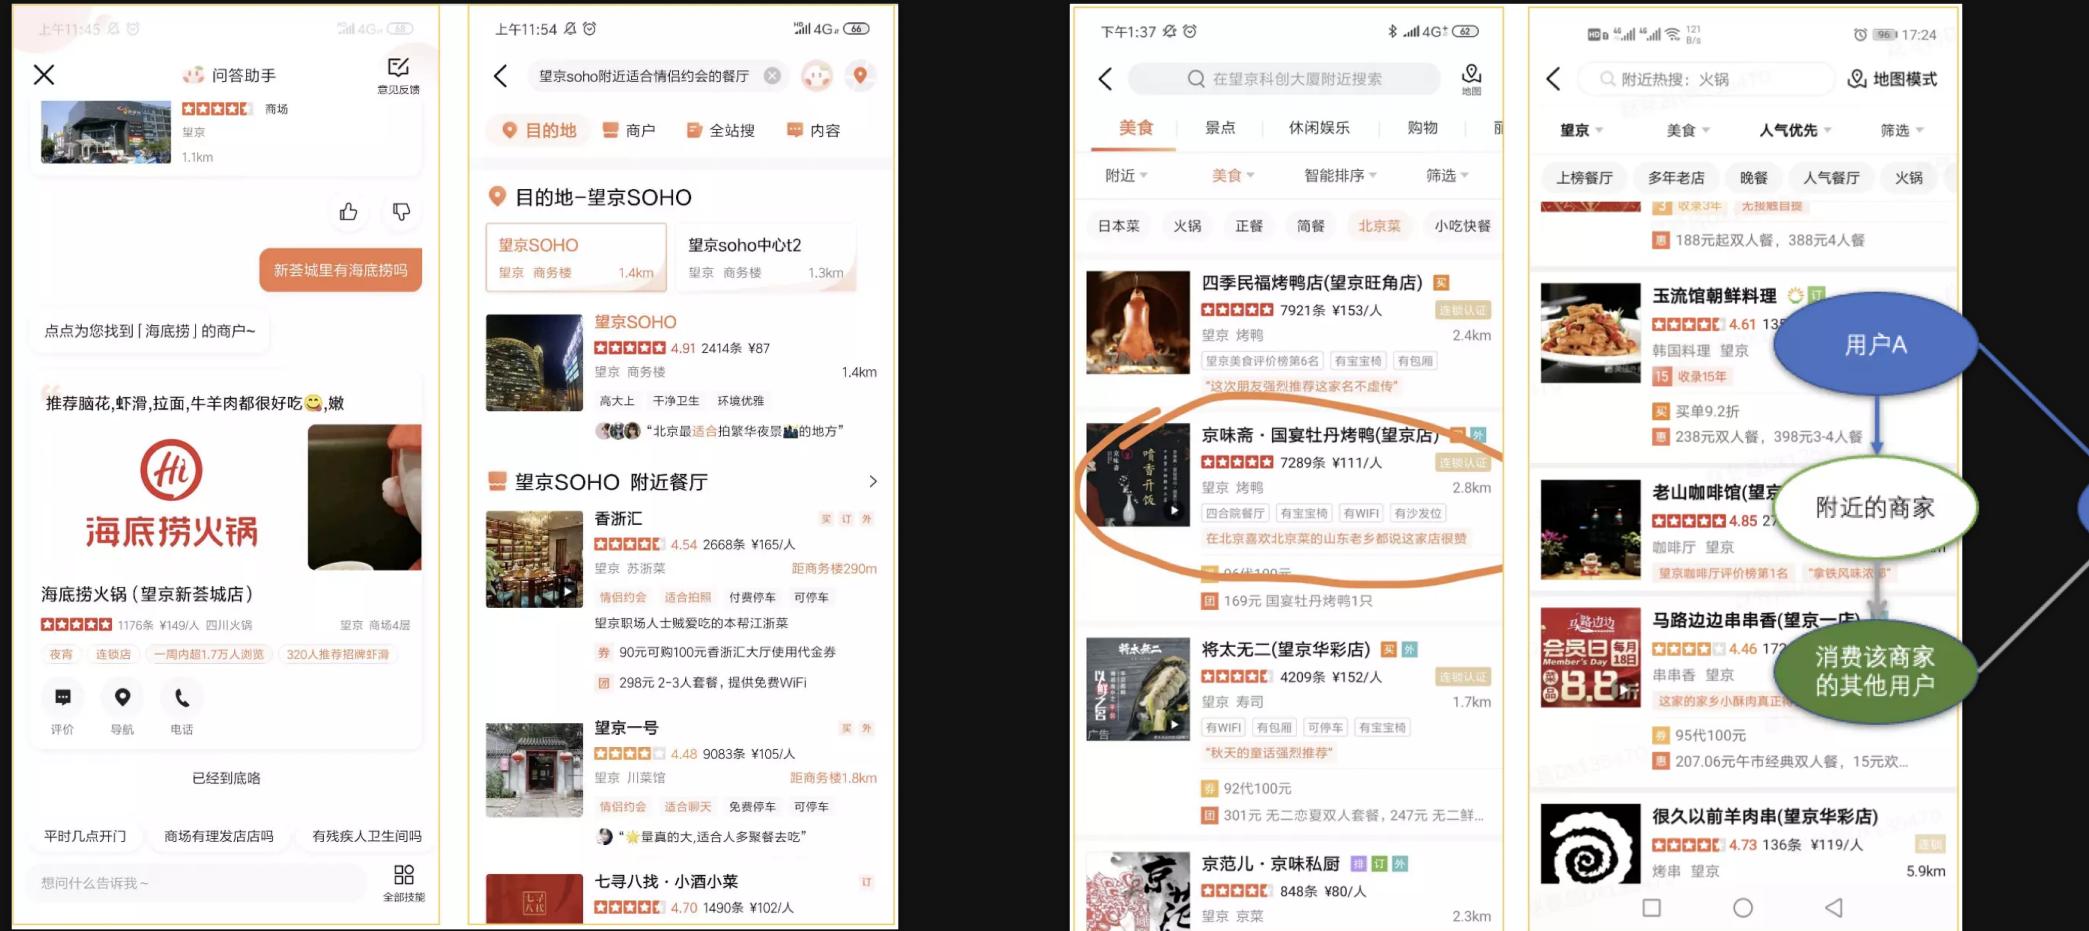
图数据库的应用场景



图数据库的应用场景



图数据库的应用场景



典型场景

社交网络(meta)

风险控制

公共安全

知识图谱

机器学习

生化制药

物联网

区块链

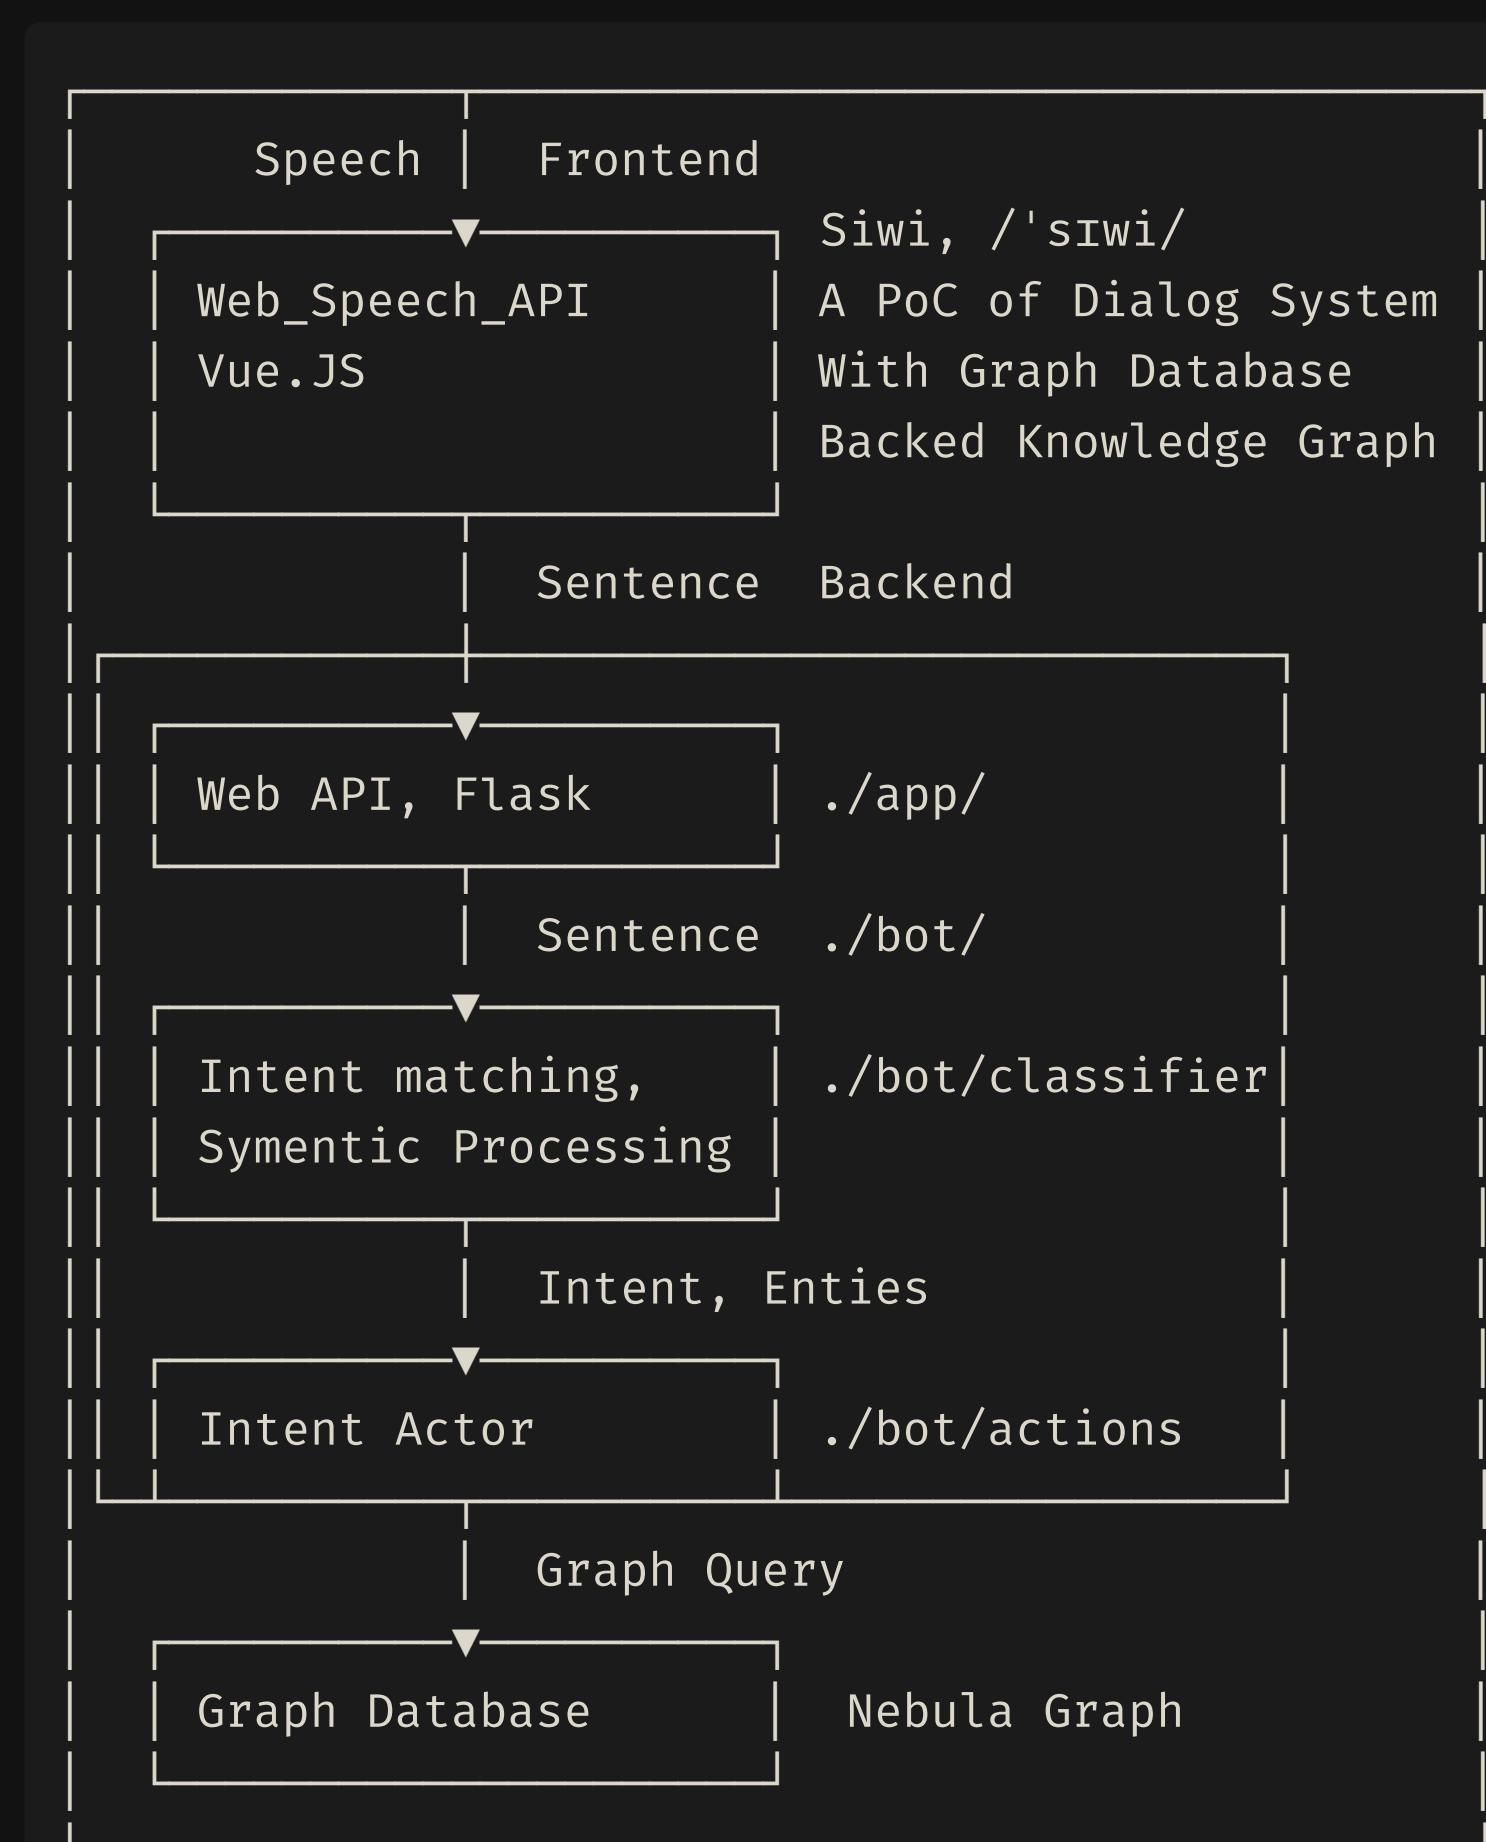
数据血缘

智能运维

Siwi

Siwi (/sIwi/), 一个知识图谱驱动的智能语音助手

Arch



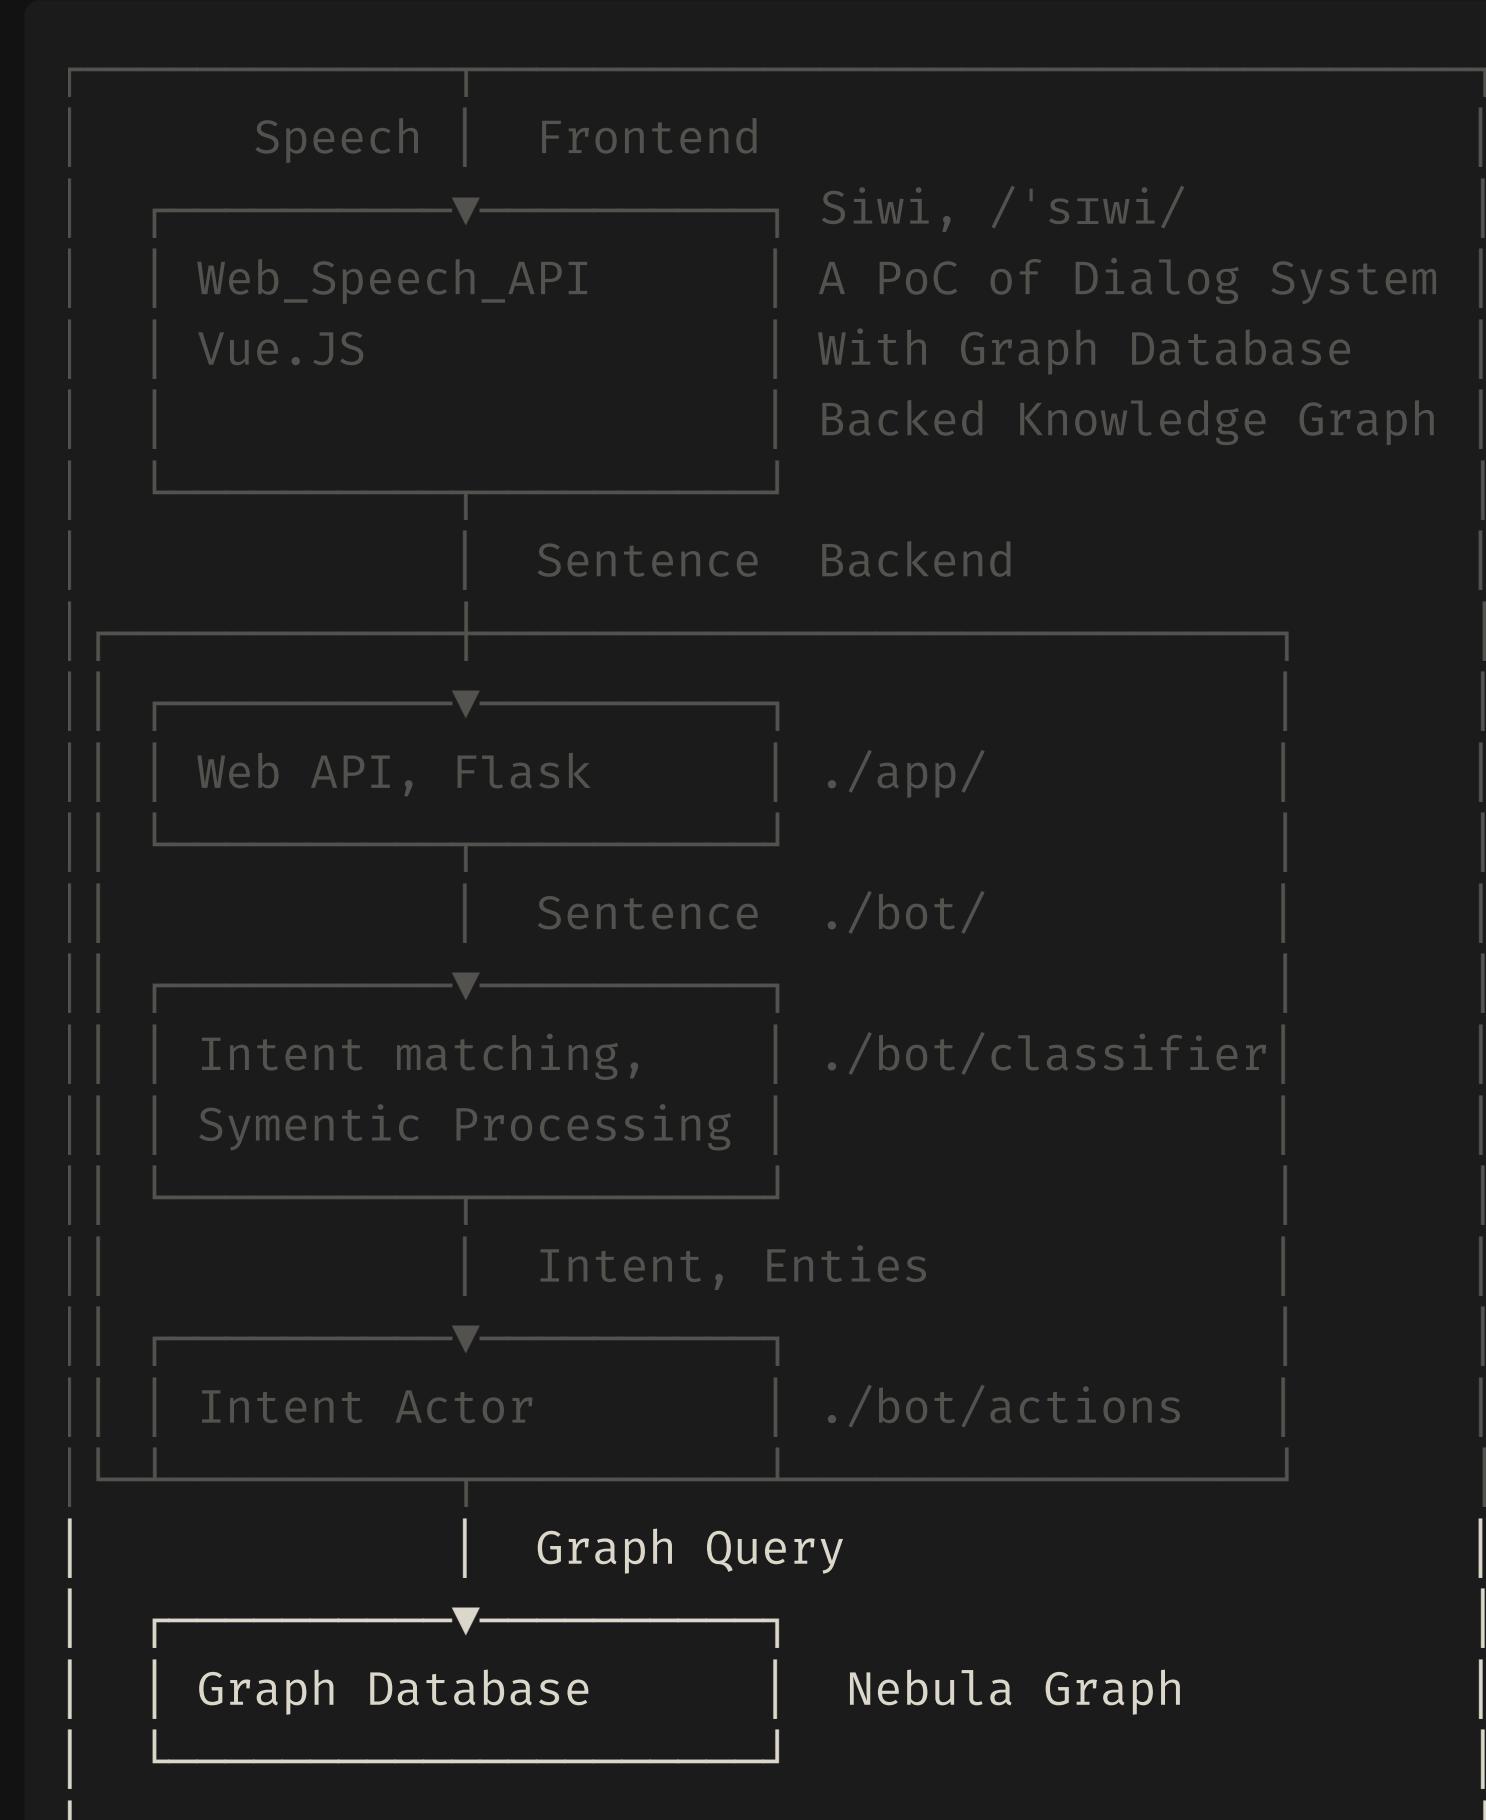
Code

```
.├── README.md  
├── src  
│   └── siwi  
│       ├── app  
│       └── bot  
│           ├── actions  
│           ├── bot  
│           ├── classifier  
│           └── test  
└── siwi_frontend  
    ├── README.md  
    ├── package.json  
    └── src  
        ├── App.vue  
        └── main.js  
wsgi.py
```



[wey-gu/nebula-siwi](https://github.com/wey-gu/nebula-siwi)

Arch



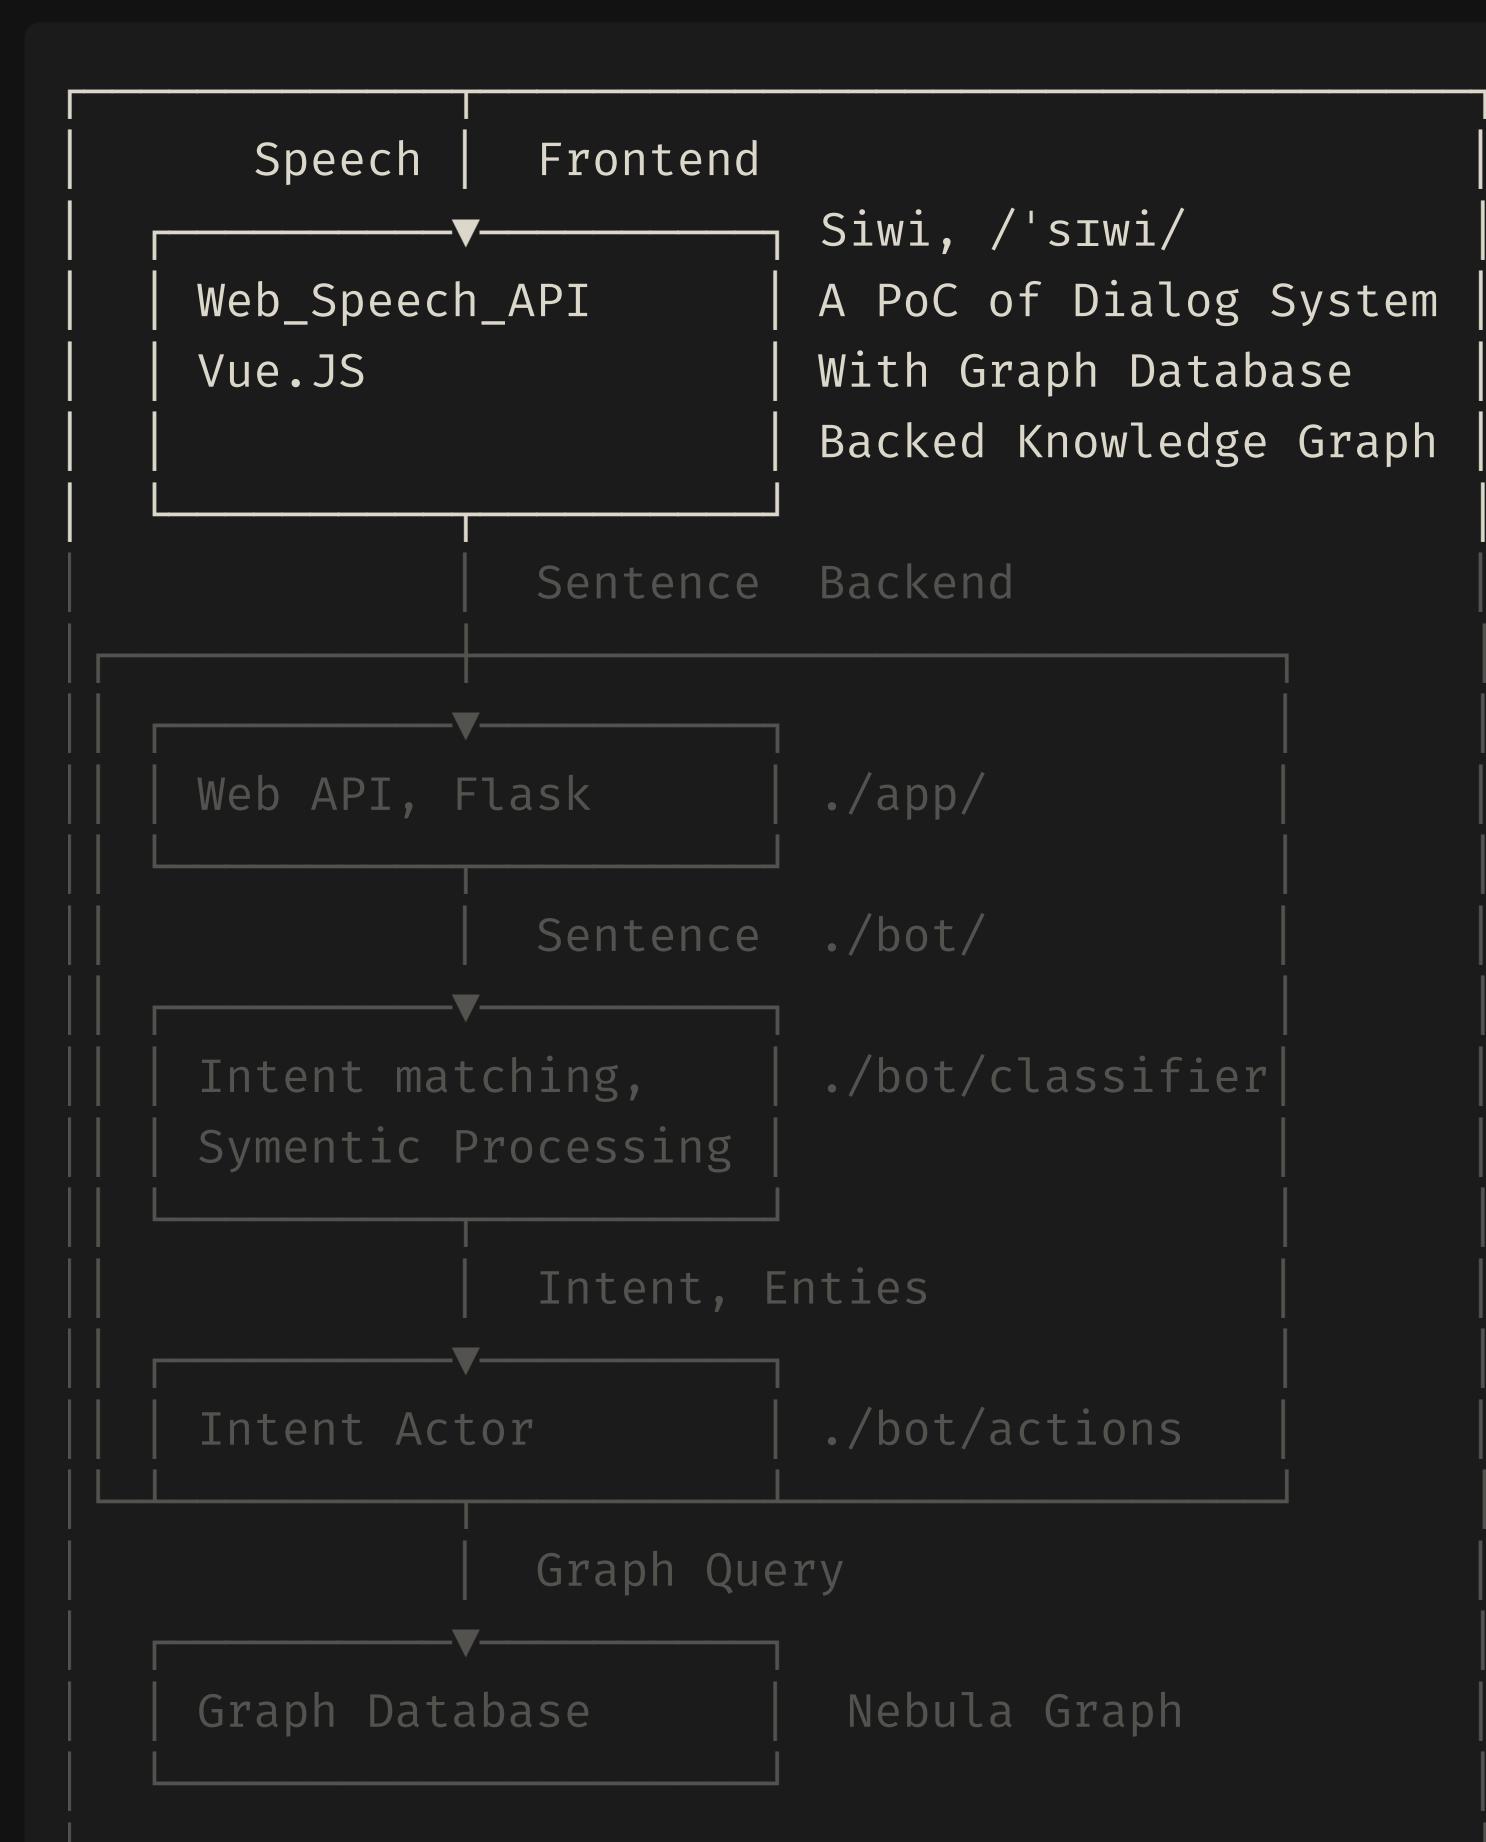
Code

```
.
├── README.md
├── src
│   ├── siwi
│   │   ├── app
│   │   └── bot
│   │       ├── actions
│   │       ├── bot
│   │       ├── classifier
│   │       │   ├── test
│   │       │   └── test_data
│   │       └── wsgi.py
│   └── siwi_frontend
│       ├── README.md
│       ├── package.json
│       └── src
│           ├── App.vue
│           └── main.js
└── wsgi.py
```



[wey-gu/nebula-siwi](https://github.com/wey-gu/nebula-siwi)

Arch



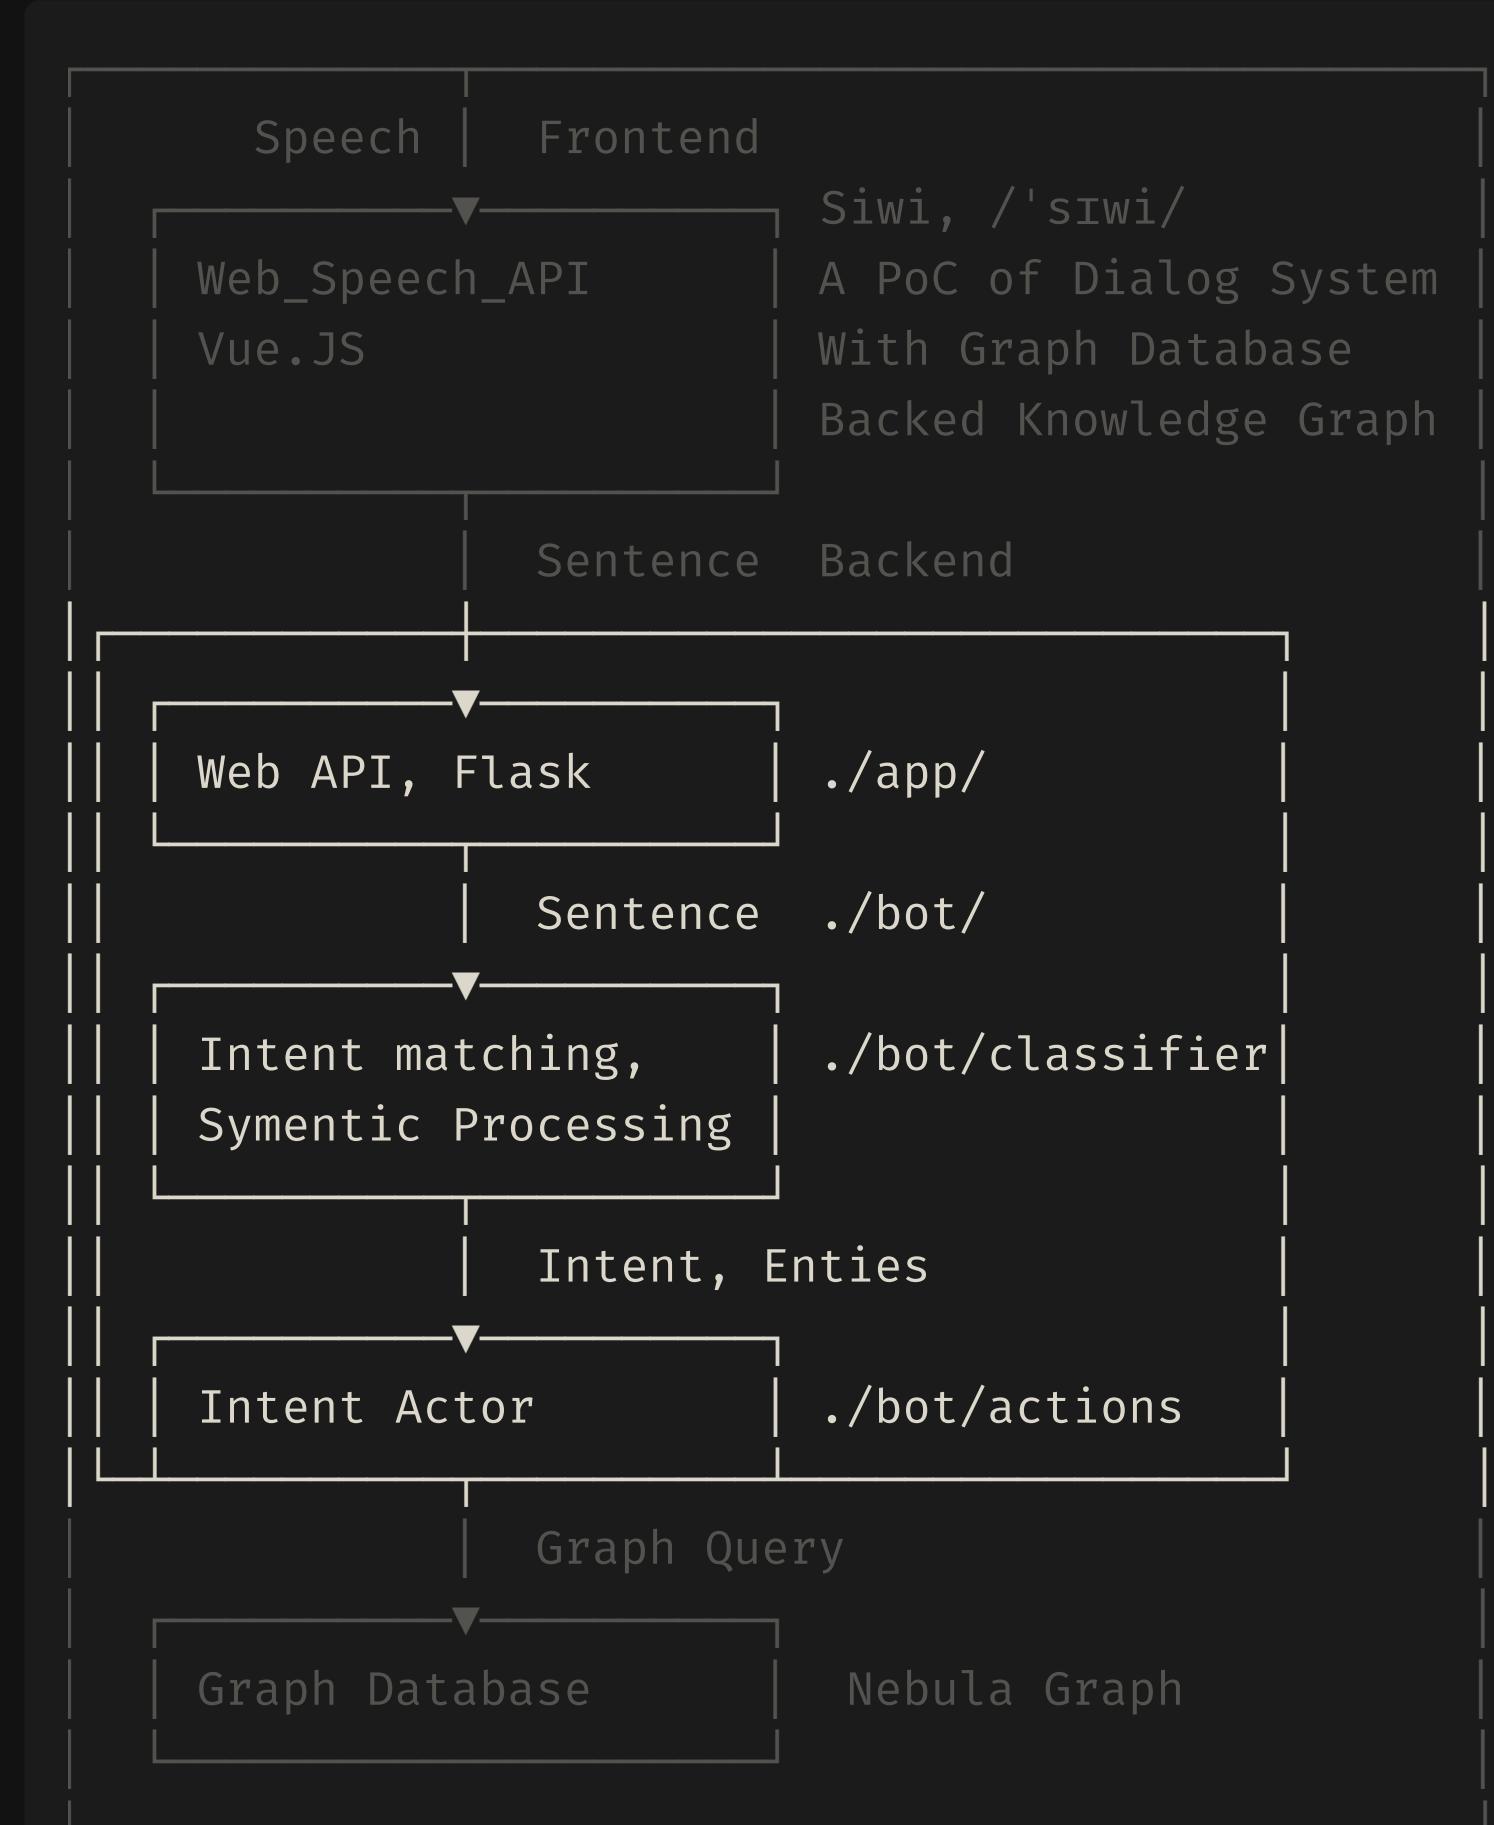
Code

```
.├── README.md  
├── src  
│   └── siwi  
│       ├── app  
│       └── bot  
│           ├── actions  
│           ├── bot  
│           ├── classifier  
│           └── test  
└── siwi_frontend  
    ├── README.md  
    ├── package.json  
    └── src  
        ├── App.vue  
        └── main.js  
wsgi.py
```



[wey-gu/nebula-siwi](https://github.com/wey-gu/nebula-siwi)

Arch



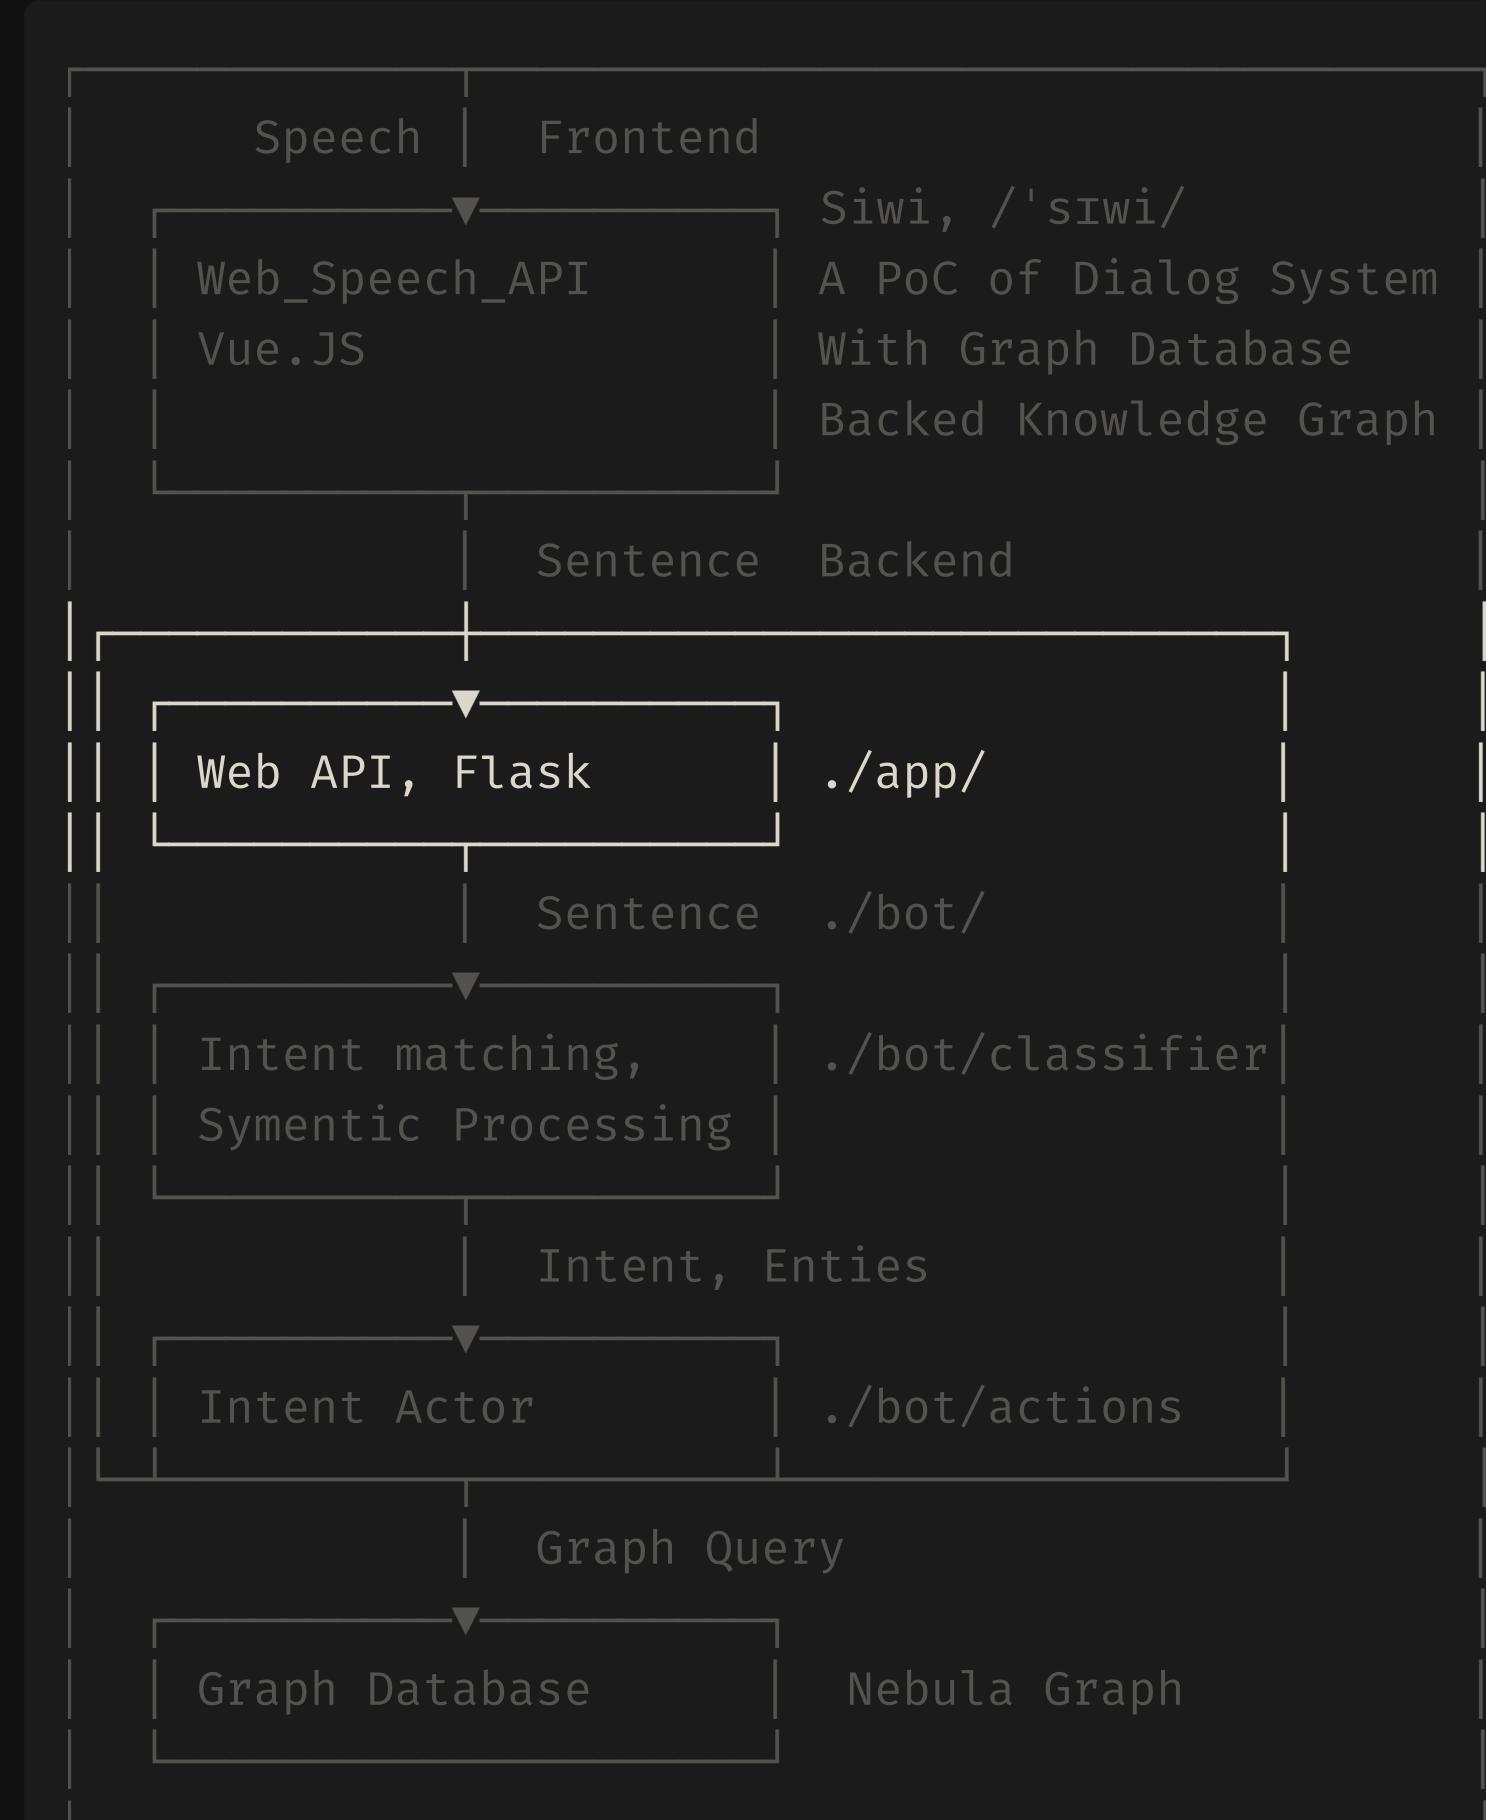
Code

```
.
├── README.md
└── src
    ├── siwi
        # Siwi-API Backend
        ├── app
        # Web Server, take HTTP req > call
        └── bot
            # Bot API
            ├── actions
            # Take Intent, Slots, Query KG here
            ├── bot
            # Entrypoint of the Bot API
            ├── classifier
            # Symentic Parse, Intent Match, etc
            └── test
            # Example Data as equivalent/mock
    └── siwi_frontend
        # Browser End
        ├── README.md
        ├── package.json
        └── src
            ├── App.vue
            # Listen to user and pass Qs to
            └── main.js
wsgi.py
```



[wey-gu/nebula-siwi](https://github.com/wey-gu/nebula-siwi)

Arch



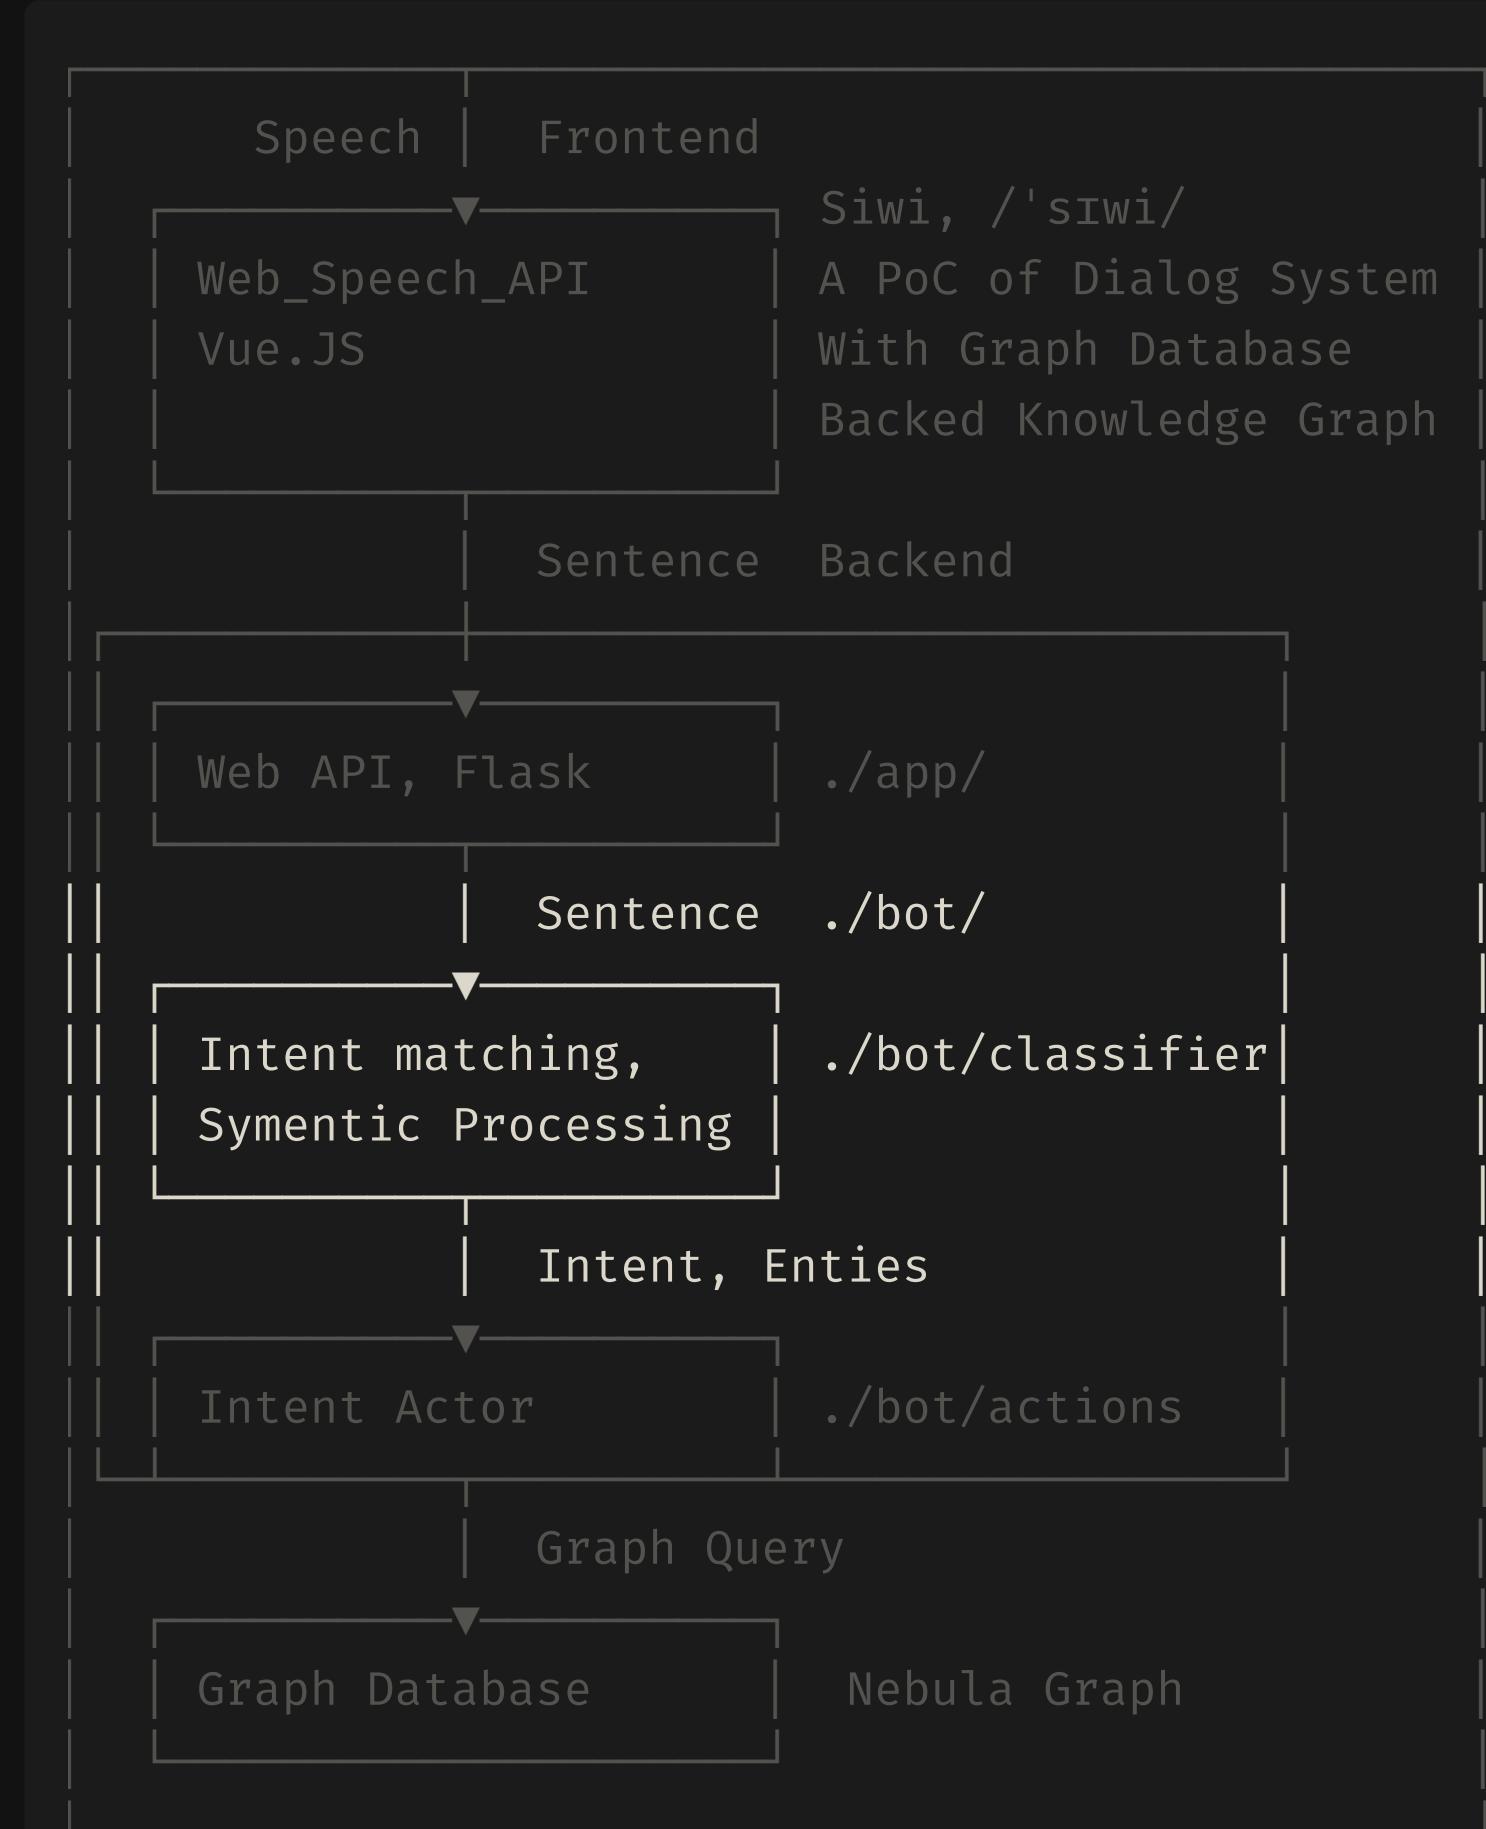
Code

```
.
├── README.md
├── src
│   └── siwi
│       ├── app
│       │   # Siwi-API Backend
│       │   # Web Server, take HTTP req > call
│       │   # Bot API
│       ├── bot
│       │   ├── actions
│       │   │   # Take Intent, Slots, Query KG here
│       │   │   # Example Data as equivalent/mock
│       │   ├── bot
│       │   │   # Entrypoint of the Bot API
│       │   ├── classifier
│       │   │   # Symentic Parse, Intent Match, etc
│       │   │   # Example Data as equivalent/mock
│       │   └── test
│           # Browser End
└── siwi_frontend
    ├── README.md
    ├── package.json
    └── src
        ├── App.vue
        │   # Listen to user and pass Qs to bot
        └── main.js
wsgi.py
```



[wey-gu/nebula-siwi](https://github.com/wey-gu/nebula-siwi)

Arch



Code

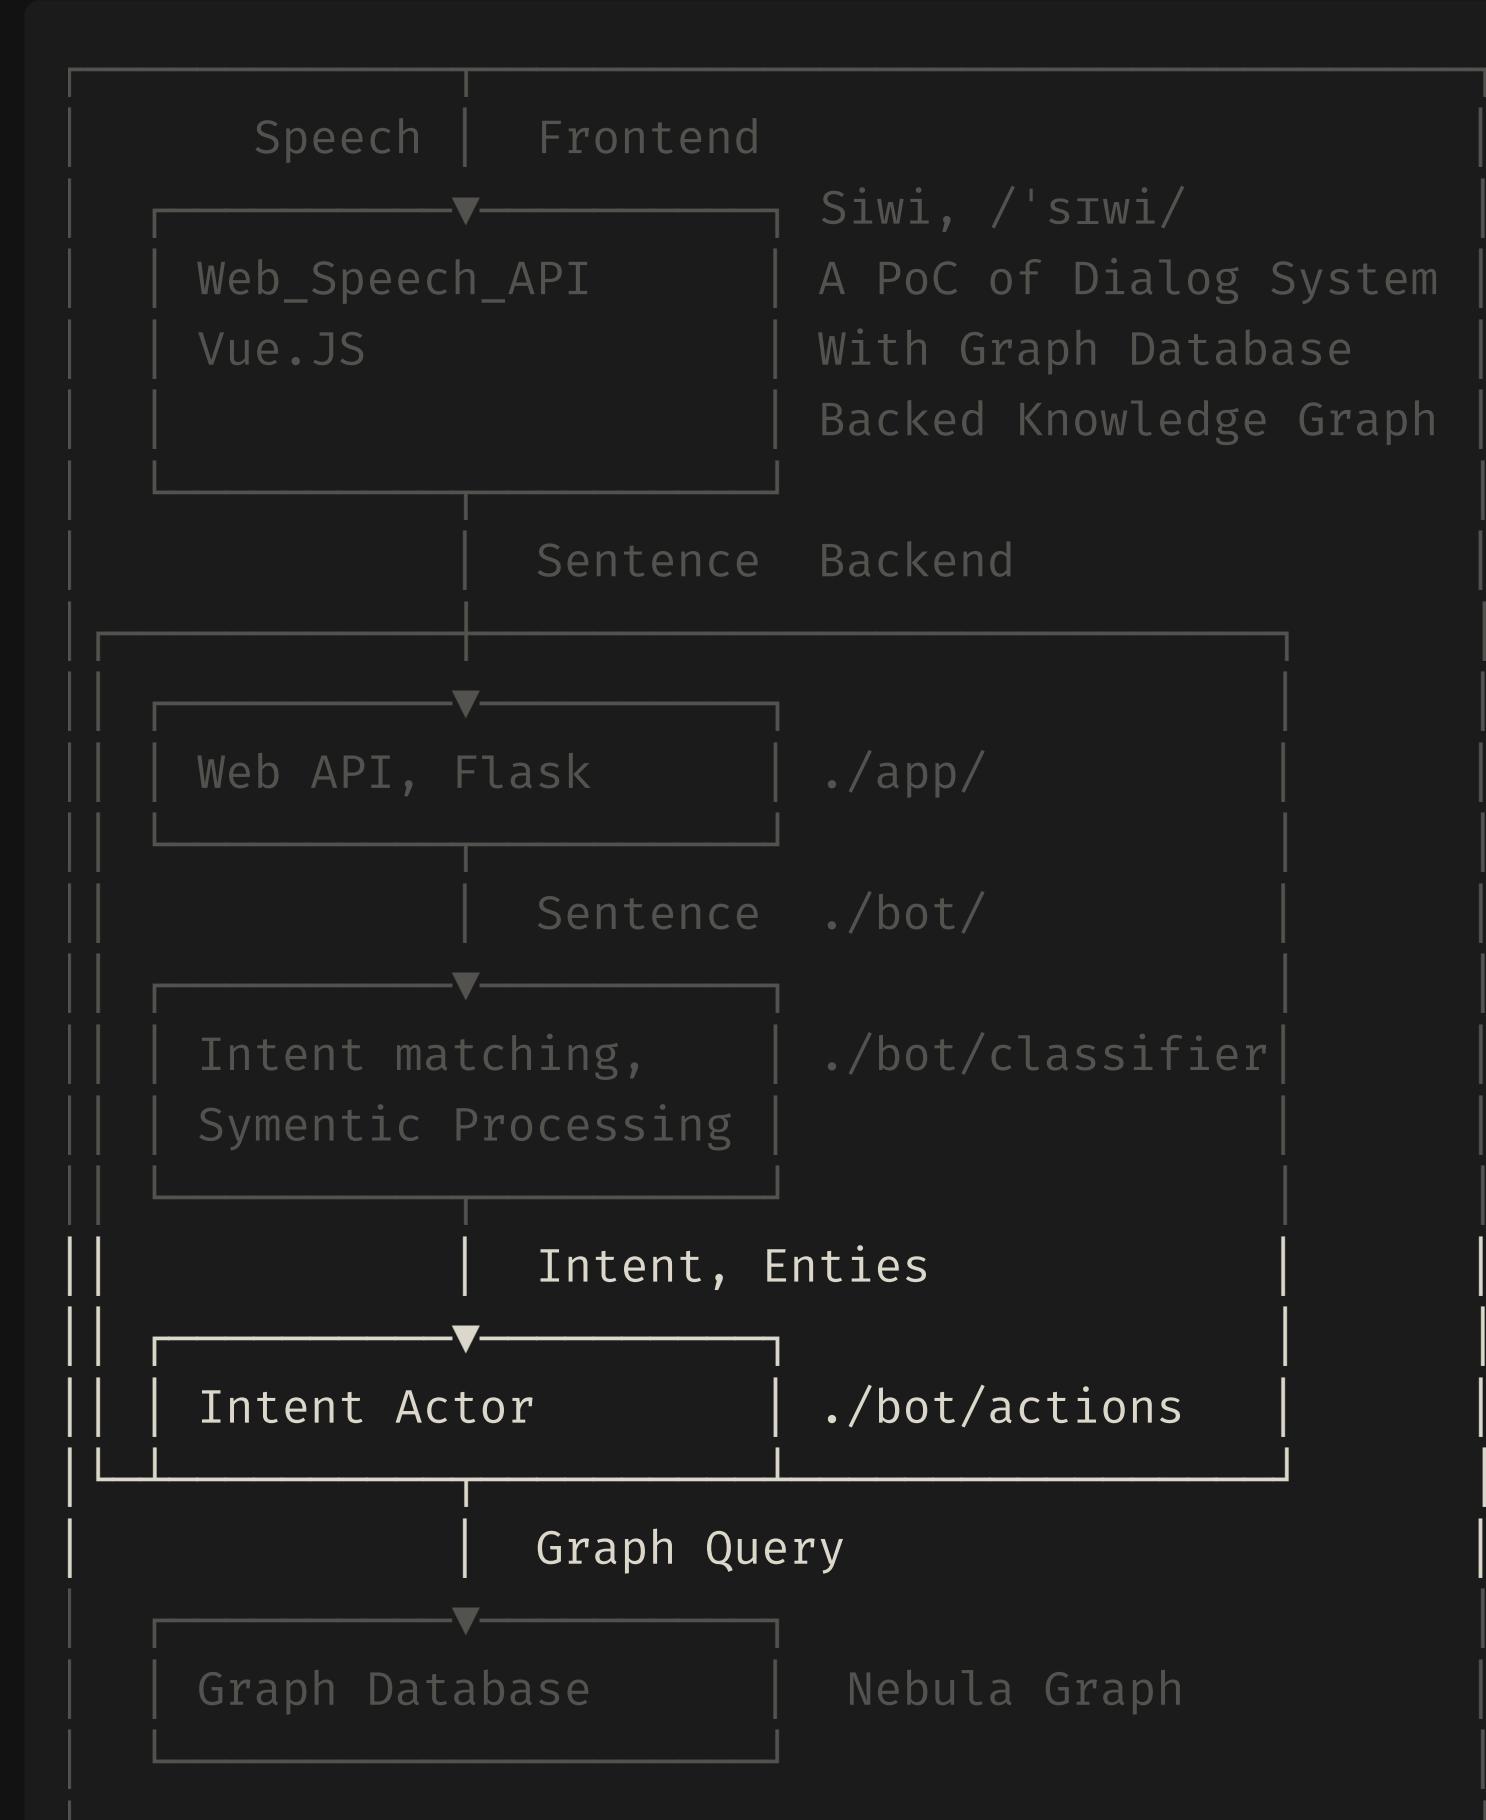
```
.
├── README.md
├── src
│   ├── siwi
│   │   ├── app
│   │   └── bot
│   │       ├── actions
│   │       ├── bot
│   │       ├── classifier
│   │       └── test
│   └── siwi_frontend
│       ├── README.md
│       ├── package.json
│       └── src
│           ├── App.vue
│           └── main.js
└── wsgi.py
```

The code structure corresponds to the architecture diagram. It includes a top-level README file and a 'src' directory. The 'src' directory contains a 'siwi' folder which holds an 'app' (the web server), a 'bot' (the bot API) which includes 'actions' (for handling intents and slots), and a 'classifier' (for semantic parsing and intent matching). It also contains a 'test' folder with example data. Additionally, there is a 'siwi_frontend' folder containing its own README, package.json, and a 'src' folder with 'App.vue' and 'main.js' files. A separate 'wsgi.py' file is located at the bottom level.



[wey-gu/nebula-siwi](https://github.com/wey-gu/nebula-siwi)

Arch



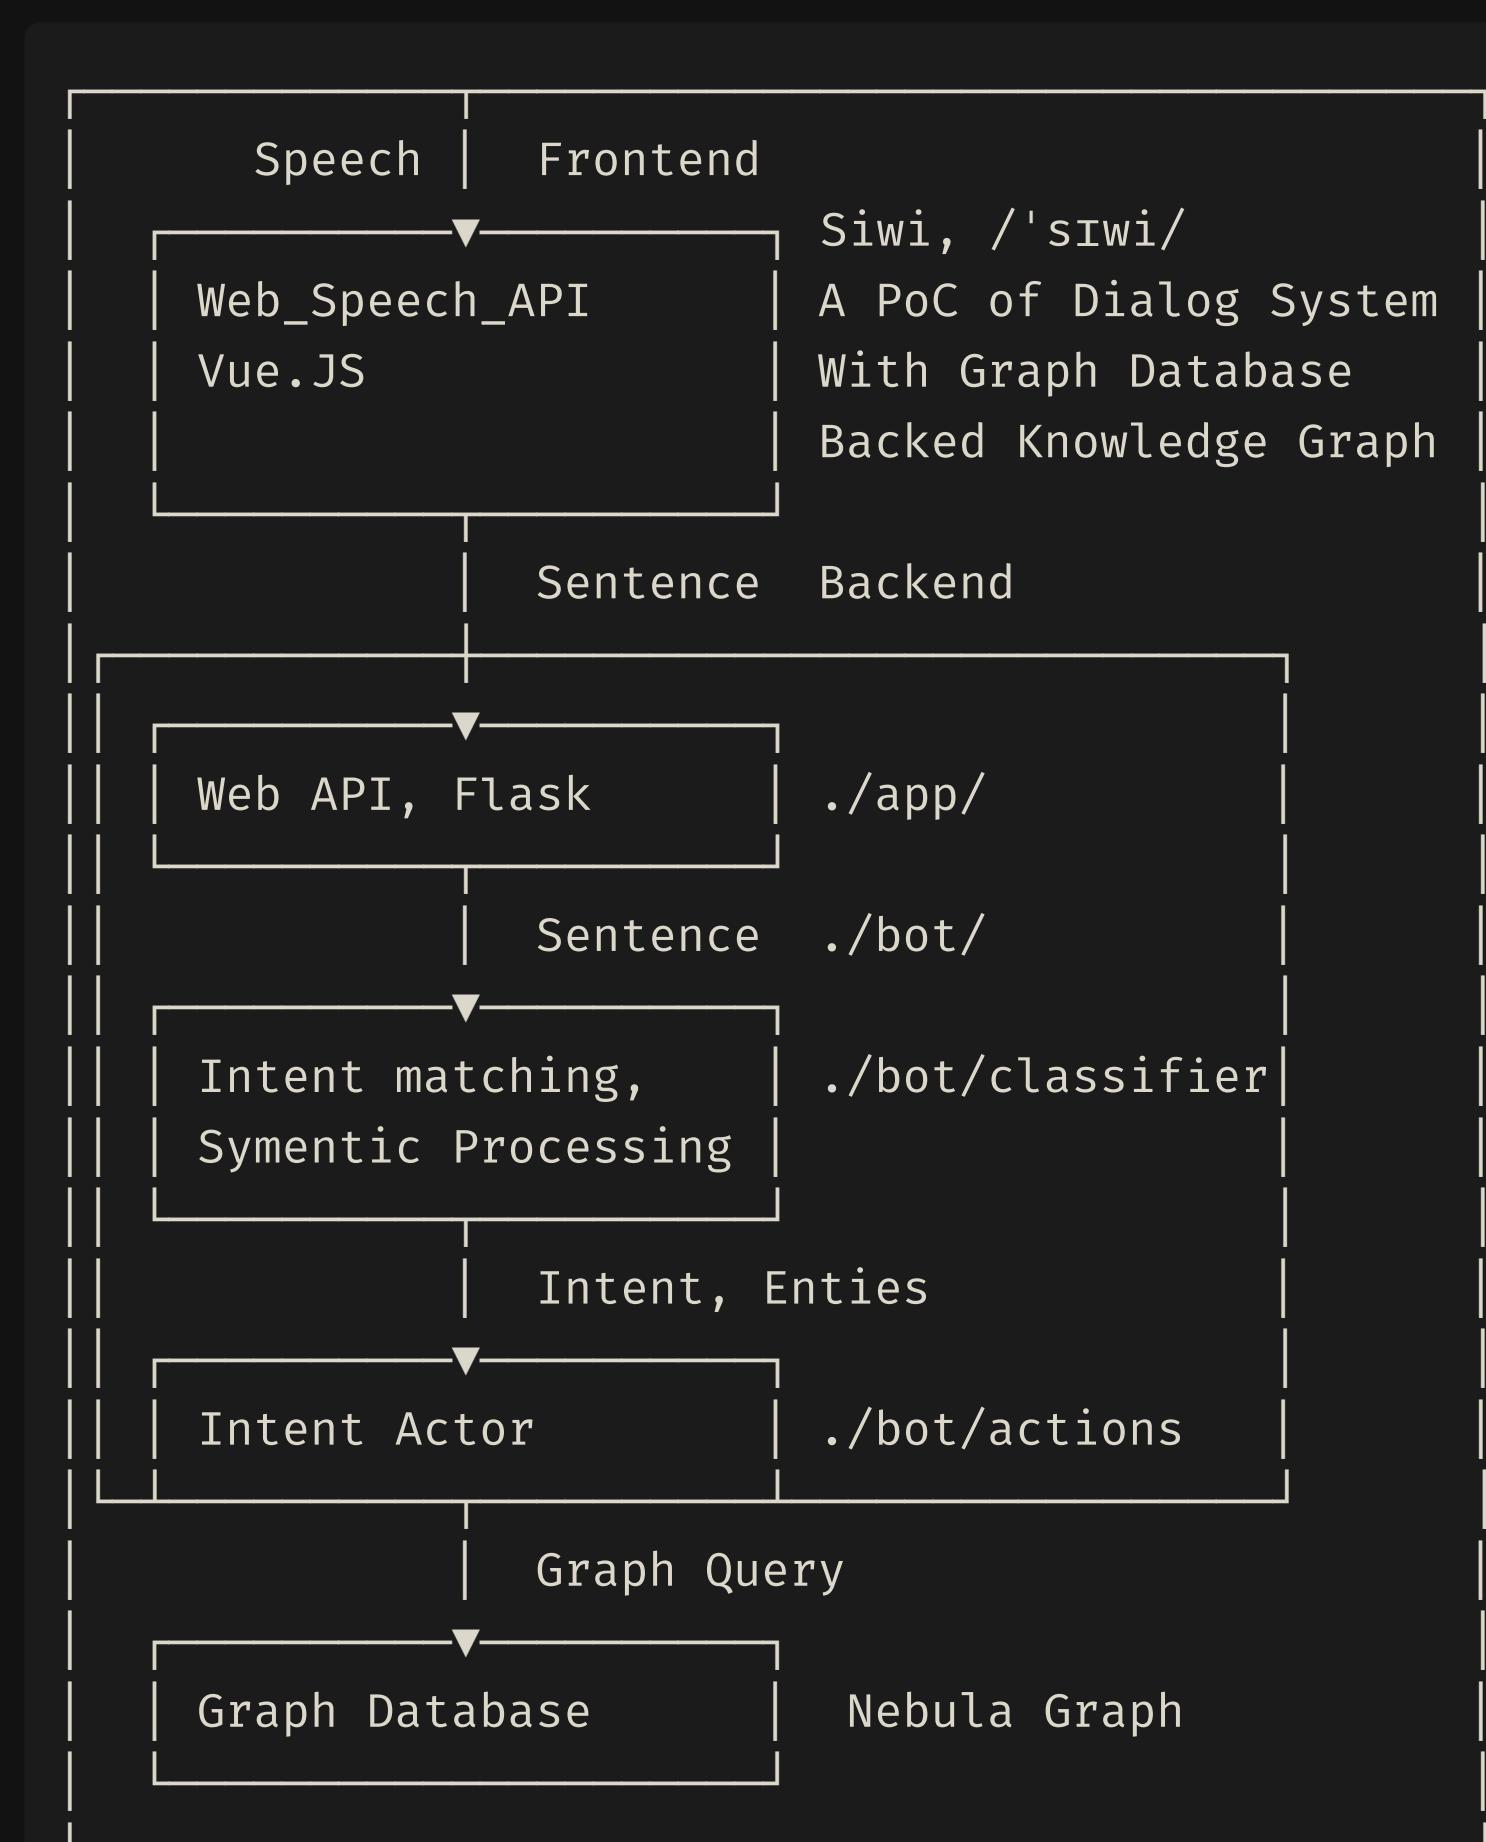
Code

```
.
├── README.md
├── src
│   └── siwi
│       ├── app
│       │   # Siwi-API Backend
│       │   # Web Server, take HTTP req > call
│       │   # Bot API
│       ├── bot
│       │   ├── actions
│       │   │   # Take Intent, Slots, Query KG here
│       │   │   # Example Data as equivalent/mock
│       │   ├── bot
│       │   │   # Entrypoint of the Bot API
│       │   ├── classifier
│       │   │   # Symentic Parse, Intent Match, etc
│       │   │   # Browser End
│       │   ├── test
│       │   │   # Example Data as equivalent/mock
│       └── siwi_frontend
│           ├── README.md
│           ├── package.json
│           └── src
│               ├── App.vue
│               │   # Listen to user and pass Qs to
│               └── main.js
└── wsgi.py
```



[wey-gu/nebula-siwi](https://github.com/wey-gu/nebula-siwi)

Arch



Code

```
.├── README.md  
├── src  
│   └── siwi  
│       ├── app  
│       └── bot  
│           ├── actions  
│           ├── bot  
│           ├── classifier  
│           └── test  
└── siwi_frontend  
    ├── README.md  
    ├── package.json  
    └── src  
        ├── App.vue  
        └── main.js  
wsgi.py
```



[wey-gu/nebula-siwi](https://github.com/wey-gu/nebula-siwi)

一个接地气的小 Demo

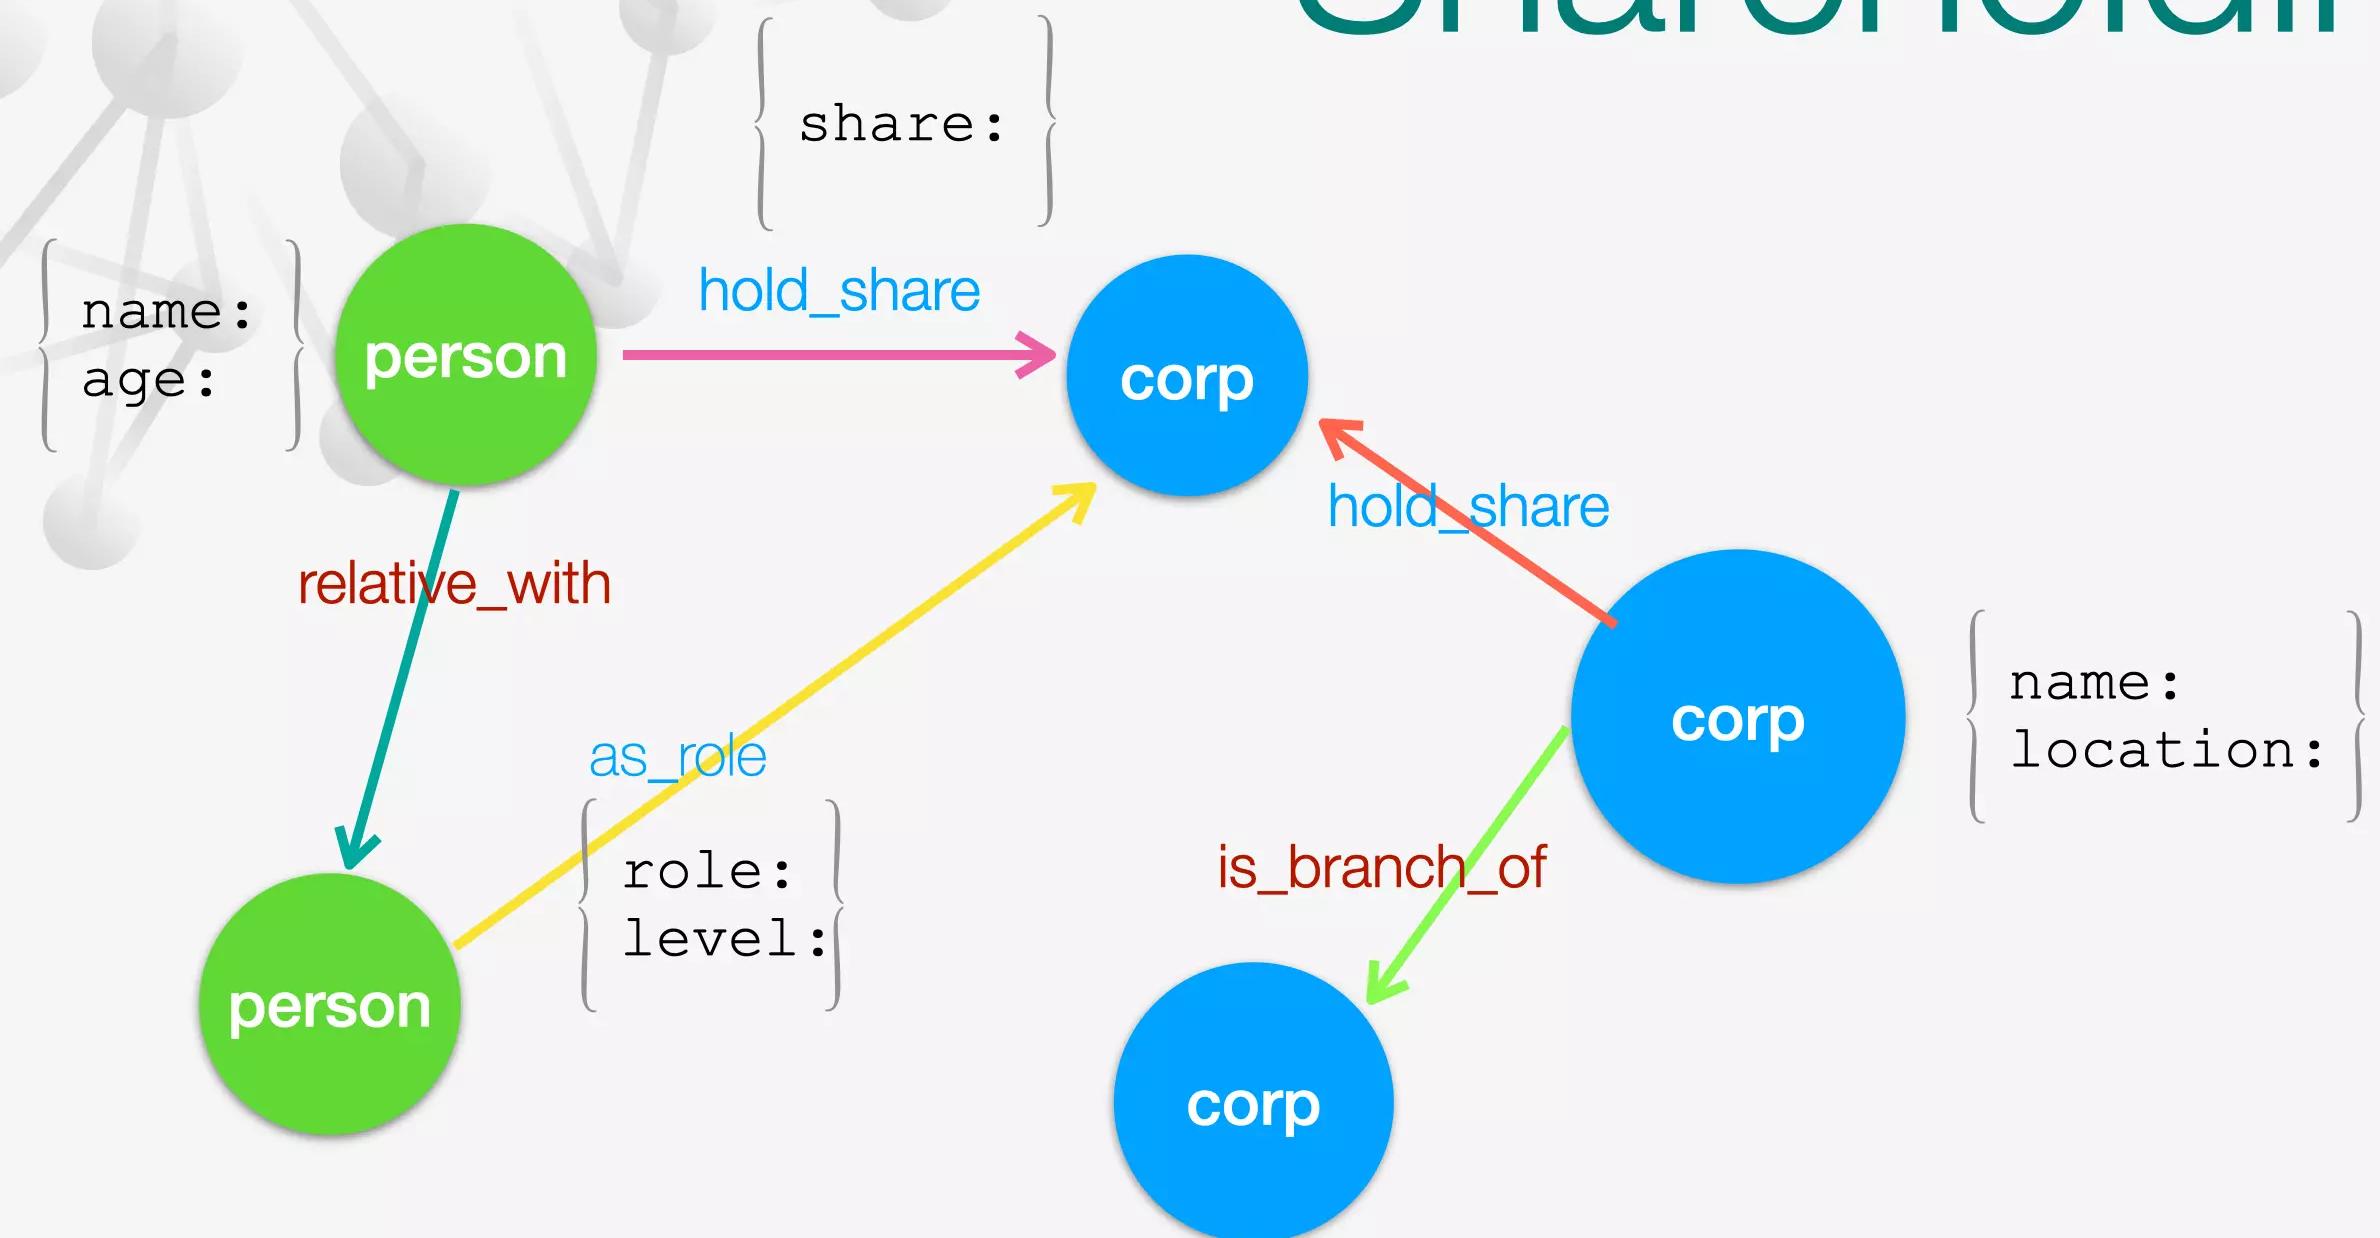
Siwi 智能语音助手

Corp-Rel-Search

一个企业股权关系分析系统

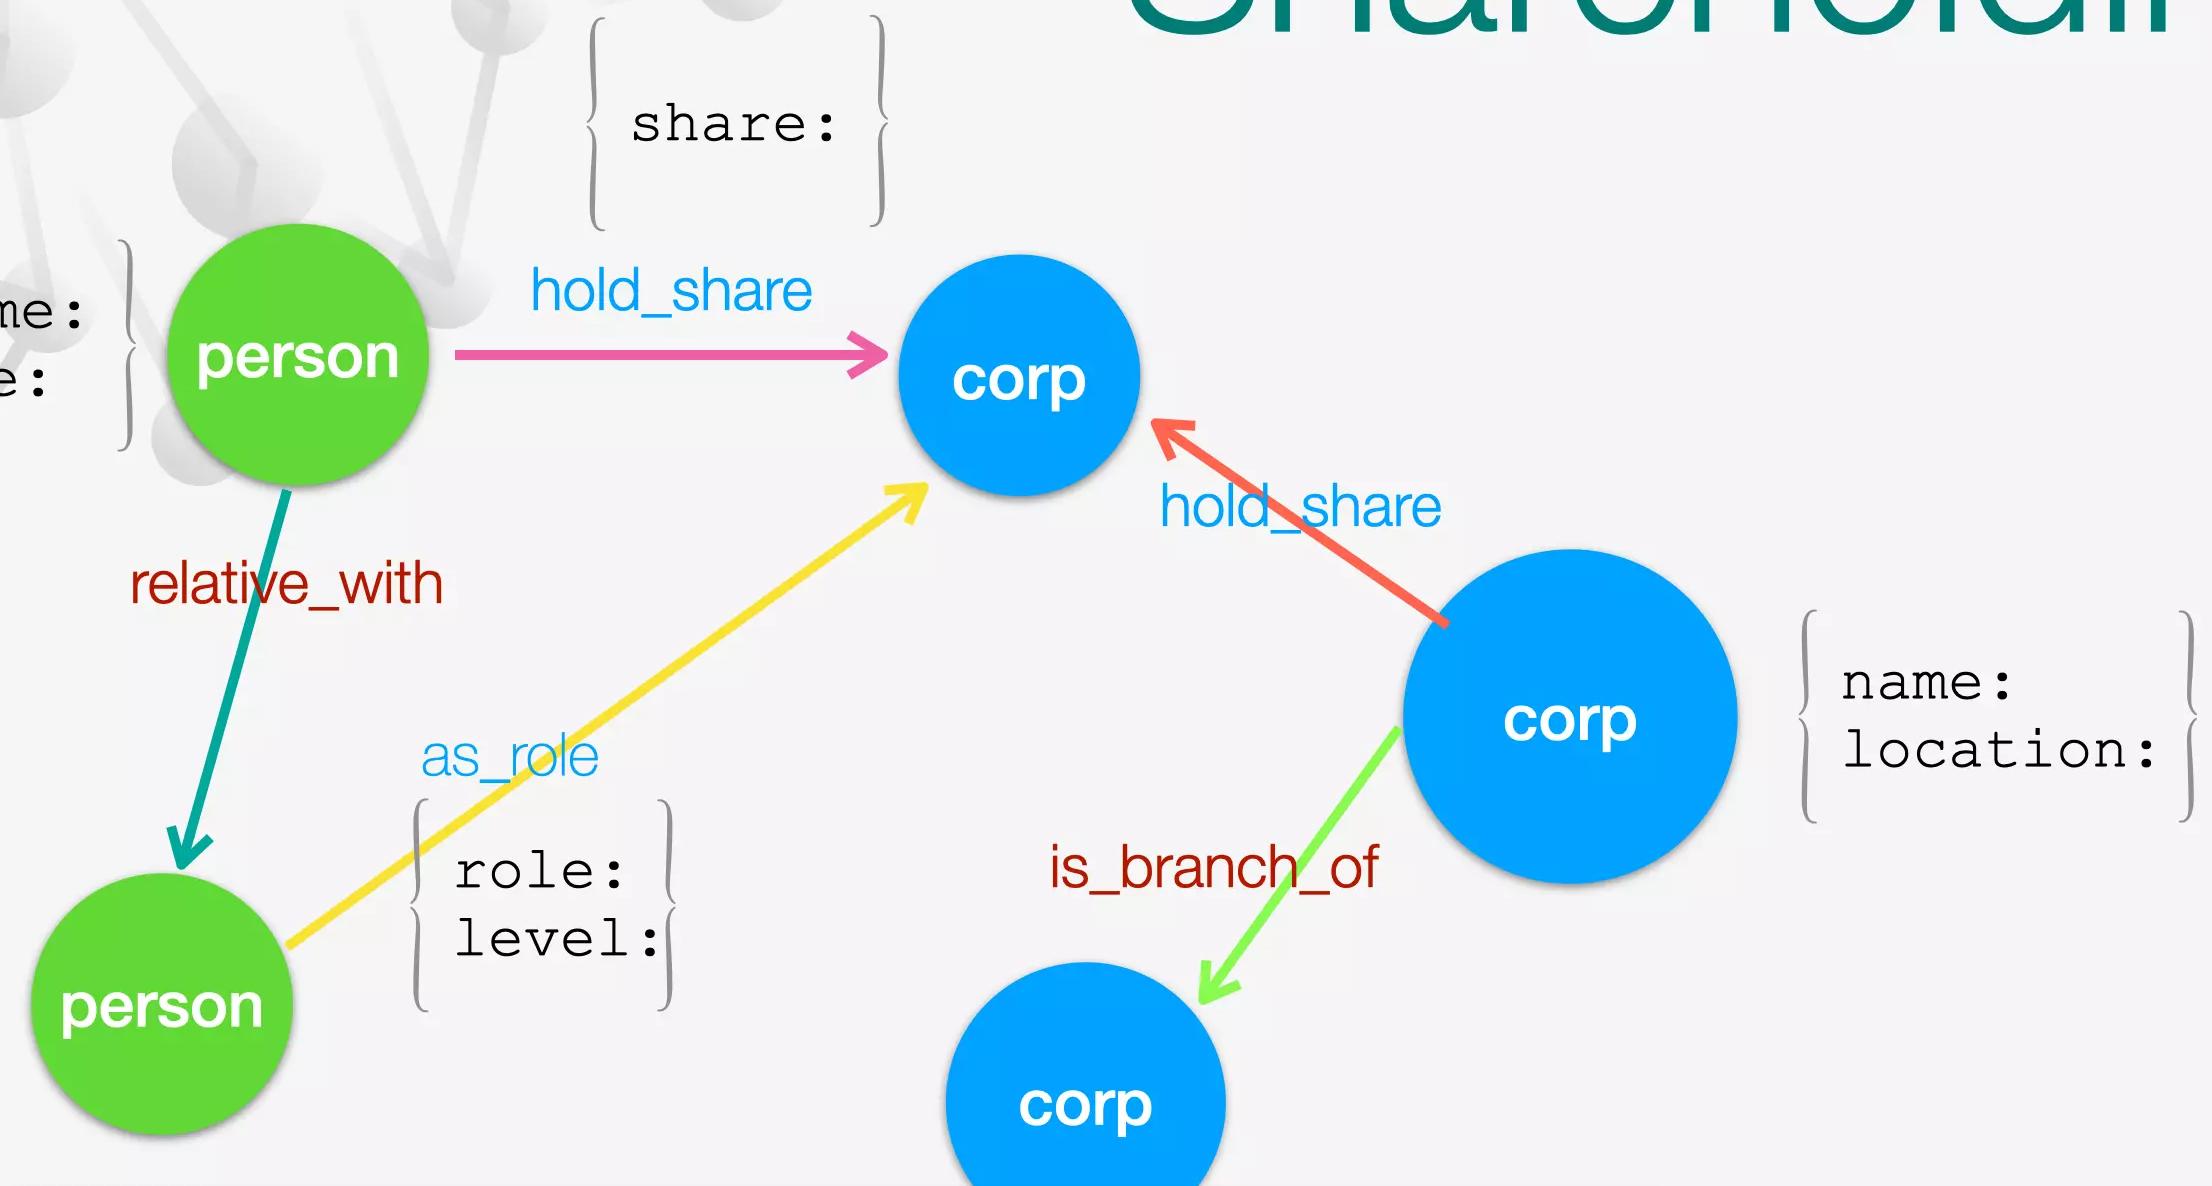
<https://nebula-graph.io>

Shareholding



[wey-gu/nebula-shareholding-example](https://github.com/wey-gu/nebula-shareholding-example)

Shareholding

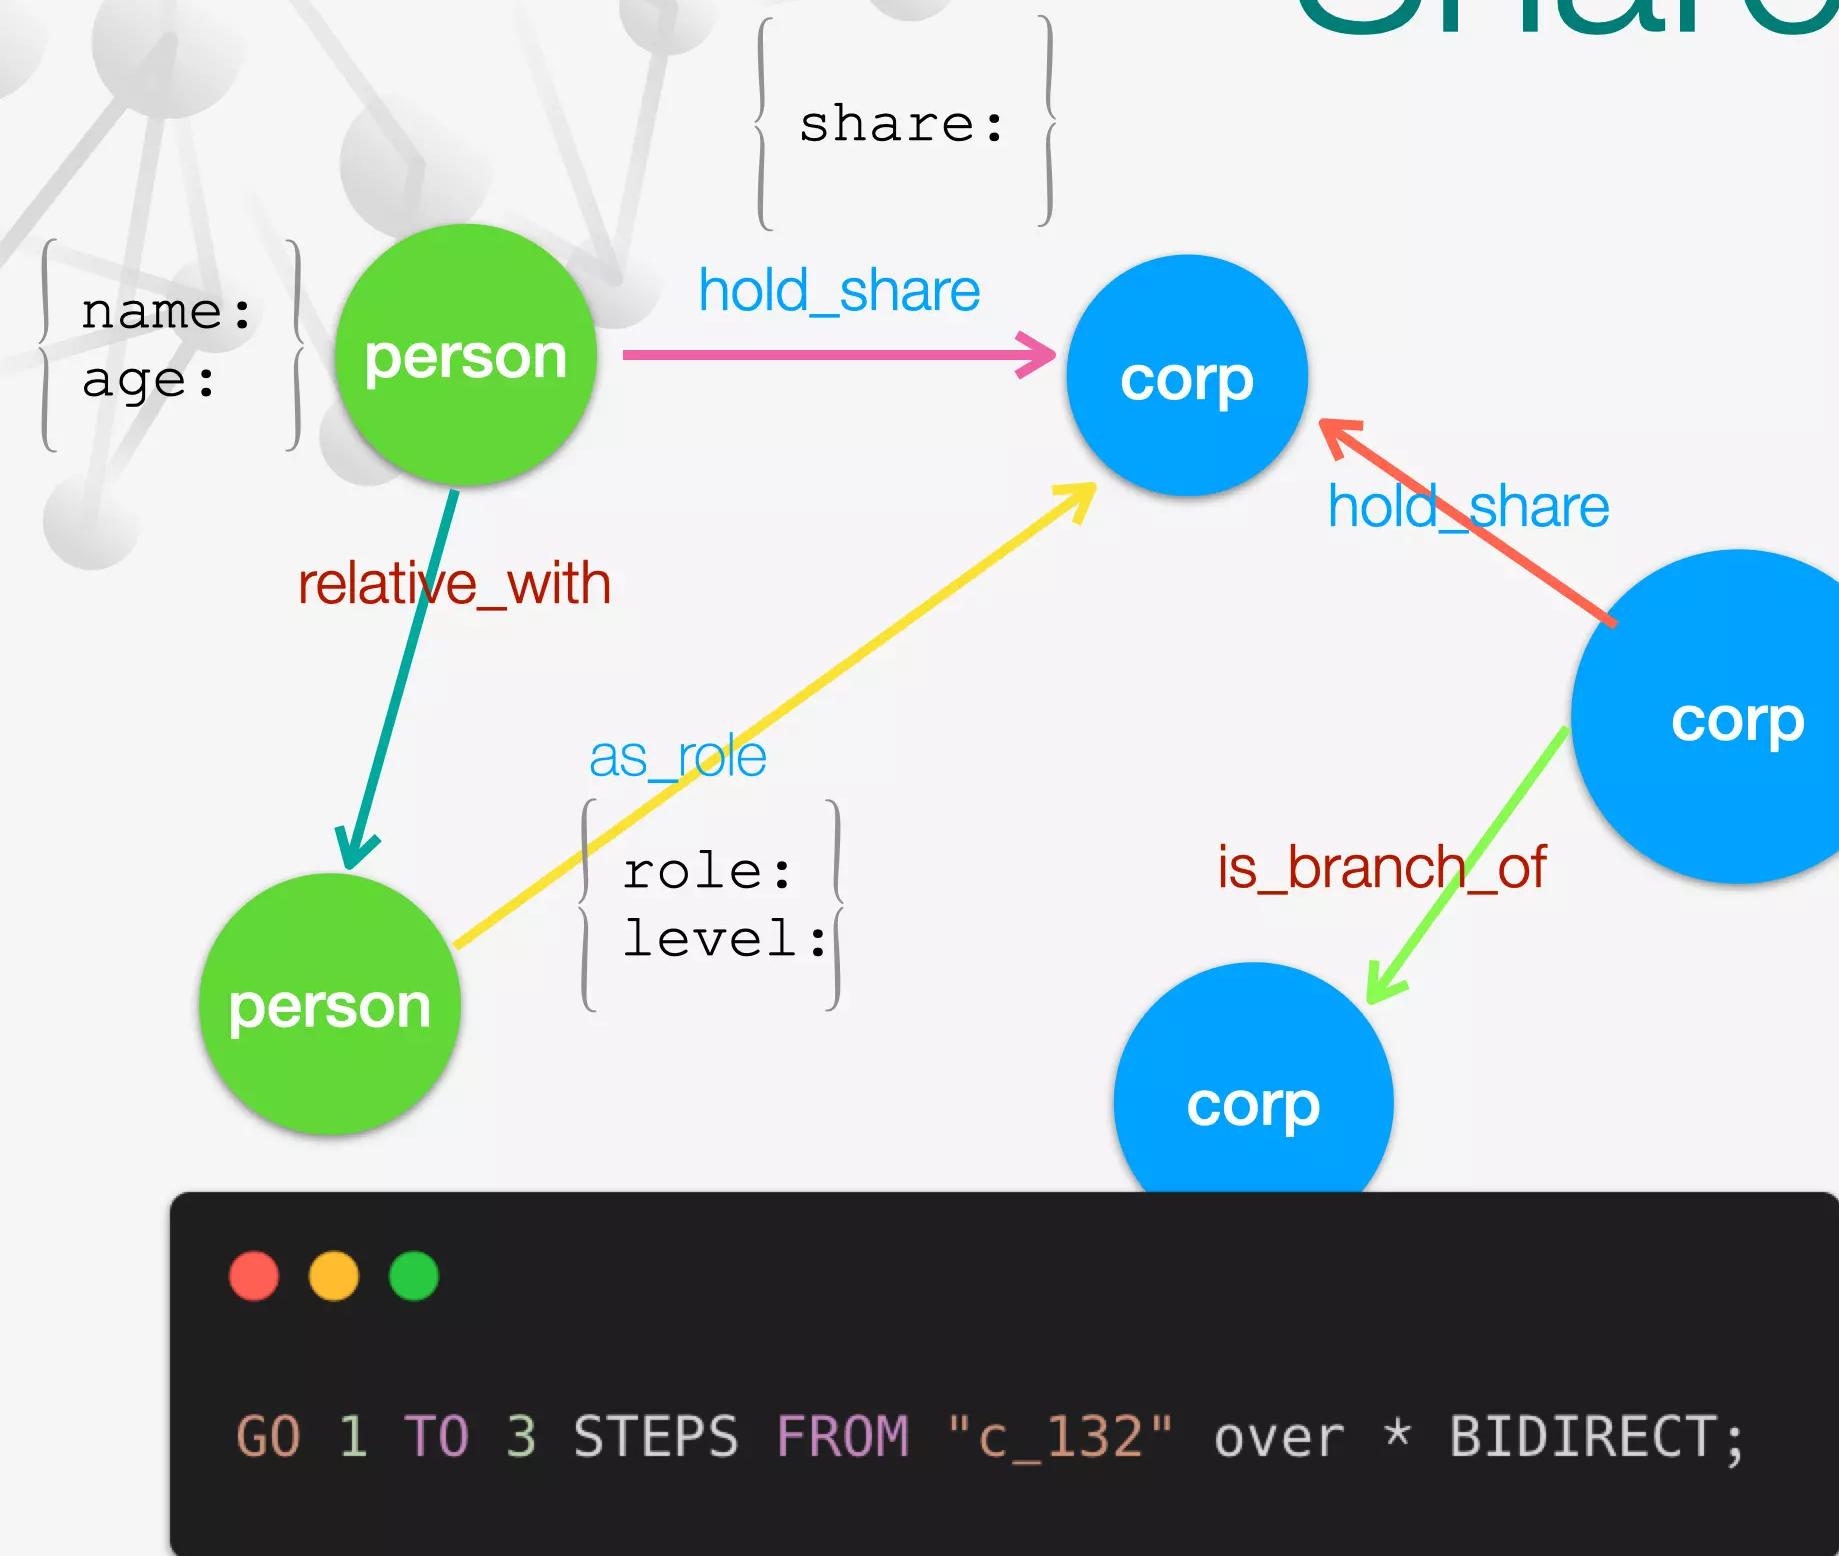


G0 1 T0 3 STEPS FROM "c_132" over * BIDIRECT;

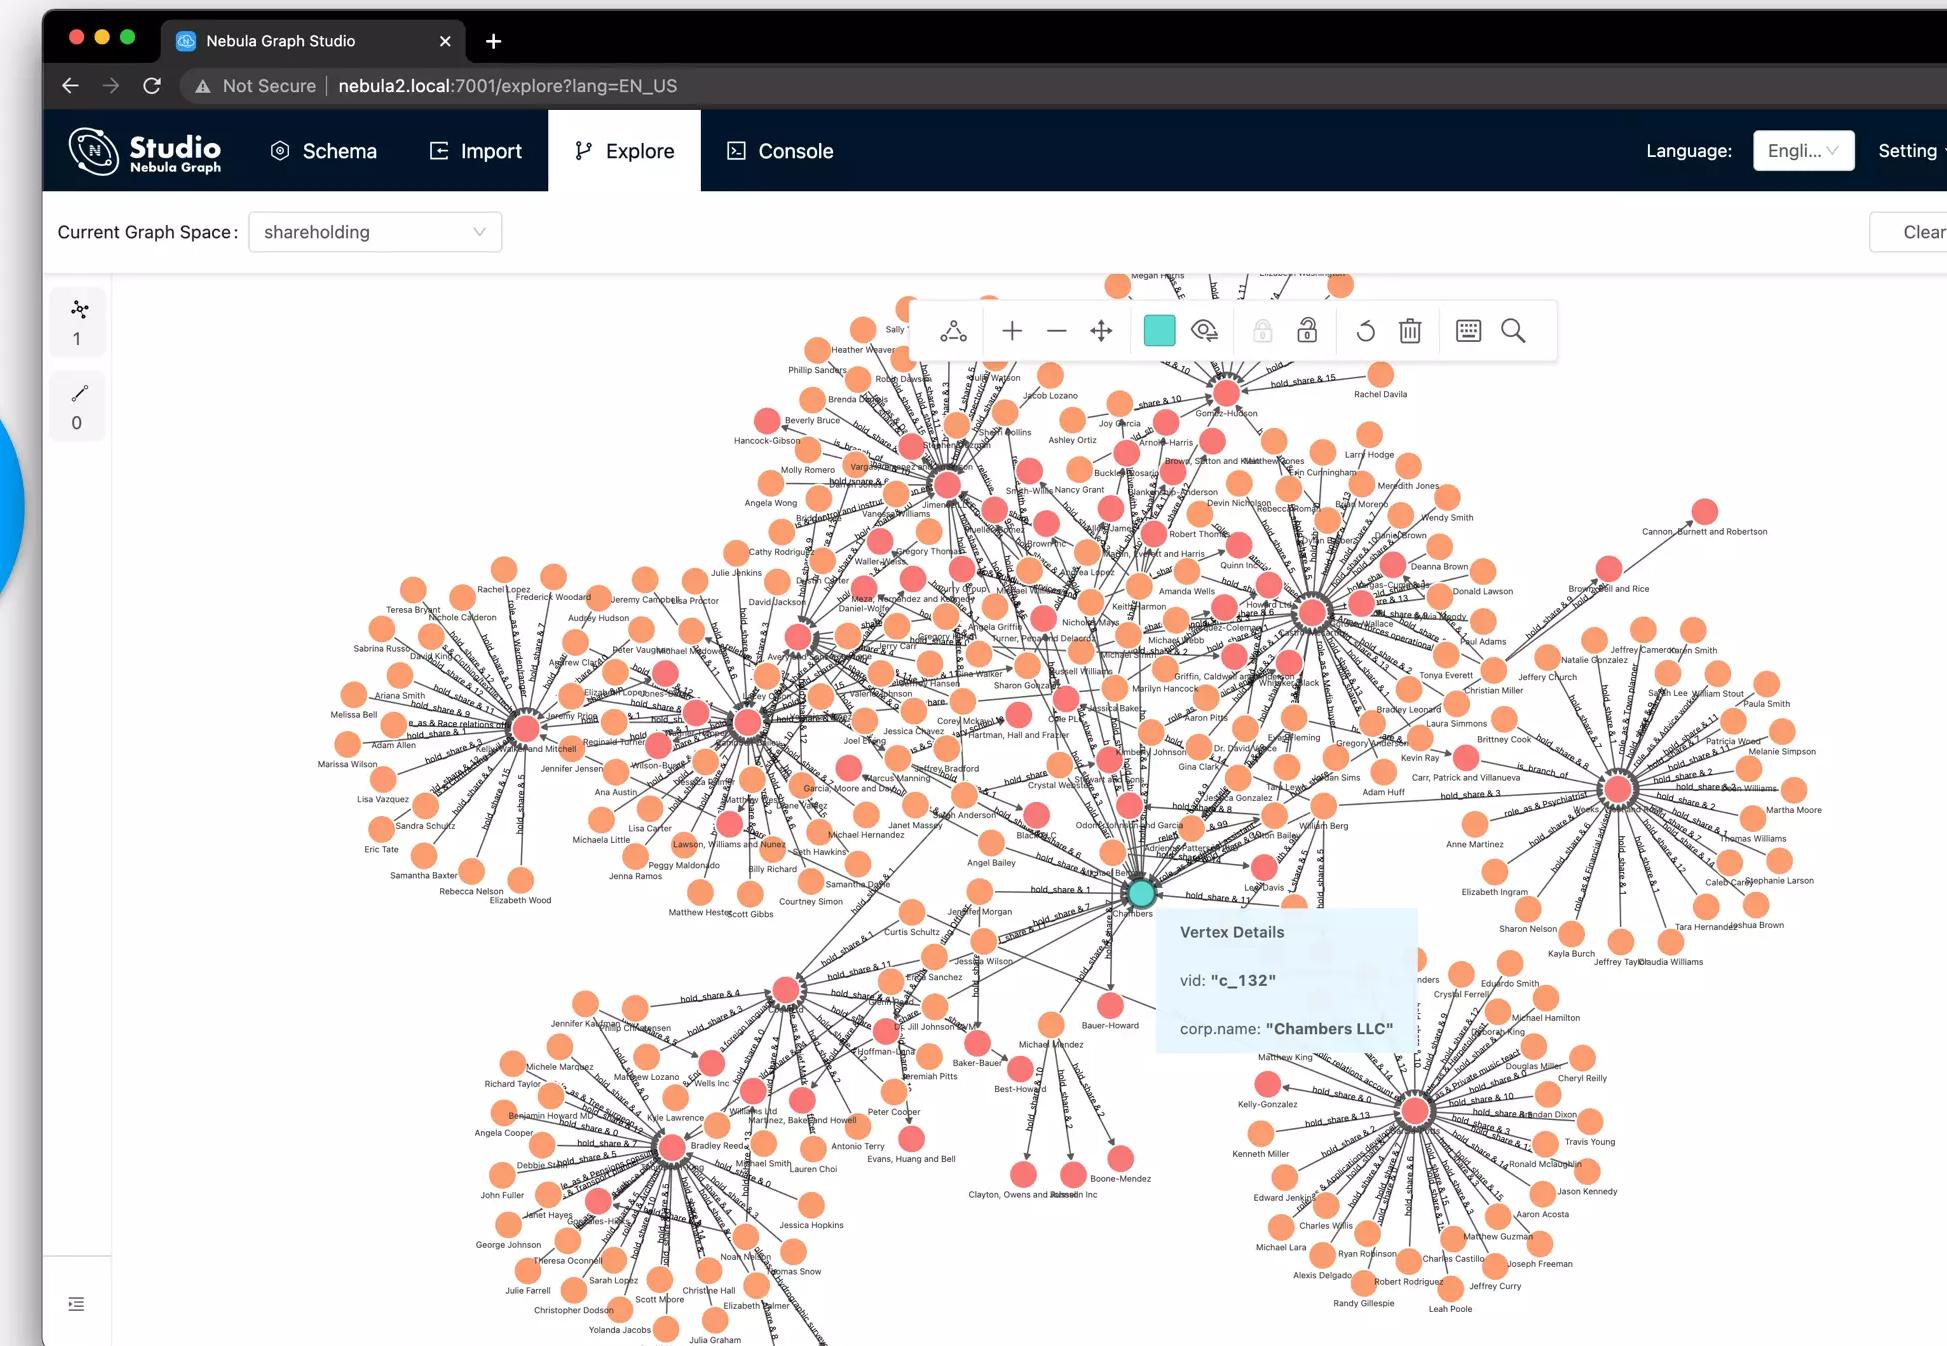


wey-gu/nebula-shareholding-example

<https://nebula-graph.io>

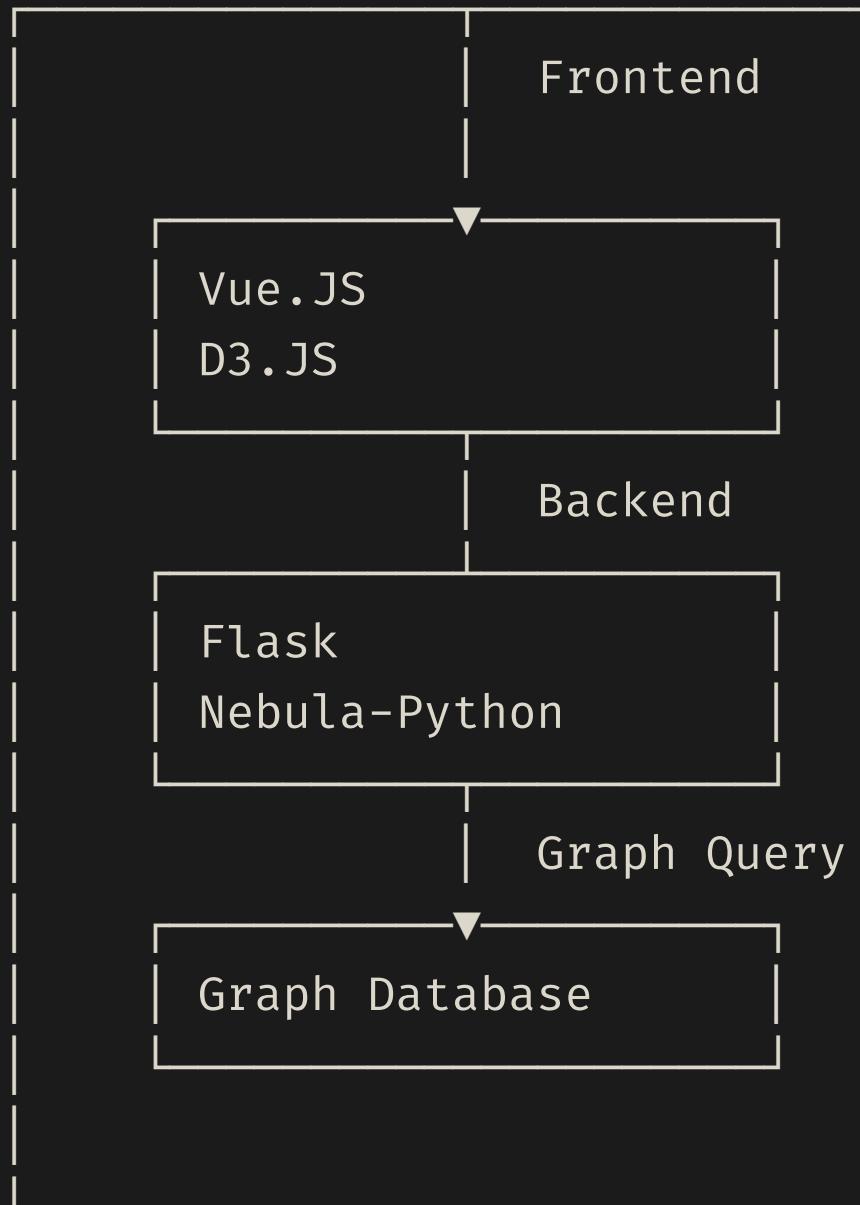


Shareholding



[wey-gu/nebula-shareholding-example](https://github.com/wey-gu/nebula-shareholding-example)

架构 Corp-Rel-Search



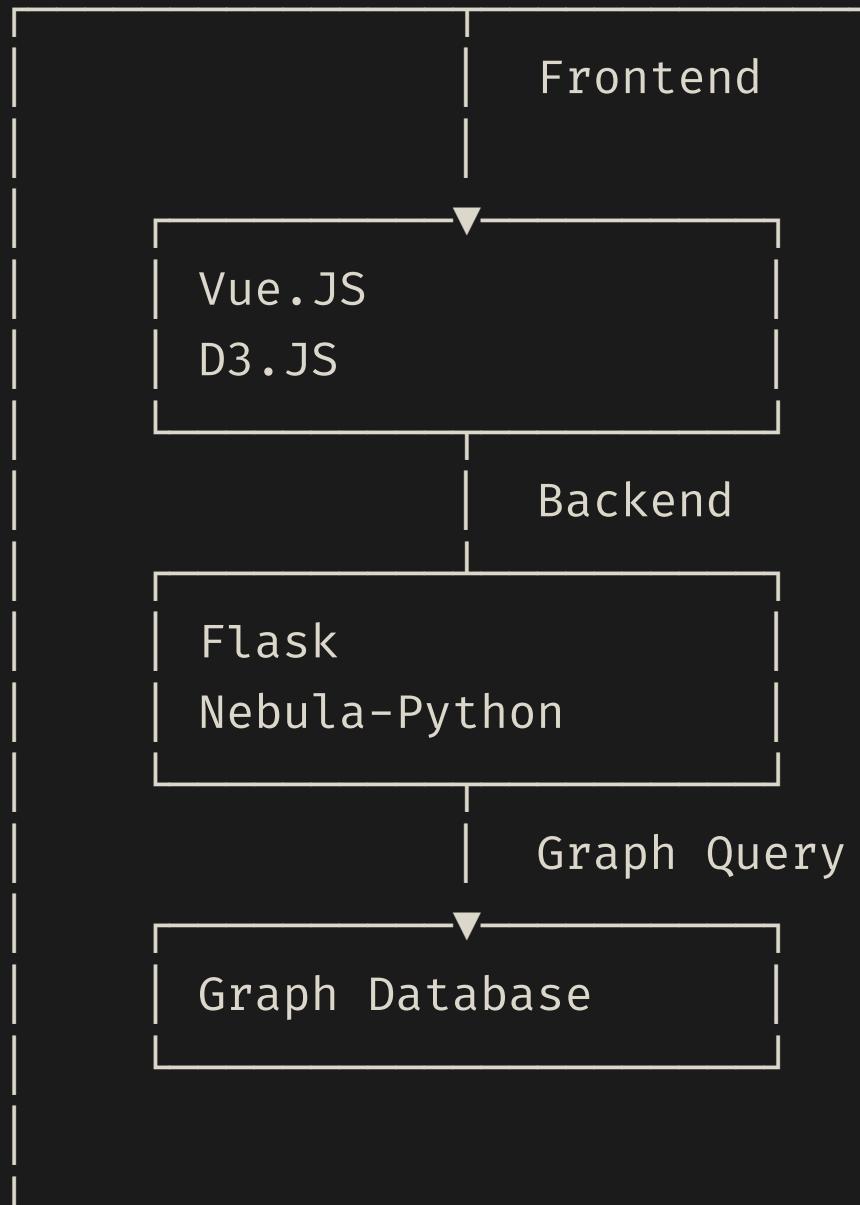
The Code

```
├── README.md      # You could find Design Logs here
├── corp-rel-backend
│   └── app.py      # Flask App to handle Request and call
├── corp-rel-frontend
│   └── src
│       ├── App.vue
│       └── main.js    # Vue App to call Flask App and Render
└── requirements.txt
```

```
@app.route("/api", methods=["POST"])
def api():
    entity = request.get_json().get("entity", "")
    if entity:
        resp = query_shareholding(entity)
        data = make_graph_response(resp)
        #
    return jsonify(data)

def query_shareholding(entity):
    query_string = (
        f"MATCH p=(v)-[e:hold_share|:is_branch_of|:reletive"
        f"WHERE id(v) IN [{entity}] RETURN p LIMIT 100"
    )
    session = connection_pool.get_session('root', 'nebula')
    resp = session.execute(query_string)
    return resp
```

架构 Corp-Rel-Search



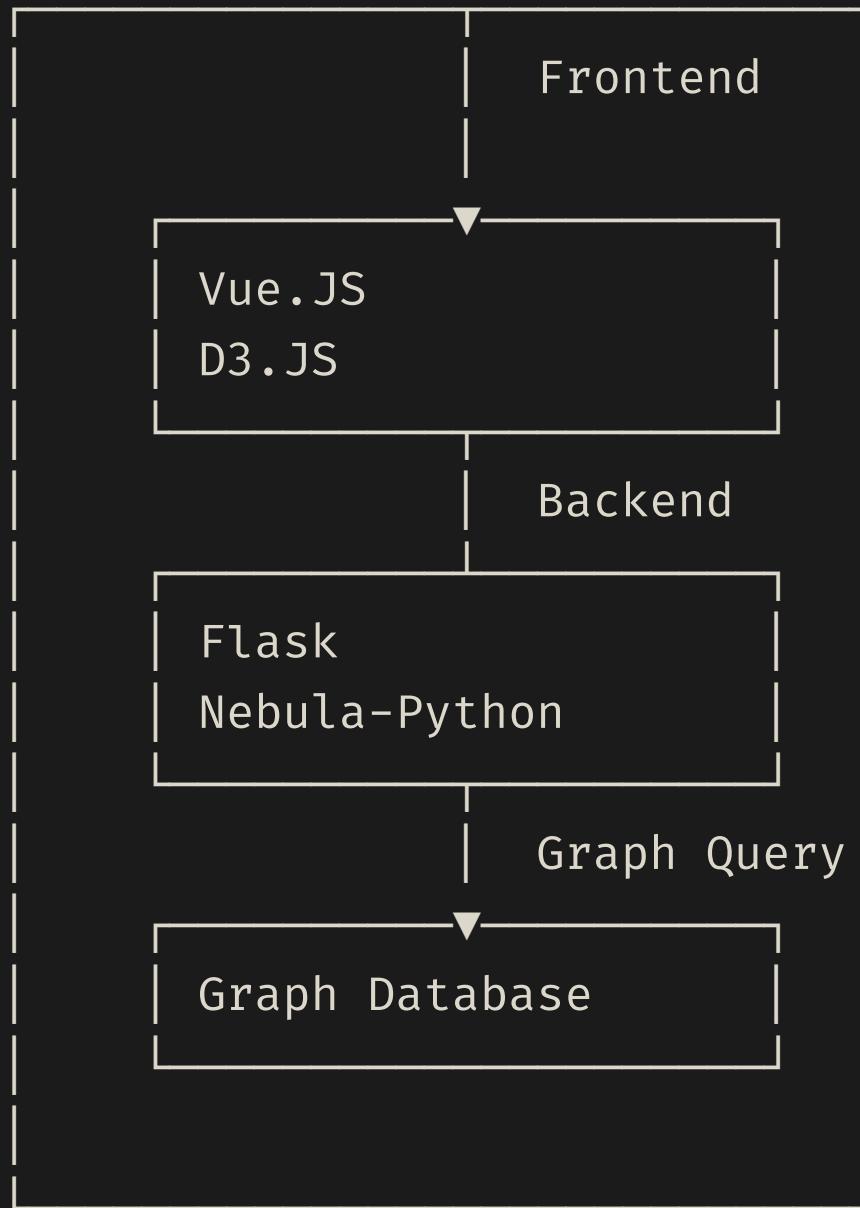
The Code

```
├── README.md      # You could find Design Logs here
├── corp-rel-backend
│   └── app.py      # Flask App to handle Request and call
├── corp-rel-frontend
│   └── src
│       ├── App.vue
│       └── main.js    # Vue App to call Flask App and Render
└── requirements.txt
```

```
@app.route("/api", methods=["POST"])
def api():
    entity = request.get_json().get("entity", "")
    if entity:
        resp = query_shareholding(entity)
        data = make_graph_response(resp)
        #
    return jsonify(data)

def query_shareholding(entity):
    query_string = (
        f"MATCH p=(v)-[e:hold_share|:is_branch_of|:reletive"
        f"WHERE id(v) IN [{entity}] RETURN p LIMIT 100"
    )
    session = connection_pool.get_session('root', 'nebula')
    resp = session.execute(query_string)
    return resp
```

架构 Corp-Rel-Search



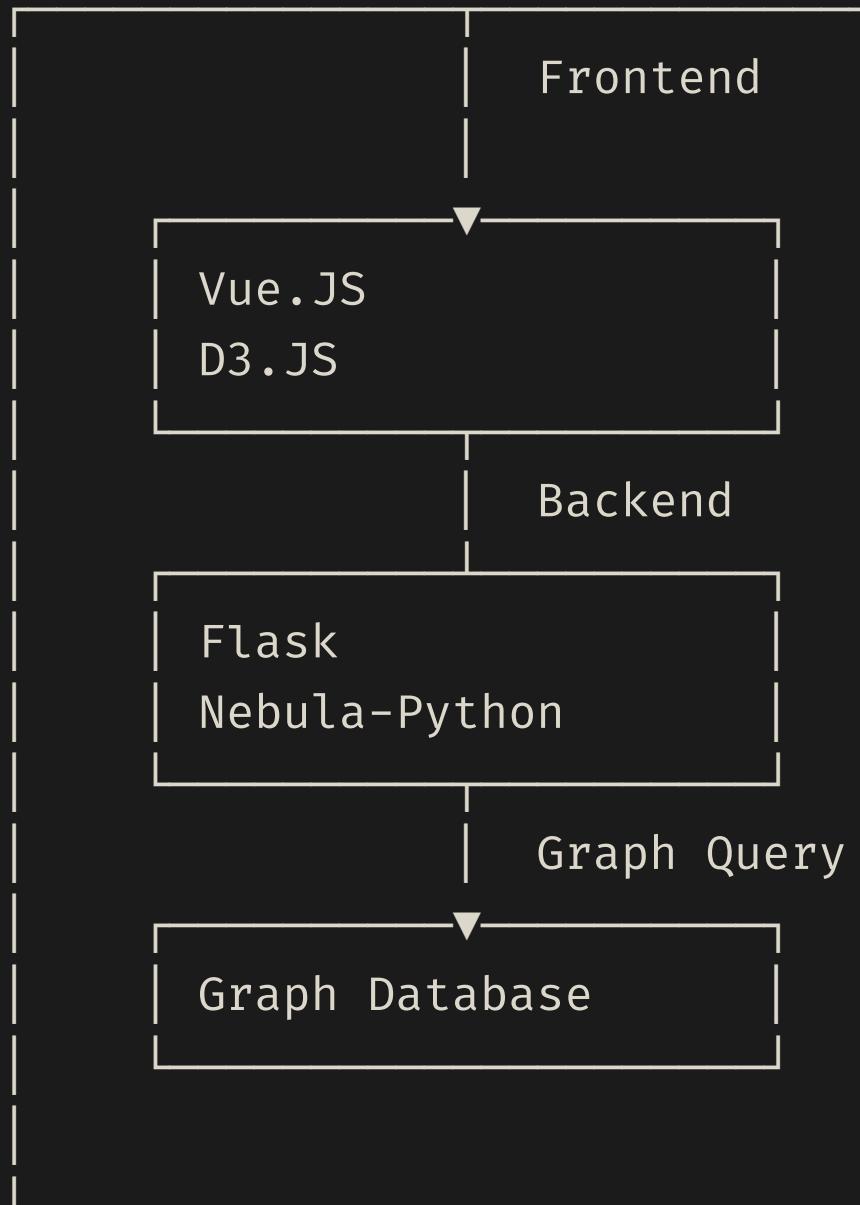
The Code

```
├── README.md      # You could find Design Logs here
├── corp-rel-backend
│   └── app.py      # Flask App to handle Request and call
├── corp-rel-frontend
│   └── src
│       ├── App.vue
│       └── main.js    # Vue App to call Flask App and Render
└── requirements.txt
```

```
@app.route("/api", methods=["POST"])
def api():
    entity = request.get_json().get("entity", "")
    if entity:
        resp = query_shareholding(entity)
        data = make_graph_response(resp)
        #
    return jsonify(data)

def query_shareholding(entity):
    query_string = (
        f"MATCH p=(v)-[e:hold_share|:is_branch_of|:reletive"
        f"WHERE id(v) IN [{entity}] RETURN p LIMIT 100"
    )
    session = connection_pool.get_session('root', 'nebula')
    resp = session.execute(query_string)
    return resp
```

架构 Corp-Rel-Search



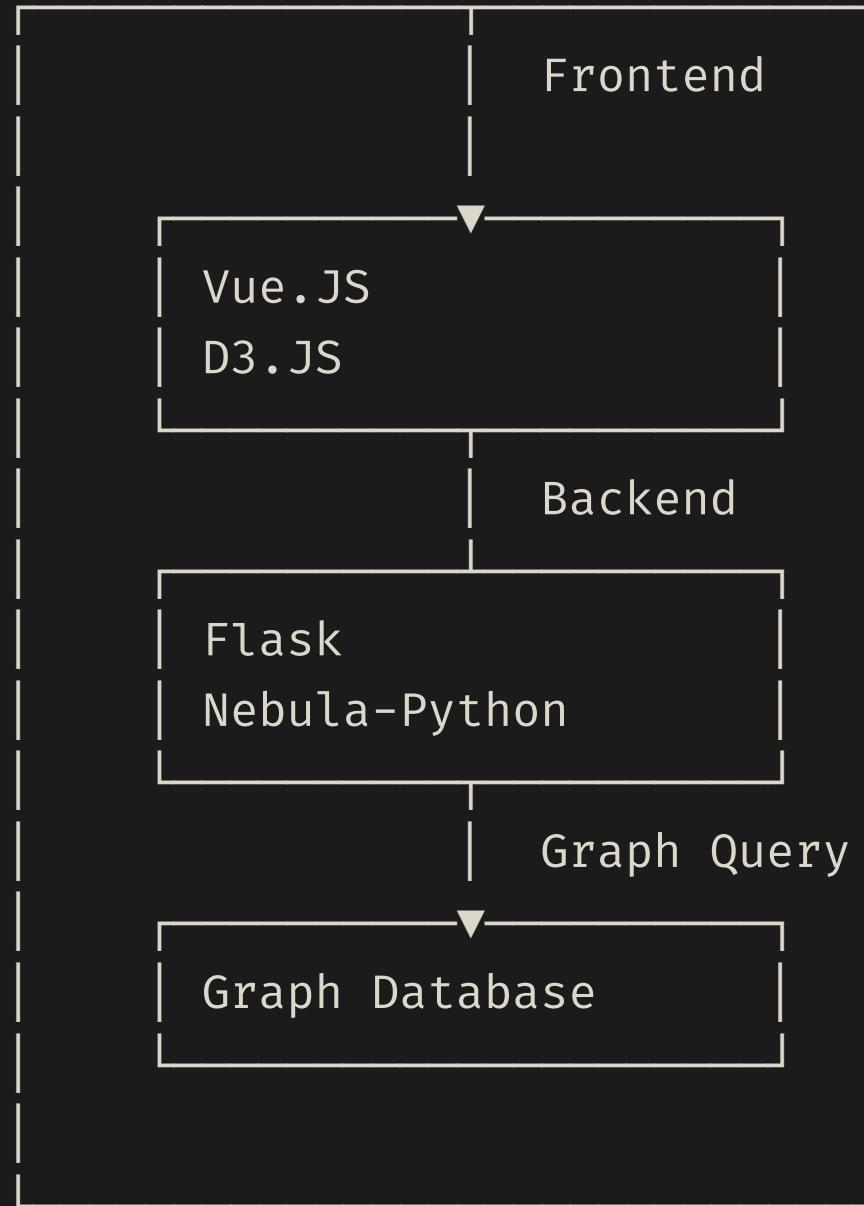
The Code

```
├── README.md      # You could find Design Logs here
├── corp-rel-backend
│   └── app.py      # Flask App to handle Request and call
├── corp-rel-frontend
│   └── src
│       ├── App.vue
│       └── main.js    # Vue App to call Flask App and Render
└── requirements.txt
```

```
@app.route("/api", methods=["POST"])
def api():
    entity = request.get_json().get("entity", "")
    if entity:
        resp = query_shareholding(entity)
        data = make_graph_response(resp)
        #
    return jsonify(data)

def query_shareholding(entity):
    query_string = (
        f"MATCH p=(v)-[e:hold_share|:is_branch_of|:reletive"
        f"WHERE id(v) IN [{entity}] RETURN p LIMIT 100"
    )
    session = connection_pool.get_session('root', 'nebula')
    resp = session.execute(query_string)
    return resp
```

架构 Corp-Rel-Search



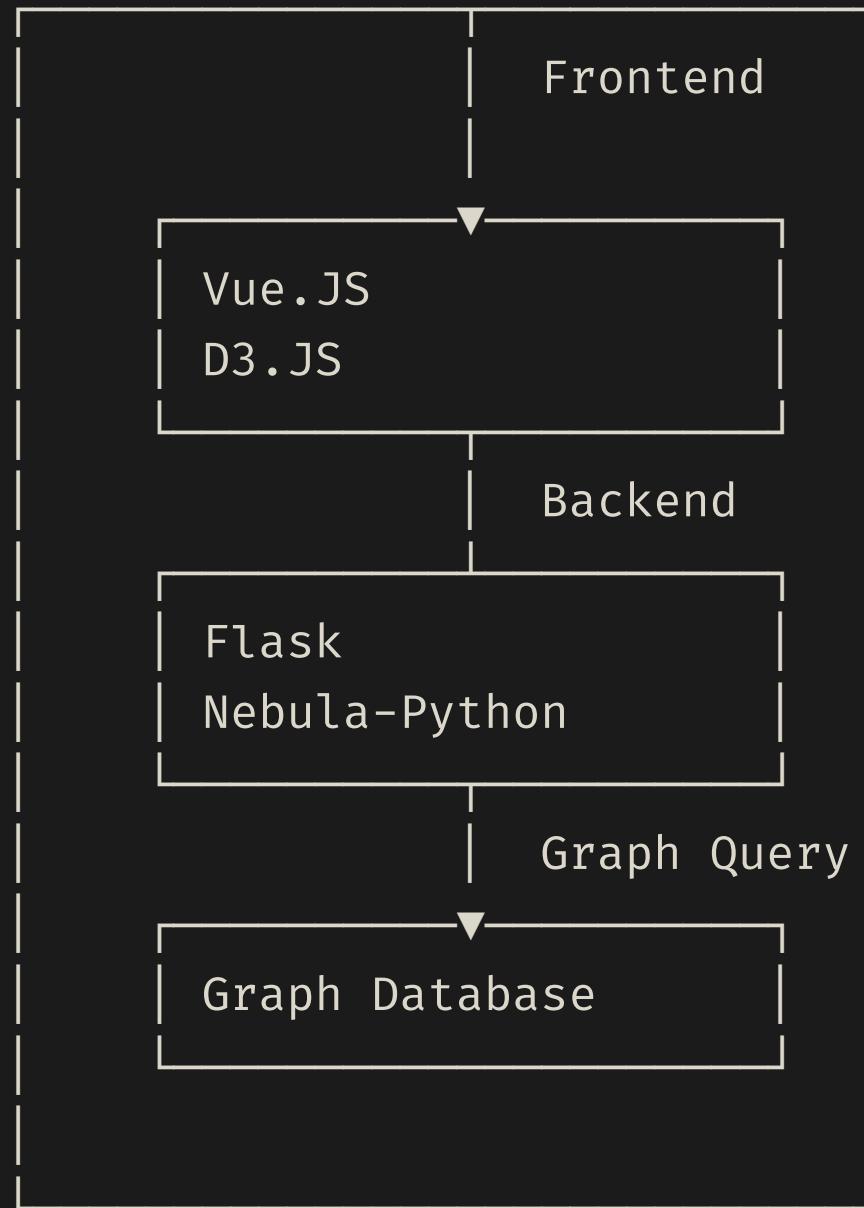
The Code

```
├── README.md      # You could find Design Logs here
├── corp-rel-backend
│   └── app.py      # Flask App to handle Request and call
├── corp-rel-frontend
│   └── src
│       ├── App.vue
│       └── main.js    # Vue App to call Flask App and Render
└── requirements.txt
```

```
@app.route("/api", methods=["POST"])
def api():
    entity = request.get_json().get("entity", "")
    if entity:
        resp = query_shareholding(entity)
        data = make_graph_response(resp)
        #
    return jsonify(data)

def query_shareholding(entity):
    query_string = (
        f"MATCH p=(v)-[e:hold_share|:is_branch_of|:relative"
        f"WHERE id(v) IN [{entity}] RETURN p LIMIT 100"
    )
    session = connection_pool.get_session('root', 'nebula')
    resp = session.execute(query_string)
    return resp
```

架构 Corp-Rel-Search



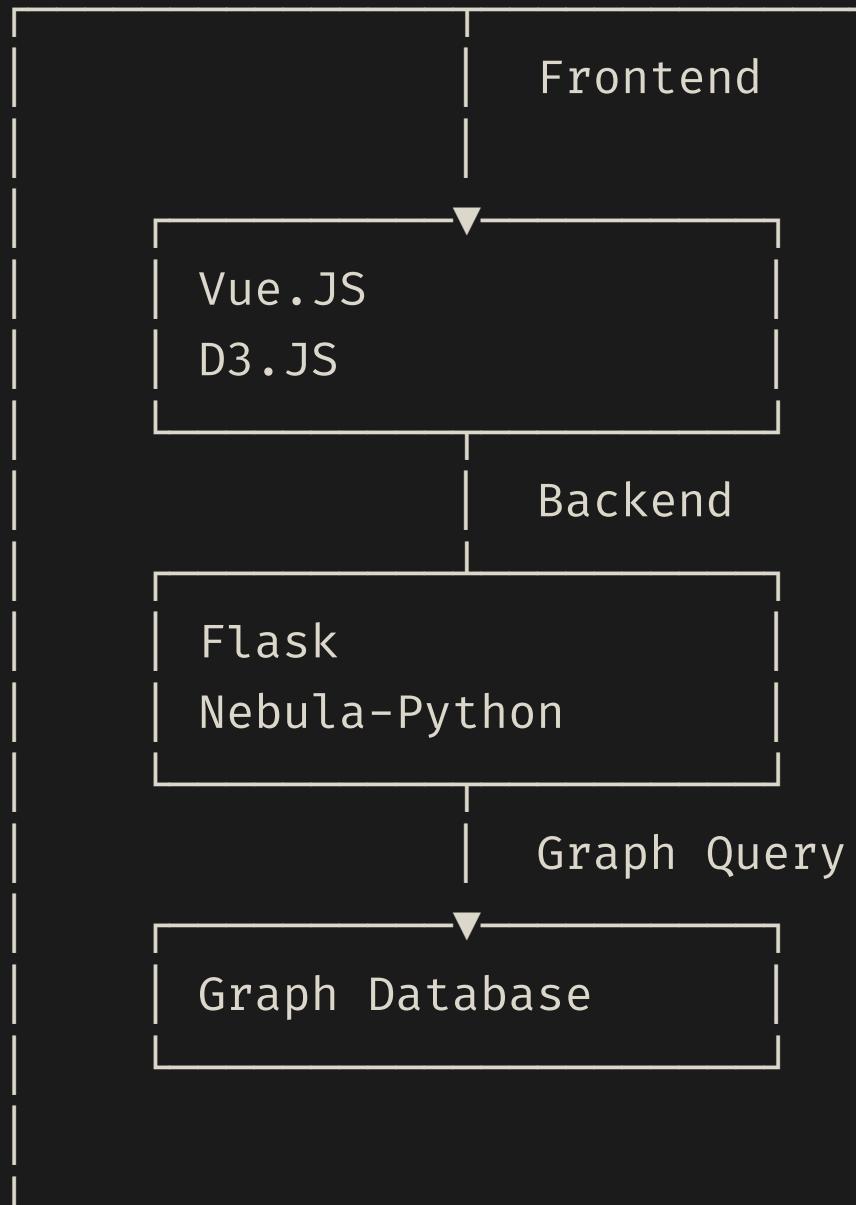
The Code

```
├── README.md      # You could find Design Logs here
├── corp-rel-backend
│   └── app.py      # Flask App to handle Request and call
├── corp-rel-frontend
│   └── src
│       ├── App.vue
│       └── main.js    # Vue App to call Flask App and Render
└── requirements.txt
```

```
@app.route("/api", methods=["POST"])
def api():
    entity = request.get_json().get("entity", "")
    if entity:
        resp = query_shareholding(entity)
        data = make_graph_response(resp)
        # ...
    return jsonify(data)

def query_shareholding(entity):
    query_string = (
        f"MATCH p=(v)-[e:hold_share|:is_branch_of|:relative"
        f"WHERE id(v) IN [{entity}] RETURN p LIMIT 100"
    )
    session = connection_pool.get_session('root', 'nebula')
    resp = session.execute(query_string)
    return resp
```

架构 Corp-Rel-Search



The Code

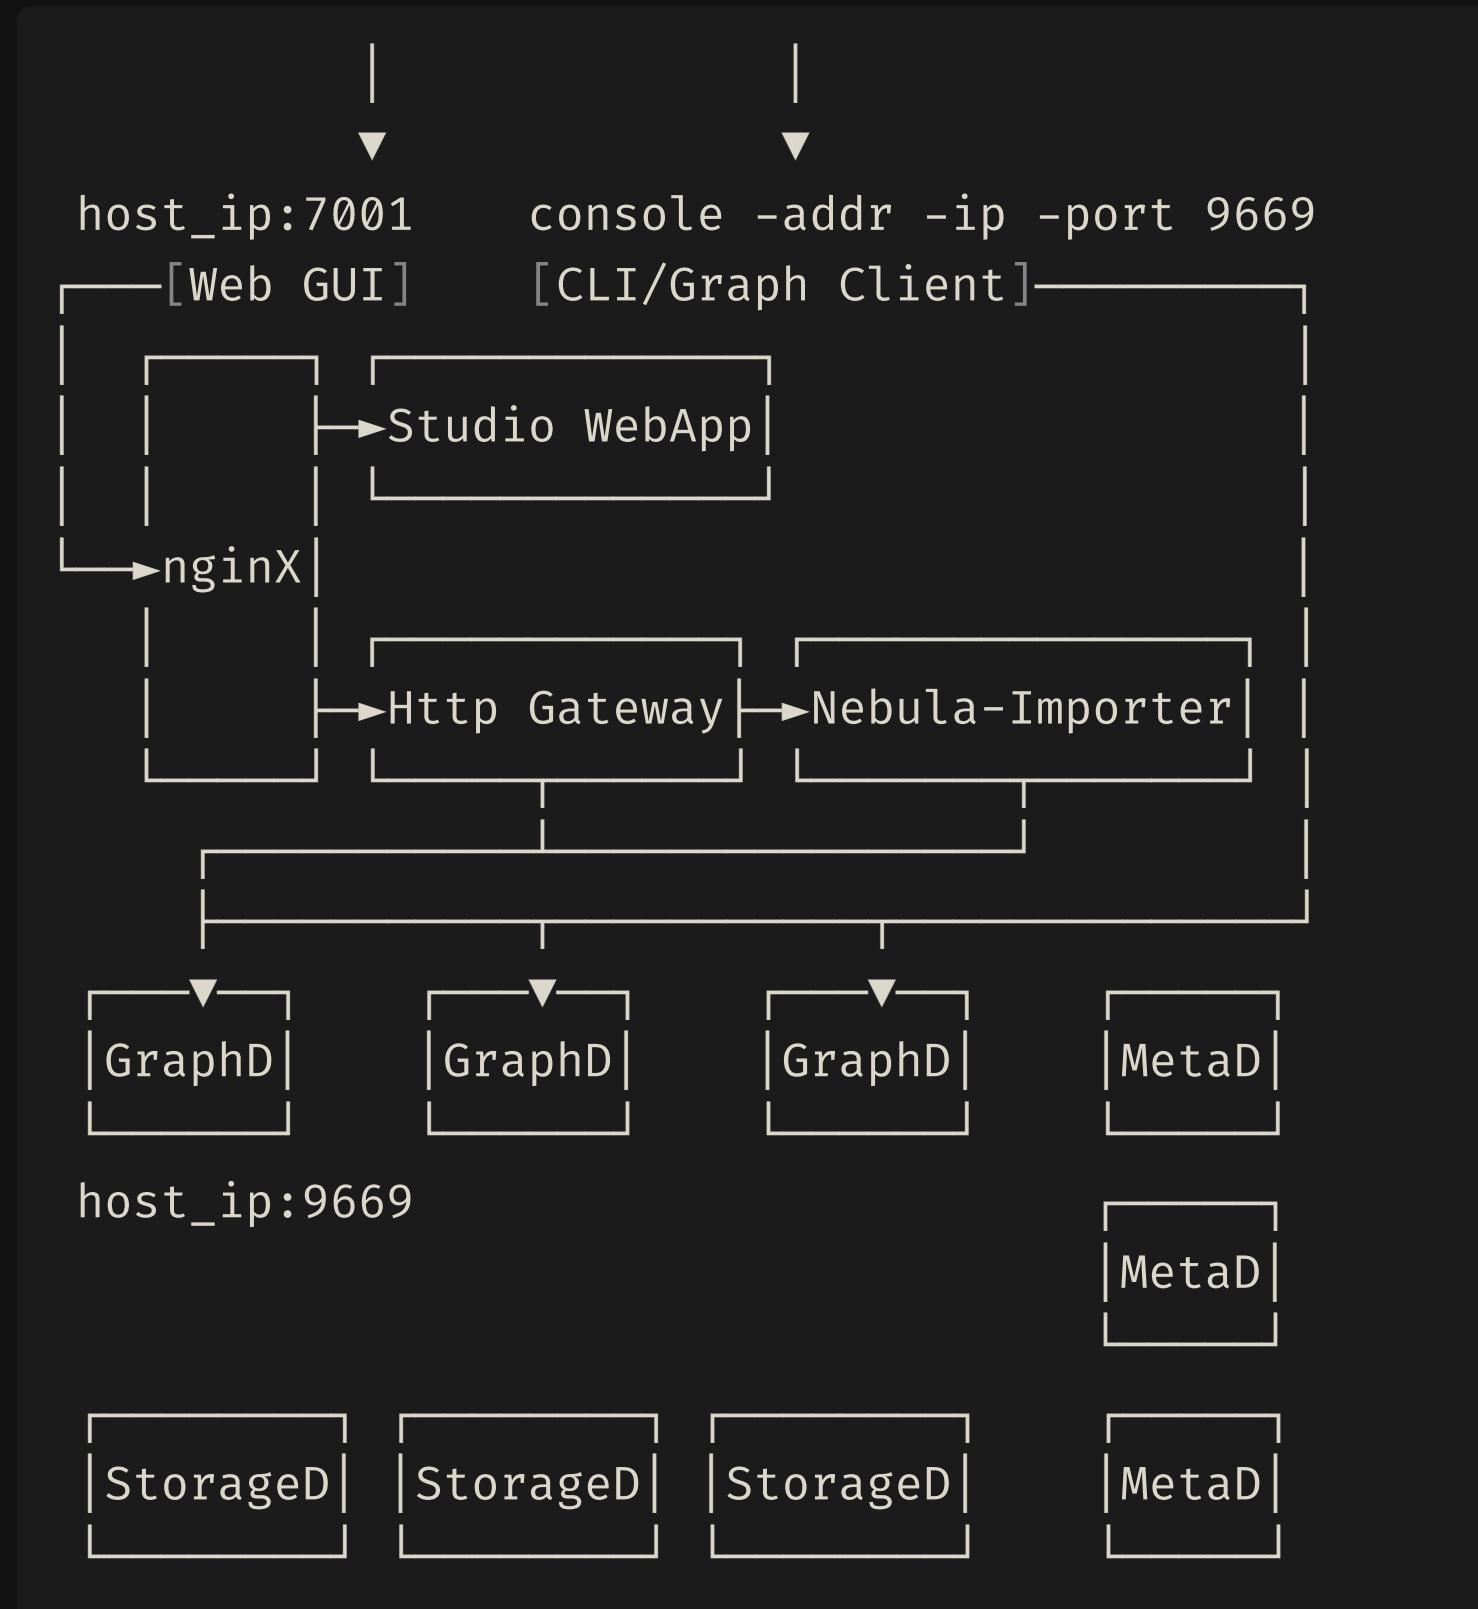
```
├── README.md      # You could find Design Logs here
├── corp-rel-backend
│   └── app.py      # Flask App to handle Request and call
├── corp-rel-frontend
│   └── src
│       ├── App.vue
│       └── main.js    # Vue App to call Flask App and Render
└── requirements.txt
```

```
@app.route("/api", methods=["POST"])
def api():
    entity = request.get_json().get("entity", "")
    if entity:
        resp = query_shareholding(entity)
        data = make_graph_response(resp)
        # ...
    return jsonify(data)

def query_shareholding(entity):
    query_string = (
        f"MATCH p=(v)-[e:hold_share|:is_branch_of|:reletive"
        f"WHERE id(v) IN [{entity}] RETURN p LIMIT 100"
    )
    session = connection_pool.get_session('root', 'nebula')
    resp = session.execute(query_string)
    return resp
```

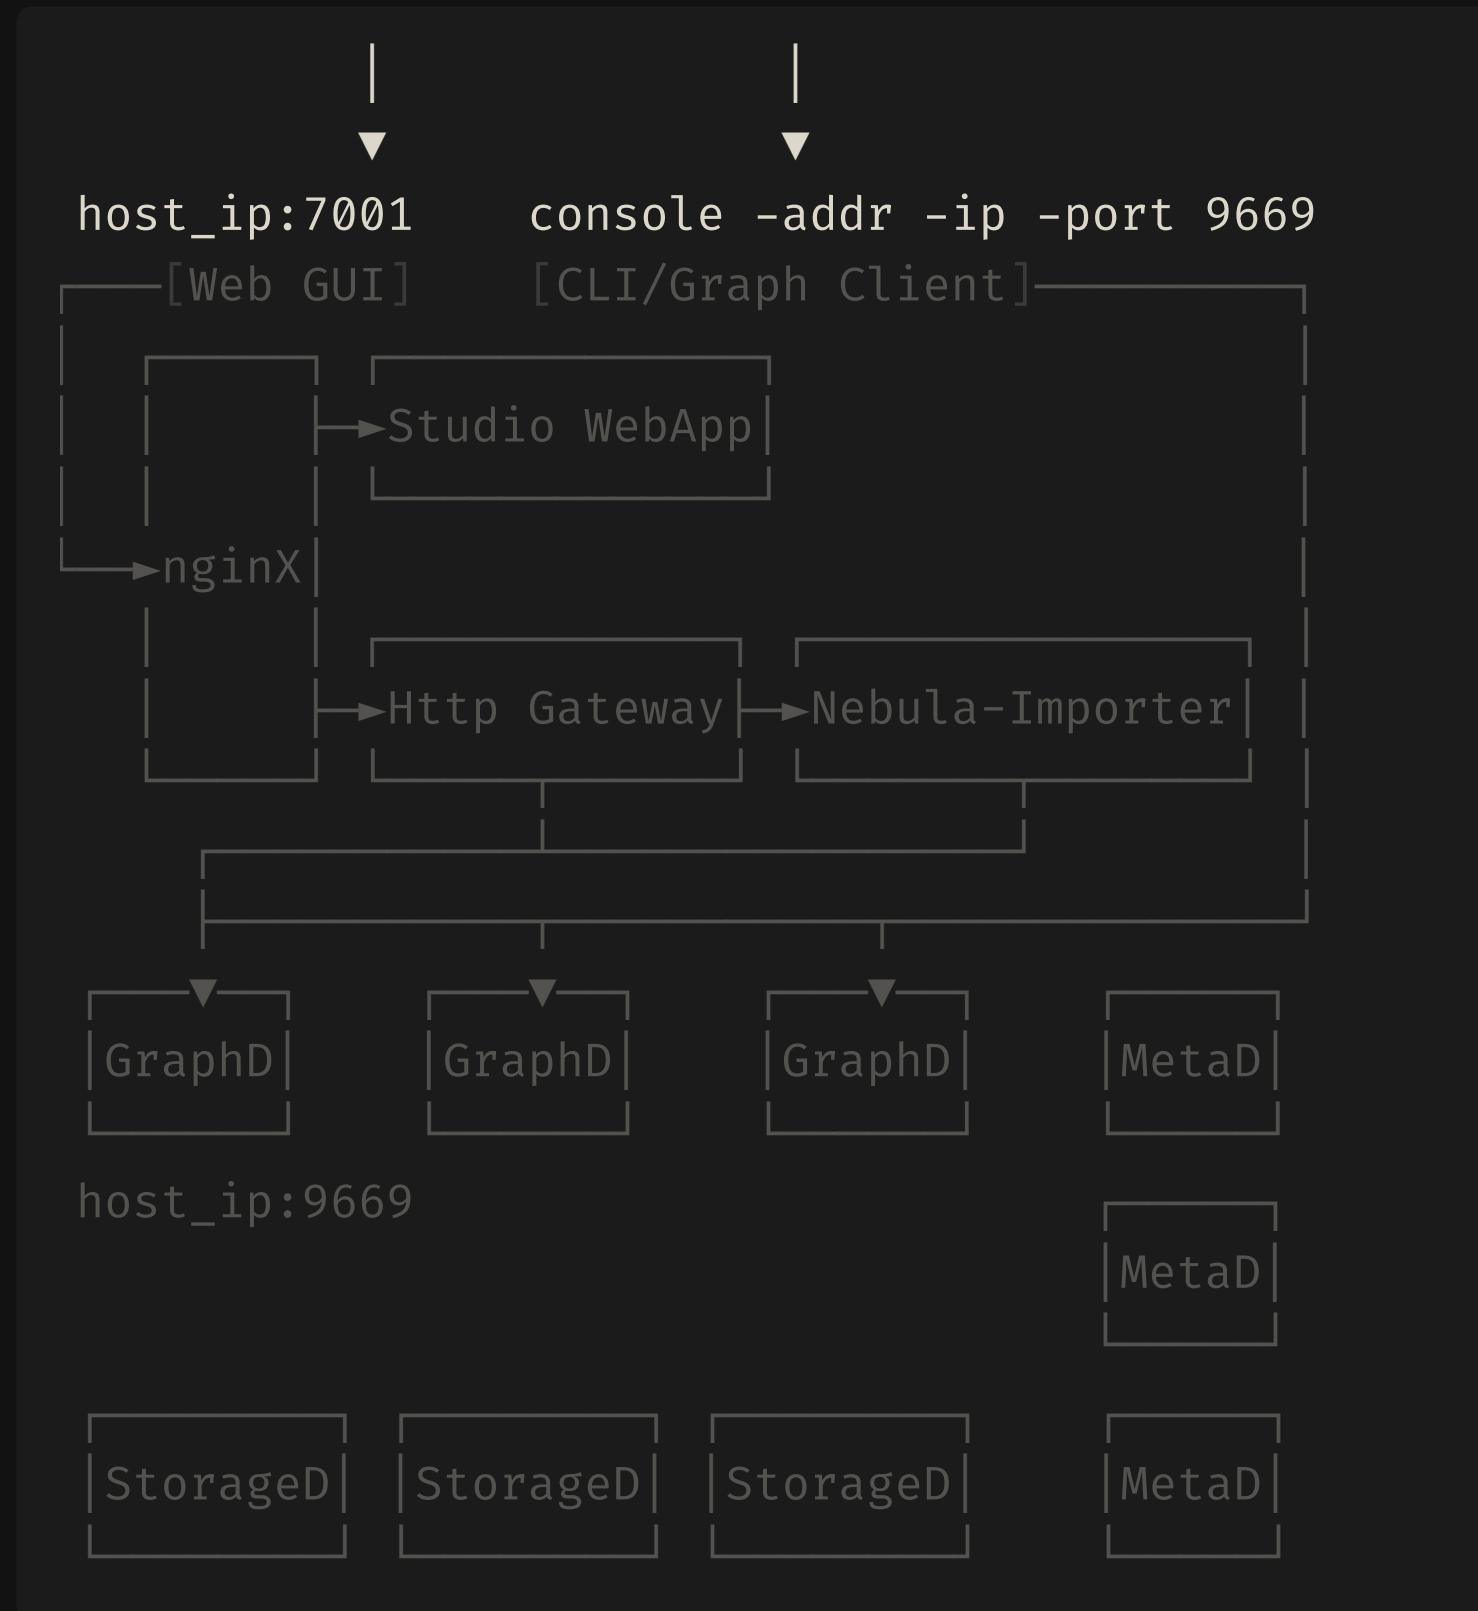
Nebula Graph in Docker and Dataset Import

```
curl -fsSL nebula-up.siwei.io/install.sh | bash
```



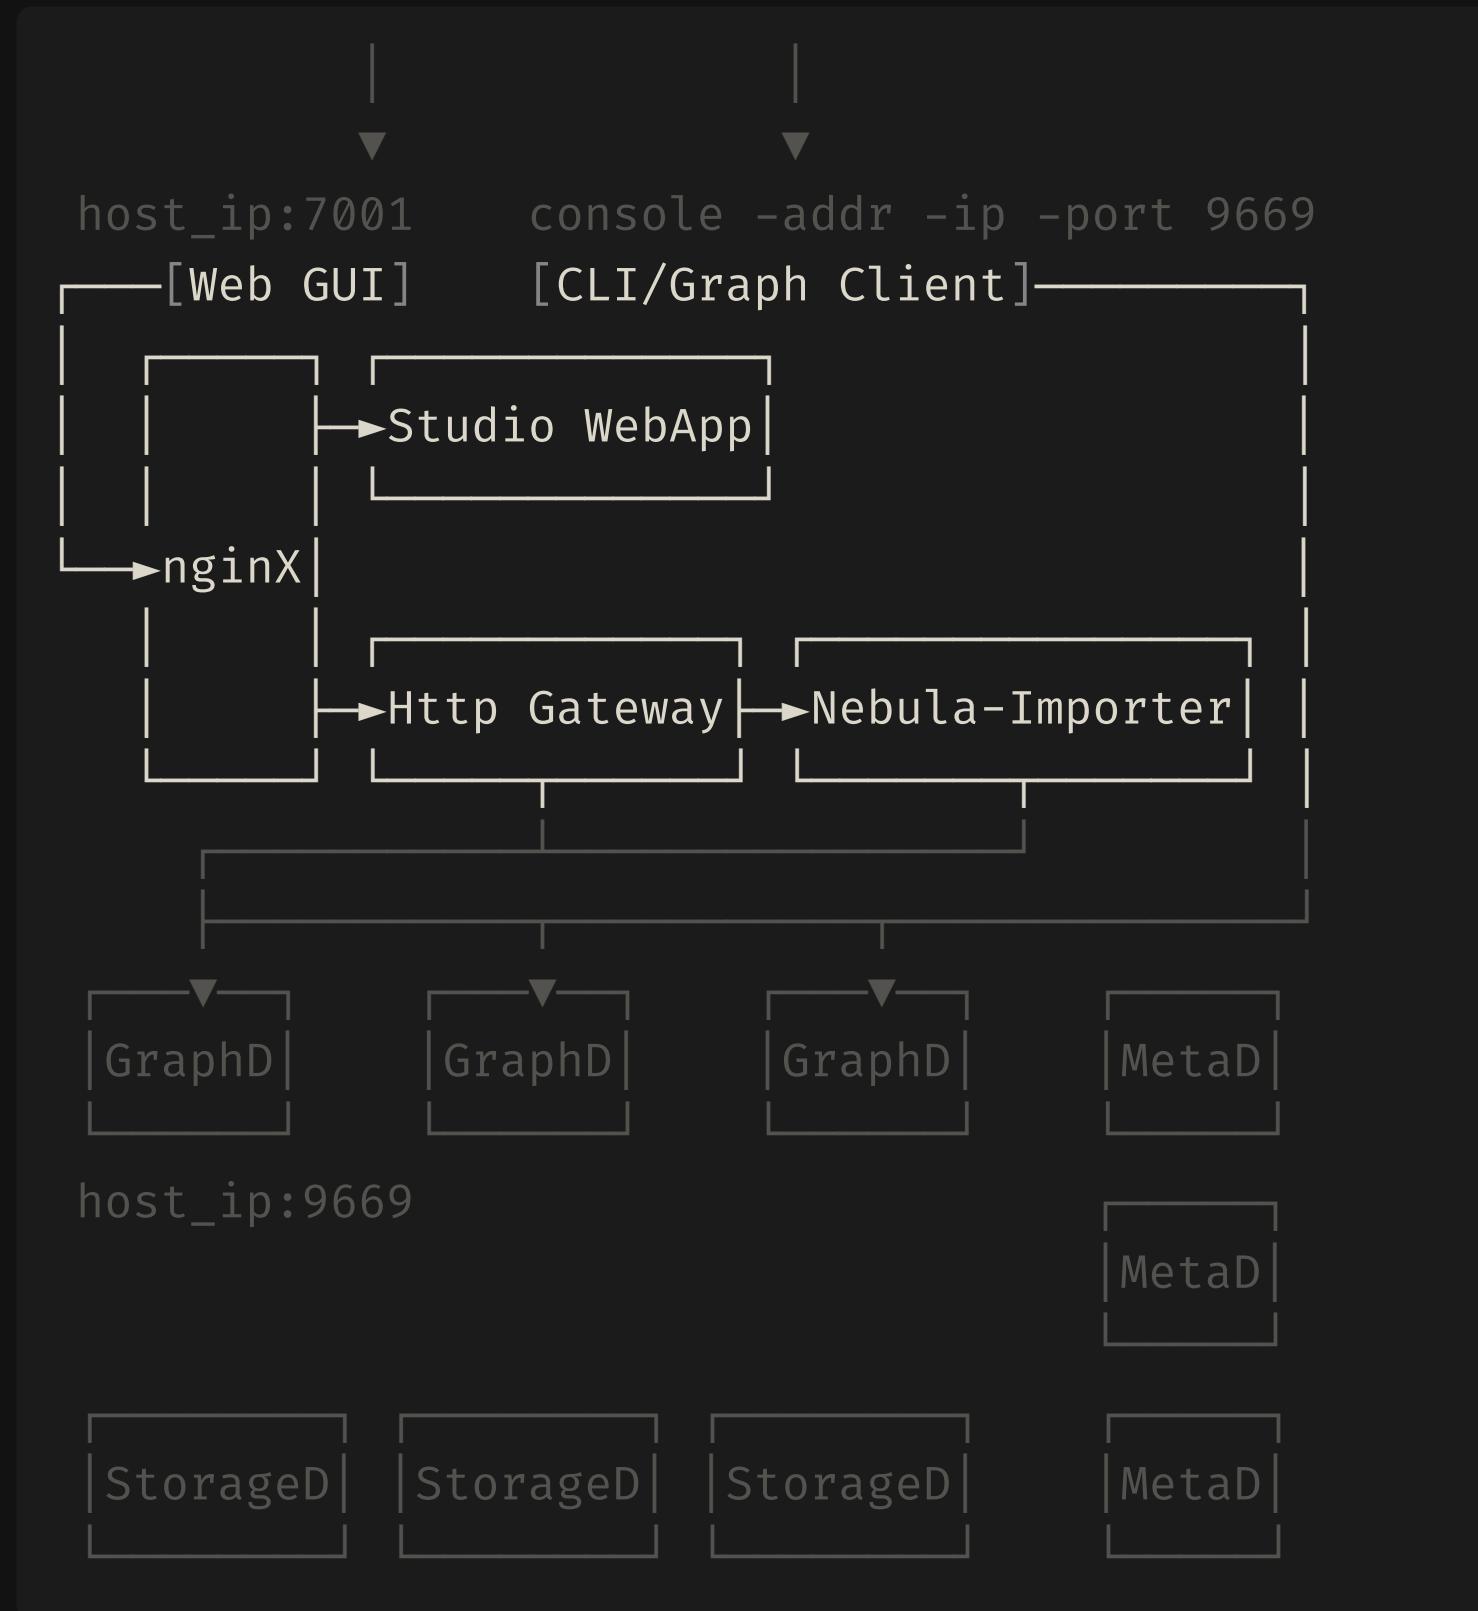
Nebula Graph in Docker and Dataset Import

```
curl -fsSL nebula-up.siwei.io/install.sh | bash
```



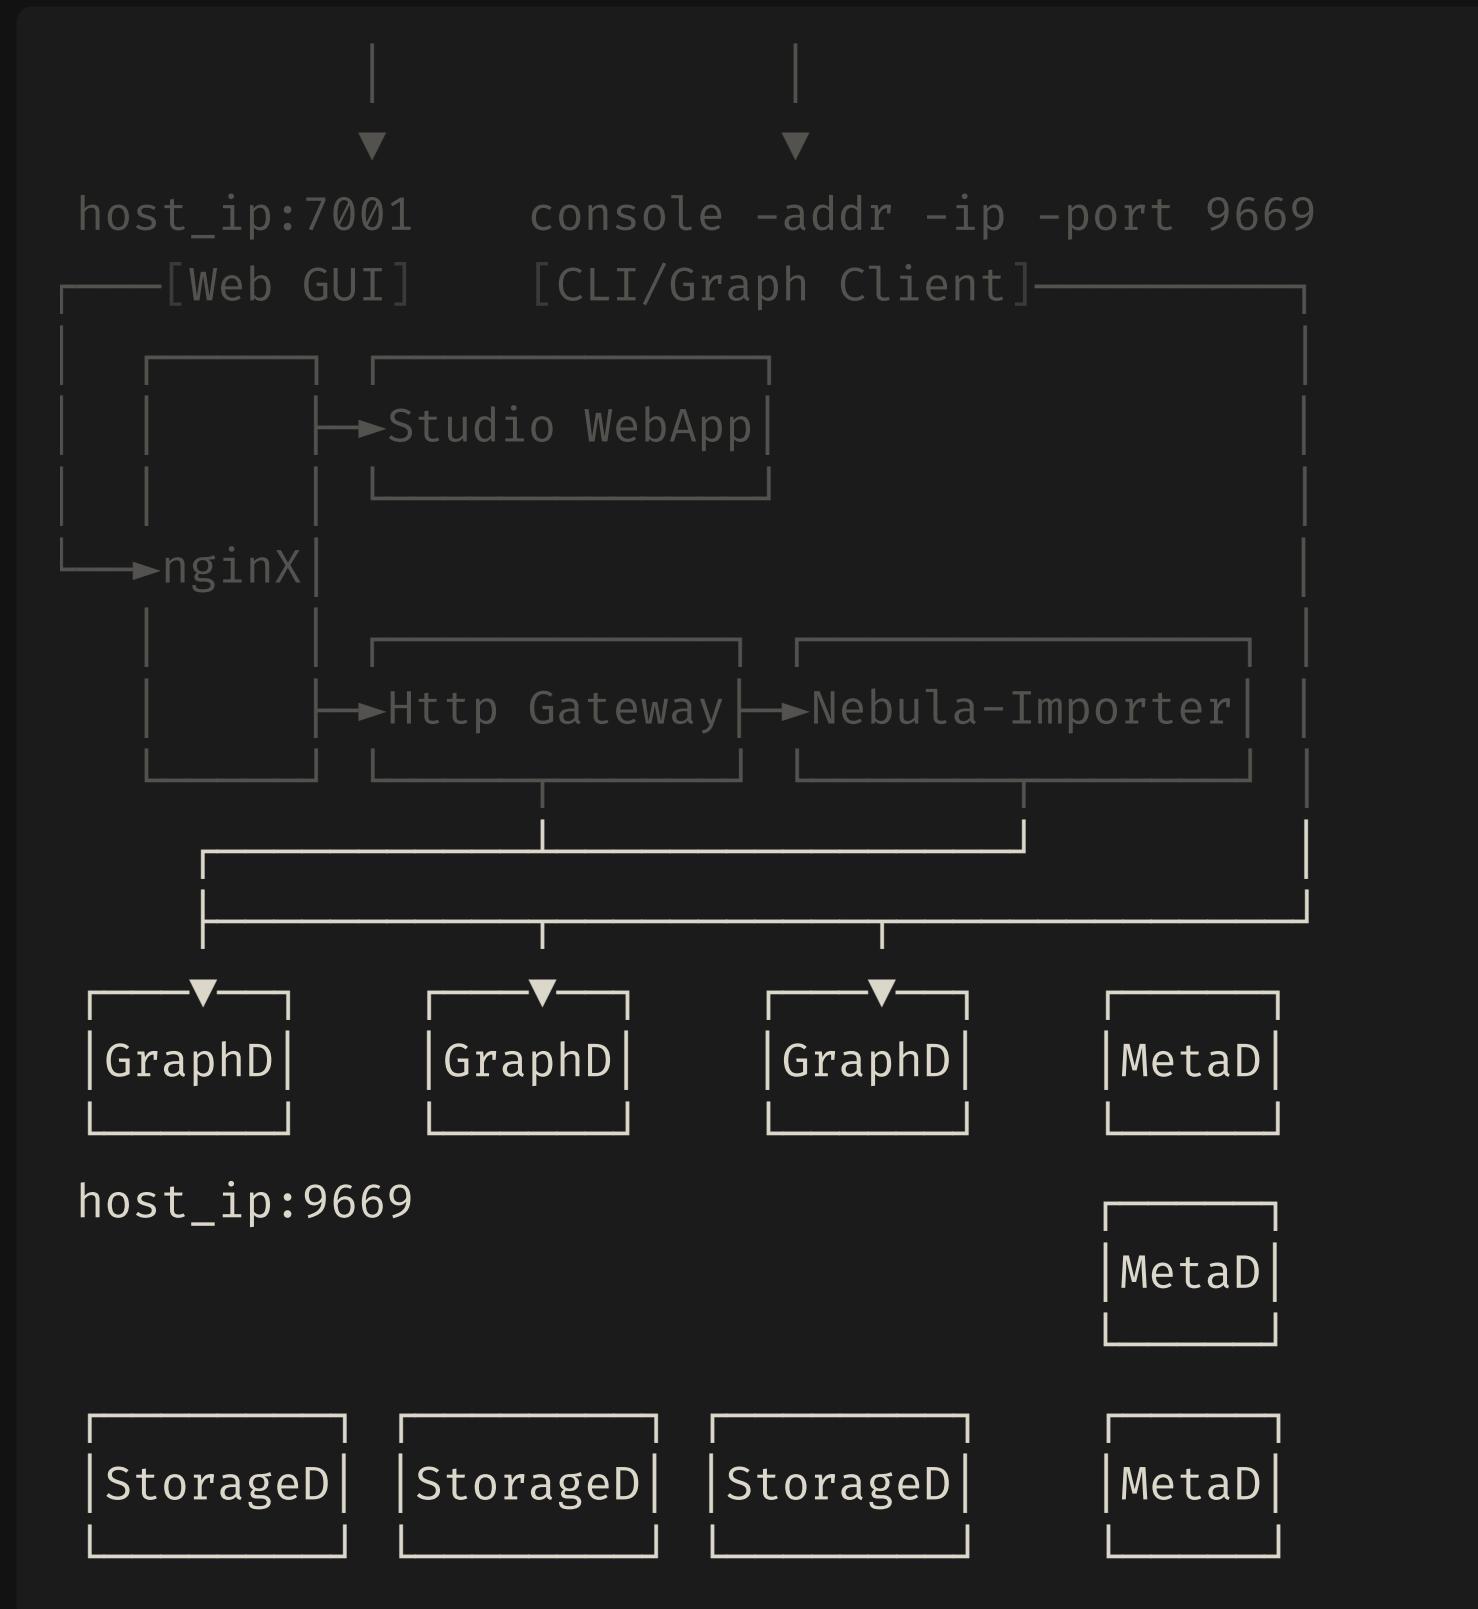
Nebula Graph in Docker and Dataset Import

```
curl -fsSL nebula-up.siwei.io/install.sh | bash
```



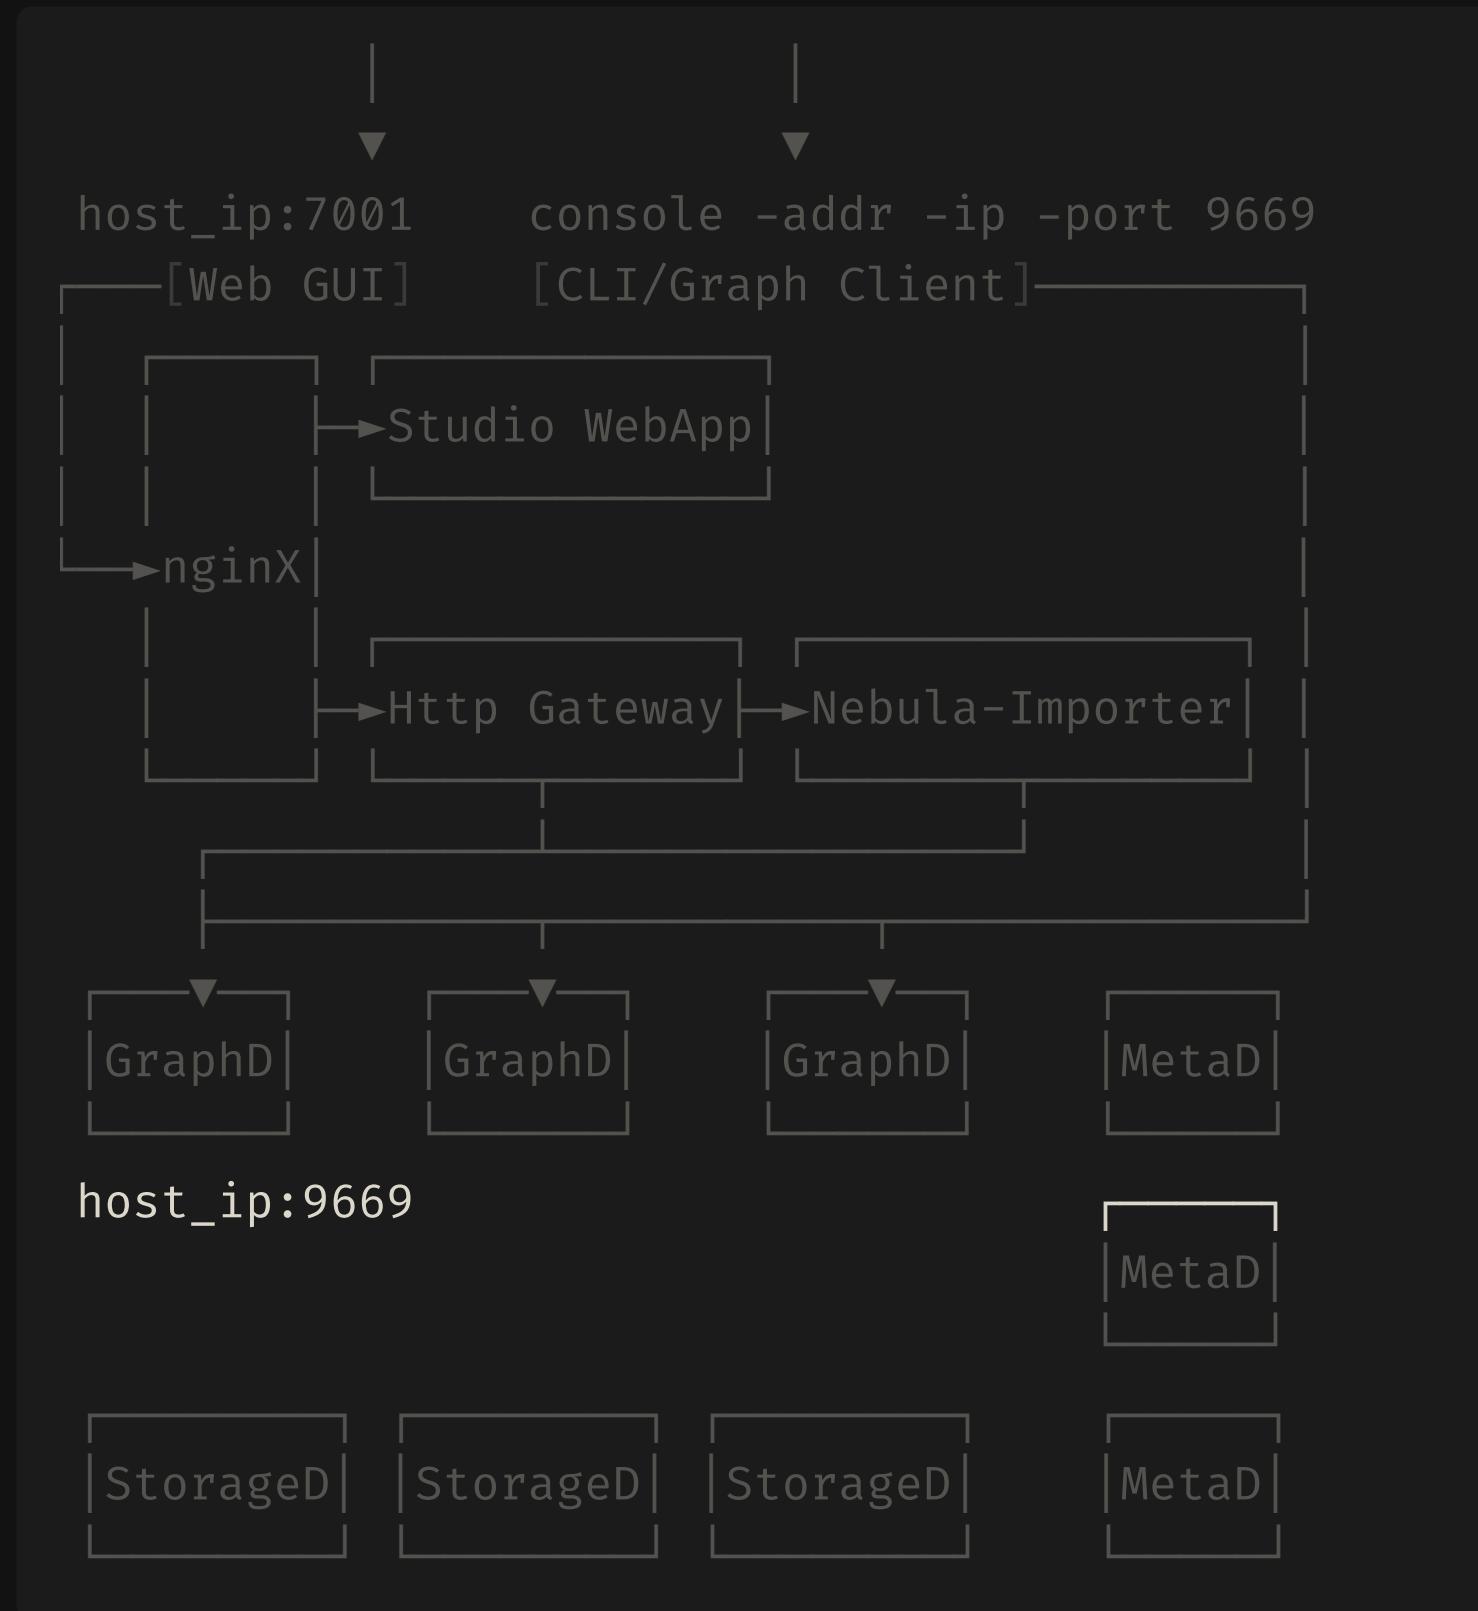
Nebula Graph in Docker and Dataset Import

```
curl -fsSL nebula-up.siwei.io/install.sh | bash
```



Nebula Graph in Docker and Dataset Import

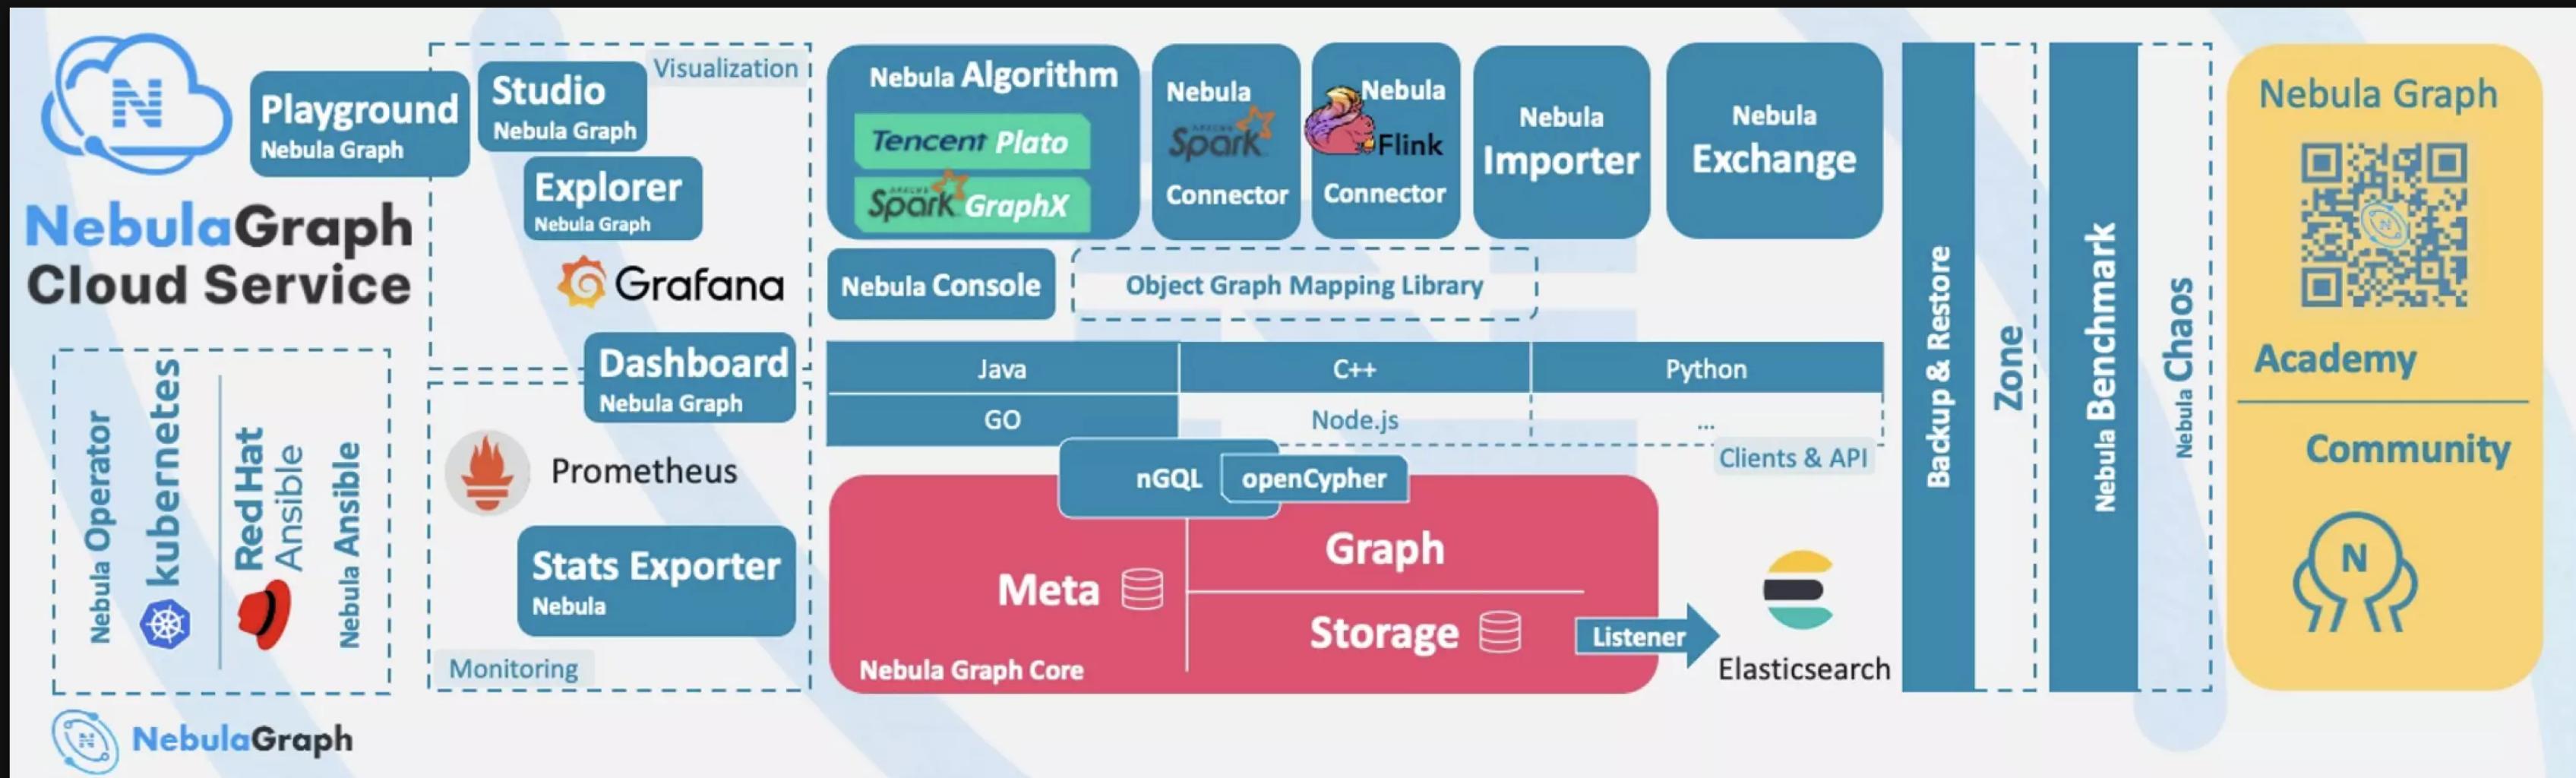
```
curl -fsSL nebula-up.siwei.io/install.sh | bash
```



Nebula Landscape

Nebula 社区生态非常丰富，并且还在日益拓展，欢迎同学们了解、参与贡献。

- Deployment, Monitoring
- Data Visualization
- Algorithm, Analytic
- Clients, Connectors, ETL



为什么选择 Nebula?

为什么选择 Nebula?

- ❤️ Open Source

为什么选择 Nebula?

- ❤️ Open Source
- 🚀 生态、性能、只有它能 hold 得住、拥抱标准、领先

为什么选择 Nebula?

- ❤️ Open Source
- 🚶 生态、性能、只有它能 hold 得住、拥抱标准、领先
- 🐋 Scale matters, Nebula Nailed it! Nebula 擅长超大规模图数据处理

① x-lab/2020 开源报告

① db-engines.com/en/ranking/graph+dbms



Wey Gu 古思为
@wey_gu



Scale makes differences
@NebulaGraph



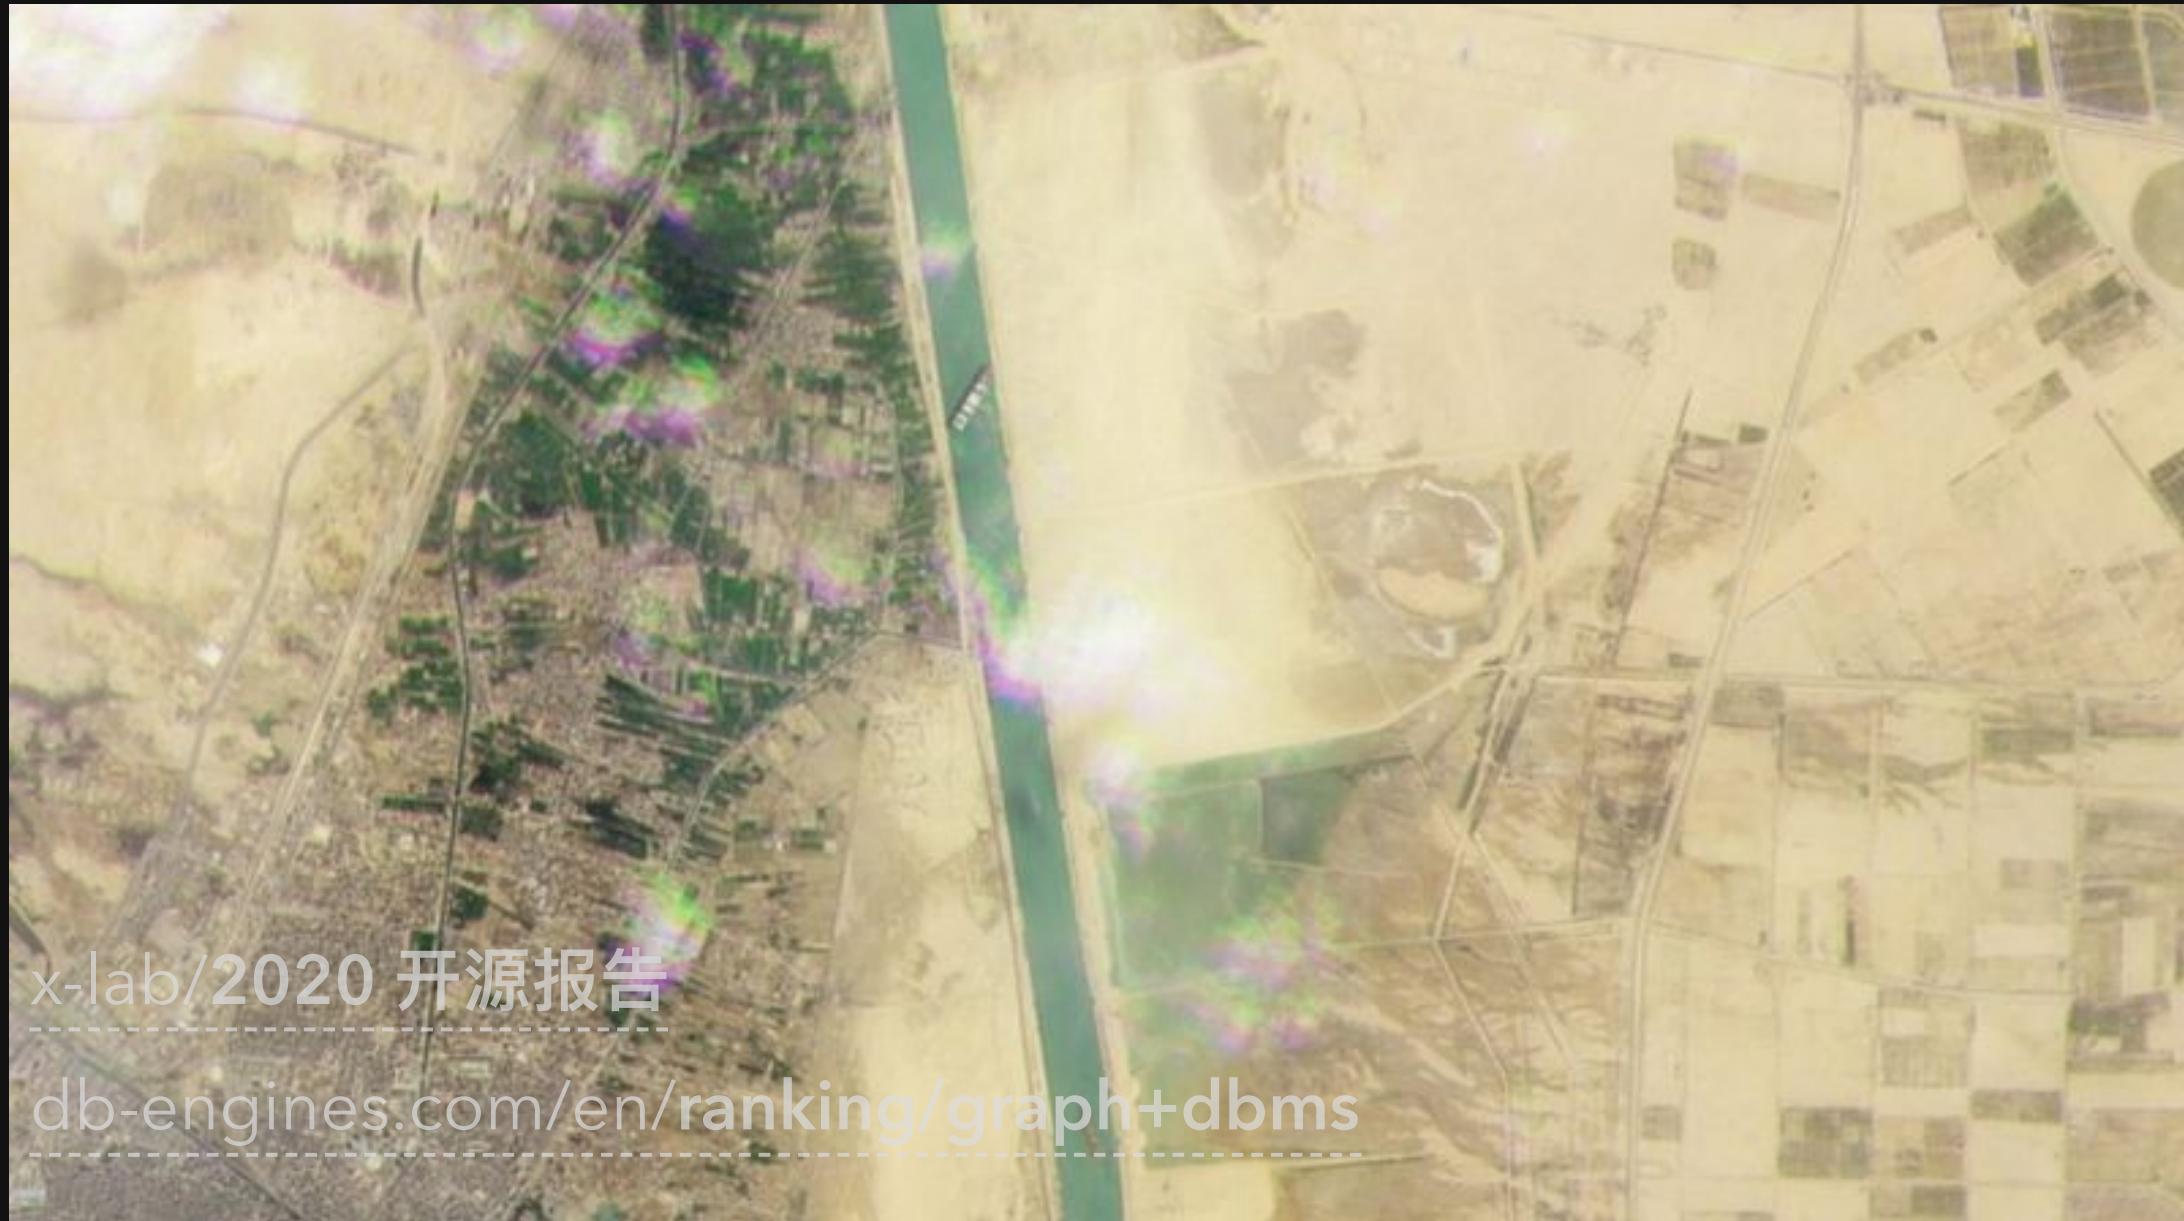
Martin Beeby @thebeebs

Small and medium businesses may have similar problems to big tech companies. But should never assume they need to use the same solution. The answer is not always a shiny new library, a complicated frontend build process, or the latest orchestrator. It is all a matter of scale.



为什么选择 Nebula?

- ❤️ Open Source
- 🚶 生态、性能、只有它能 hold 得住、拥抱标准、领先
- 🐋 Scale matters, Nebula Nailed it! Nebula 擅长超大规模图数据处理



Wey Gu 古思为
@wey_gu



Scale makes differences
@NebulaGraph



Martin Beeby @thebeebs

Small and medium businesses may have similar problems to big tech companies. But should never assume they need to use the same solution. The answer is not always a shiny new library, a complicated frontend build process, or the latest orchestrator. It is all a matter of scale.



为什么选择 Nebula?

- ❤️ Open Source
- 🚶 生态、性能、只有它能 hold 得住、拥抱标准、领先
- 🐋 Scale matters, Nebula Nailed it! Nebula 擅长超大规模图数据处理



Wey Gu 古思为
@wey_gu



Scale makes differences
@NebulaGraph



Martin Beeby @thebeebs

Small and medium businesses may have similar problems to big tech companies. But should never assume they need to use the same solution. The answer is not always a shiny new library, a complicated frontend build process, or the latest orchestrator. It is all a matter of scale.



为什么选择 Nebula?

- ❤️ Open Source
- 🚶 生态、性能、只有它能 hold 得住、拥抱标准、领先
- 🐋 Scale matters, Nebula Nailed it! Nebula 擅长超大规模图数据处理



Wey Gu 古思为
@wey_gu



Scale makes differences
@NebulaGraph



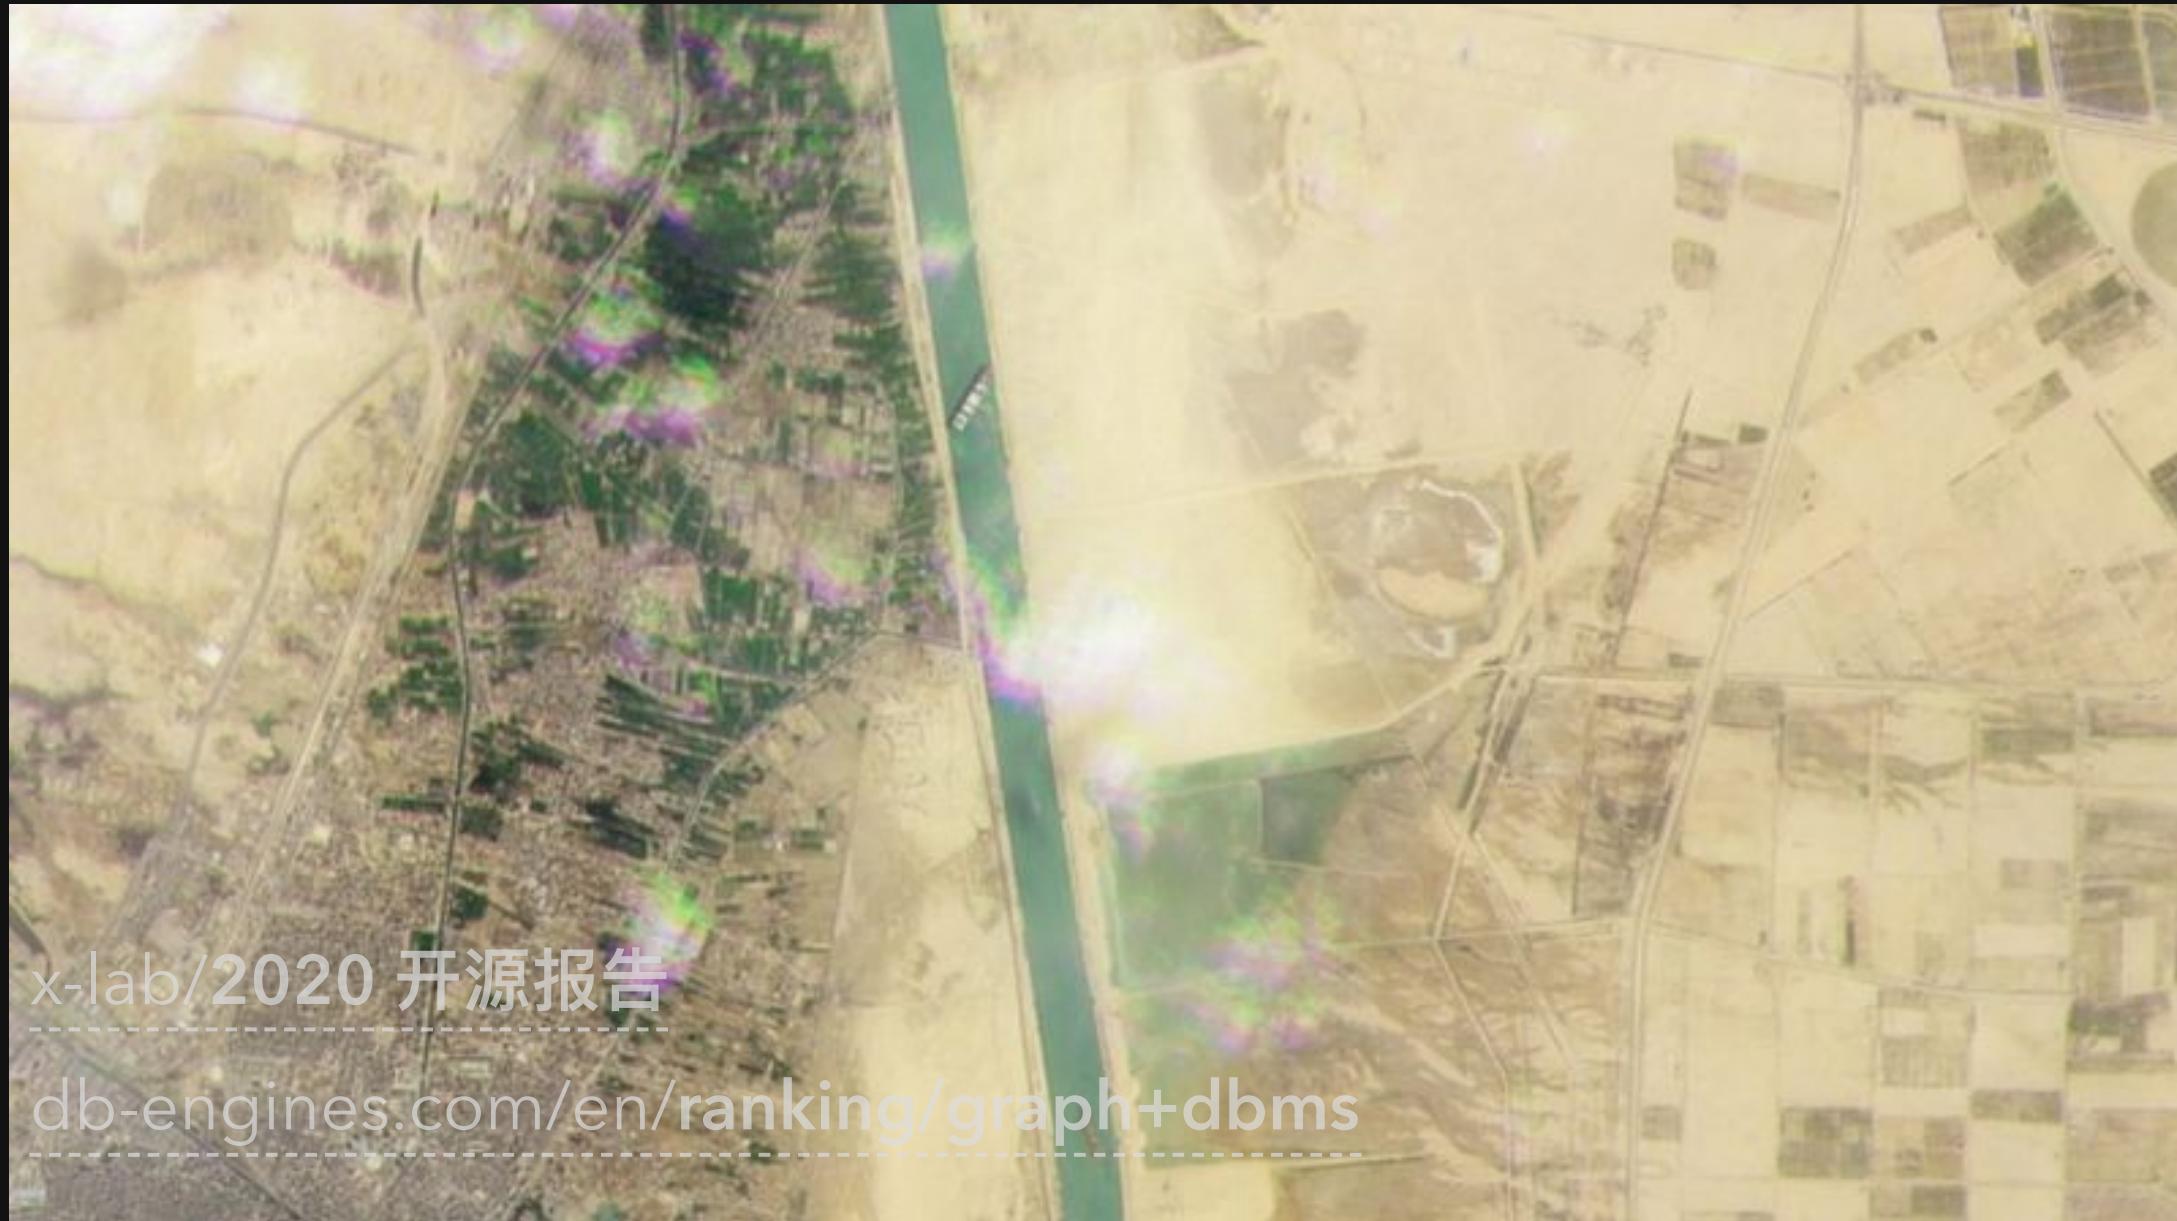
Martin Beeby @thebeebs

Small and medium businesses may have similar problems to big tech companies. But should never assume they need to use the same solution. The answer is not always a shiny new library, a complicated frontend build process, or the latest orchestrator. It is all a matter of scale.



为什么选择 Nebula?

- ❤️ Open Source
- 🚶 生态、性能、只有它能 hold 得住、拥抱标准、领先
- 🐋 Scale matters, Nebula Nailed it! Nebula 擅长超大规模图数据处理



Wey Gu 古思为
@wey_gu



Scale makes differences
@NebulaGraph



Martin Beeby @thebeebs

Small and medium businesses may have similar problems to big tech companies. But should never assume they need to use the same solution. The answer is not always a shiny new library, a complicated frontend build process, or the latest orchestrator. It is all a matter of scale.



回顾

siwei.io/talks/2021-modb

回顾

-  基础：什么是图？图数据库？

siwei.io/talks/2021-modb

回顾

-  基础：什么是图？图数据库？
-  疗效：为什么需要图数据库？

siwei.io/talks/2021-modb



线上直播

Dec. 9th, 2021

回顾

-  基础：什么是图？图数据库？
-  疗效：为什么需要图数据库？
-  适应症：到底图数据库可以做什么？

siwei.io/talks/2021-modb



线上直播

Dec. 9th, 2021

回顾

-  基础：什么是图？图数据库？
-  疗效：为什么需要图数据库？
-  适应症：到底图数据库可以做什么？
-  Why Nebula (星云)：处理超大规模图数据、活跃的开源社区、丰富的生态、分布式架构、领先、快速进化、云原生...

siwei.io/talks/2021-modb
